NoC-Lab Research Consortium

Network-on-Chip and Advanced Interconnects

Architecture, Co-Design & Emerging Frontiers

November 6, 2026

Springer Nature

To the researchers, architects, and silicon engineers

whose relentless curiosity pushes the physical boundaries of on-chip communication,

and to the students who will design the trillion-transistor computing systems of tomorrow.

Preface

For nearly four decades, the computing revolution drew its momentum from the harmonic cadence of Moore's Law and Dennard Scaling: transistors shrank, clock frequencies accelerated into the multi-gigahertz regime, and single-threaded execution performance soared without threatening silicon thermal limits. Around 2005, the irreversible breakdown of Dennard scaling halted clock frequency scaling, forcing a historic pivot toward spatial parallelism: multi-core processors, massively parallel graphics processors, and heterogeneous Systems-on-Chip (SoCs).

Yet scaling compute elements from tens to hundreds, and now thousands of cores on a single die or across 2.5D/3D multi-chiplet assemblies, immediately collided with an unyielding physical reality: the **interconnect crisis**. Global shared multi-drop buses collapsed under quadratic RC wire delay and serialization contention; full crossbar matrices exploded in $O(N^2)$ silicon area and routing track congestion; and deep-submicron copper interconnects below 14 nm succumbed to severe electron scattering and resistive degradation.

The **Network-on-Chip (NoC)** paradigm answered this crisis by transposing packet-switched networking principles directly onto silicon. Today, on-chip and die-to-die communication fabrics form the nervous system of modern computing—from mobile application processors and high-performance server SoCs to wafer-scale accelerators and neuromorphic tensor processors.

Scope and Structure of This Book

This textbook offers a rigorous, first-principles exploration of on-chip and multi-chiplet interconnection networks. Rather than treating interconnects as abstract graph topologies, this volume bridges the chasm between physical silicon reality, router microarchitecture, mathematical deadlock theory, and emerging technology frontiers.

The book is organized into five cohesive parts:

- **Part I: Foundations & Router Microarchitecture (Chapters 1–3)** establishes the physical interconnect crisis, mathematical delay formulations, switching techniques, credit and token flow control, and pipelined router microarchitecture with virtual channel allocation and crossbar arbitration.
- **Part II: Topologies, Routing & Deadlock Freedom (Chapters 4–5)** investigates direct, indirect, and high-radix networks, Dally's channel dependency graph (CDG) theorems, Duato's adaptive routing theory, and turn models.
- **Part III: Physical Constraints: Power, Thermal & Reliability (Chapters 6–7)** addresses dark silicon dynamics, voltage/frequency scaling, compact thermal RC modeling, hotspot mitigation, and fault-tolerant routing.
- **Part IV: Advanced Packaging, 3D NoC & Chiplets (Chapters 8–9)** explores through-silicon via (TSV) vertical architectures, 2.5D silicon interposers, and the Universal Chiplet Interconnect Express (UCIe) die-to-die standard.
- **Part V: Emerging Frontiers (Chapter 10)** covers silicon nanophotonic interconnects (ONoC), micro-ring resonators, and on-chip millimeter-wave/sub-THz wireless communication (WNoC).

Grounding in Premier Literature & Silicon Implementations

Every theoretical formulation, architectural trade-off, and silicon case study in this monograph is grounded directly in premier peer-reviewed literature from IEEE and ACM venues, including *IEEE Transactions on Very Large Scale Integration (VLSI) Systems* (TVLSI), *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* (TCAD), *IEEE Journal of Solid-State Circuits* (JSSC), and *IEEE Transactions on Circuits and Systems* (TCAS-I/II). Landmark silicon achievements—such as the 2,048-core HammerBlade SoC, the FlooNoC open-source AXI4 fabric, and the SWIFT token flow-control architecture—are dissected alongside first-principles physics.

Co-Design with the Interactive Companion Hub

To enhance pedagogy and provide hands-on experimentation, this monograph is co-designed with an interactive online platform:

NoC Learning Hub: `https://network-on-chip.github.io/learning/`

Each chapter links directly to interactive browser-based simulators, including the 2D Mesh Routing Visualizer, Arbiter Arena, Credit Commander, and UCIe Protocol Analyzer, accompanied by multi-tier assessment engines and executable Verilog/Python code repositories.

The NoC-Lab Research Consortium
November 2026

Contents

Part I Foundations & Router Microarchitecture

1 The Multi-Core Communication Crisis & Interconnect Physics 3
1.1 The Breakdown of Dennard Scaling and the Multi-Core Pivot . . 3
1.2 The Failure of Shared Multi-Drop Buses 4
1.2.1 Capacitive Loading & Elmore Wire Delay 4
1.2.2 The Serialization Contention Bottleneck 5
1.3 Sub-10nm Interconnect Physics & Wire Scaling 6
1.3.1 The Mayadas-Shatzkes & Fuchs-Sondheimer Scattering Models 6
1.3.2 Sub-3nm Transistor Architectures: GAAFET, RibbonFET, and CFET 7
1.3.3 Backside Power Delivery Networks (BSPDN / PowerVia) 8
1.3.4 Alternative Interconnect Metallization: Ruthenium, Cobalt, and Molybdenum 8
1.3.5 Historical Technology Scaling Across Semiconductor Generations 9
1.4 High-Frequency Transmission Line Effects and Inductive Peaking 10
1.4.1 The Telegrapher's Equations and the Wave Equation . . . 10
1.4.2 High-Frequency Skin Effect and Dielectric Loss 11
1.5 The Monolithic Crossbar and Its Quadratic Area Explosion 11
1.6 Optimal Repeater Insertion Theory (Bakoglu's Formulation) . . . 12
1.7 The Network-on-Chip Paradigm: Packet-Switched Modular Interconnects 13
1.8 Industrial Case Studies: From 48 Cores to Wafer-Scale Engines . 13
1.8.1 The Intel TeraFLOPS 80-Core Processor (65nm CMOS) . 13
1.8.2 The Tilera TILE64 Manycore Processor (90nm CMOS) . 14
1.8.3 The Cerebras Wafer-Scale Engine (WSE-2 & WSE-3) . . . 14
1.8.4 The 2,048-Core HammerBlade Manycore SoC (12nm FinFET) 15

1.9 Interactive Lab & Hands-on Challenges 16
1.10 Checkpoint Assessment 17

2 Switching Techniques, Flow Control & Buffer Management . 19
2.1 The Data Granularity Hierarchy: Messages, Packets, Flits & Phits 19
2.2 Silicon Network Interface: The AXI4 / NoC Protocol Translation 20
2.3 Mathematical Latency Analysis of Switching Techniques 21
2.3.1 Store-and-Forward (SAF) Switching 21
2.3.2 Virtual Cut-Through (VCT) Switching 21
2.3.3 Wormhole (WH) Switching 22
2.4 Queuing Theory Foundations of On-Chip Buffering 23
2.4.1 The $M/M/1/B$ and $M/D/1/B$ Finite Queue Models . . . 23
2.4.2 The Pollaczek-Khinchine Formula for General Service Distributions ($M/G/1$) 24
2.4.3 Open Jackson Networks and Multi-Hop End-to-End Latency 24
2.4.4 Bursty Traffic Modeling: The Markov-Modulated Poisson Process (MMPP) 25
2.5 Lossless Flow Control Mechanisms 26
2.5.1 Credit-Based Flow Control & The Credit Turnaround Loop 26
2.5.2 Flow Control Taxonomy and Hardware Trade-Offs 26
2.6 Buffer Microarchitecture: Input vs. Output vs. Shared DAMQ . . 28
2.7 Alternative Flow Control: Token Flow Control in the SWIFT NoC 28
2.8 Interactive Lab & Hands-on Challenges 29
2.9 Checkpoint Assessment 29

3 High-Performance Router Microarchitecture & Arbitration . 31
3.1 The Canonical Virtual-Channel Router Pipeline 31
3.1.1 Critical Path Timing and Logical Effort Analysis 32
3.2 Wavefront and Wrapped Wavefront Allocators (WFA & WWFA) 33
3.2.1 Systolic Token Ripple Allocation 33
3.2.2 Wrapped Wavefront Allocation (WWFA) 33
3.3 Advanced Pipeline Compression & Router Bypassing 34
3.3.1 Lookahead Routing & Speculative Switch Allocation 34
3.3.2 Buffer Reservation Scheduling: BuffeRS (TCAD 2026) . . 35
3.4 Asynchronous Router Architectures: HyPA (TVLSI 2026) 35
3.5 Single-Cycle Multi-Hop Traversal (SMART / SHARP) 37
3.6 Crossbar Arbitration Theory & Matching Algorithms 37
3.6.1 The Matrix Arbiter: Strict $O(1)$ Hardware LRU 38
3.6.2 McKeown's iSLIP Iterative Matching Algorithm 38
3.6.3 Maximum Weight Matching (MWM) and the Tassiulas-Ephremides Bound 40
3.6.4 Ping-Pong and Dual-Priority Arbiters 40

3.7 Interactive Lab & Hands-on Challenges 41
3.8 Checkpoint Assessment 41

Part II Topologies, Routing & Deadlock Freedom

4 Direct, Indirect & High-Radix Interconnect Topologies 45
4.1 Interconnect Graph Theory and Formal Performance Metrics . . 45
4.1.1 Spectral Graph Theory and Algebraic Connectivity 46
4.1.2 Classical Topological Performance Metrics 46
4.2 Direct Topologies: Mesh, Torus, and the Folded Torus 48
4.2.1 2D Mesh Networks: First-Principles Hop Derivations 48
4.2.2 Channel Load Under Adversarial Traffic Permutations . . 48
4.2.3 2D Torus and the Wire Delay Dilemma 49
4.2.4 The Folded Torus Invariance 49
4.3 Indirect Multistage Interconnects: Clos and Fat-Trees 50
4.3.1 Clos Networks and Non-Blocking Theorems 50
4.3.2 Fat-Trees and Constant Bisection Bandwidth 50
4.4 High-Radix Interconnects: Flattened Butterfly, Dragonfly, and Extremal Graphs 50
4.4.1 Flattened Butterfly Topologies 51
4.4.2 The Dragonfly Topology and Hierarchical Grouping 52
4.4.3 Moore Bound, Slim Fly, and PolarFly Extremal Topologies 54
4.5 Physical Packaging Hierarchy and Multi-Die Interconnects 55
4.6 Industrial High-Radix Interconnect Case Studies 56
4.6.1 Case Study 1: NVIDIA NVSwitch Fabric 56
4.6.2 Case Study 2: The Cray Slingshot Interconnect 57
4.7 Application-Specific Synthesis & Cascaded Crossbars 57
4.7.1 MIRO Synthesis of Cascaded Crossbar Networks 57
4.7.2 Switch Merger and Partitioning Algorithms 58
4.8 Silicon Floorplanning Limits: The Passas Radix-128 Proof 58
4.9 Quantitative Topology Evaluation Matrix 60
4.10 Interactive Lab & Hands-on Challenges 60
4.11 Exercises and Review Questions 61

5 Deadlock Theory, Turn Models & Adaptive Routing 63
5.1 The Deadlock Problem in Wormhole-Switched Fabrics 63
5.2 The Dally-Seitz Channel Dependency Graph Theorem 64
5.2.1 Formal Definition of the CDG 64
5.2.2 The Theorem 64
5.2.3 Application: Dimension-Order Routing (XY Routing) . . . 65
5.2.4 Dateline Routing in k-ary n-cube Torus Networks 65
5.2.5 Matrix Formulation and Algorithmic Cycle Detection in CDGs 66
5.3 Glass & Ni Turn Models for Adaptive Routing 67
5.3.1 Turn Analysis in a 2D Mesh 67
5.3.2 The Canonical 2D Turn Models 67

5.3.3 3D Turn Models and Combinatorial Complexity 68
5.4 The Odd-Even Turn Model 69
5.5 Duato's Theorem for Fully Adaptive Routing 69
5.5.1 Formal Mathematical Statement 70
5.5.2 Proof of Duato's Theorem 70
5.5.3 Microarchitectural Realization: Adaptive and Escape VCs 71
5.6 Protocol-Level Deadlock and the PCMT Architecture 71
5.6.1 Coherence Dependency Chains and Classical Virtual Networks 72
5.6.2 PCMT: Prioritizing Coherence Message Types 73
5.6.3 Detailed Coherence Transaction Sequences and Race Deadlocks in MOESI 73
5.7 Topology-Agnostic Deadlock-Free Routing: Up*/Down* and Segment-Based 74
5.7.1 The Up*/Down* Routing Algorithm 74
5.7.2 The Root Bottleneck and Segment-Based Routing (SR) . 75
5.8 3D Multi-Chiplet Deadlock Freedom: ScaDA-Rou 76
5.9 Routing Algorithm Comparative Matrix 76
5.10 Interactive Lab & Hands-on Challenges 77
5.11 Exercises and Review Questions 77

Part III Physical Constraints: Power, Thermal & Reliability

6 Low-Power Architectures & Dark Silicon Dynamics 81
6.1 The Power Wall and the Dark Silicon Interconnect Bottleneck . . 81
6.1.1 Dennard Breakdown and the Dark Silicon Scaling Model 82
6.1.2 Physical Power Breakdown in Nanoscale Routers 83
6.2 DVFS and Voltage-Frequency Islands (VFIs) 83
6.2.1 Asynchronous Clock Domain Crossing (CDC) and Dual-Clock FIFOs 84
6.2.2 Feedback Control Theory for Autonomous VFI Frequency Scaling 85
6.2.3 Near-Threshold Computing (NTC) Interconnects & Razor Timing Speculation 86
6.3 Ultrafine-Grained Run-Time Power Gating 87
6.3.1 Sleep Transistor Sizing and In-Rush Damping 88
6.3.2 Breakeven Idle Time Formulation 89
6.3.3 Lookahead Early Wake-Up 89
6.4 Low-Power Link Encoding: The LPCR Framework 90
6.5 Bufferless Deflection Routing 90
6.5.1 Microarchitectural Mechanics and Flit Arbitration 90
6.5.2 The Energy Crossover Threshold 91
6.6 Quantitative Evaluation Matrix 91
6.7 Interactive Lab & Hands-on Challenges 91
6.8 Exercises and Review Questions 92

7 Thermal Modeling, Hotspots & Fault-Tolerant Resilience ... 93
7.1 The Thermal Emergency in Deep Sub-Micron Silicon 93
7.2 Compact Thermal Physics: The Electro-Thermal RC Duality .. 95
7.2.1 Duality Mapping Between Thermal and Electrical Domains 95
7.2.2 Multi-Layer Physical Heat Conduction Formulation 95
7.2.3 State-Space Governing Differential Equations 96
7.3 Thermal Sensor Placement & Reconstruction Under Non-Gaussian Noise 96
7.3.1 Gaussian-Sum Filter (GSF) Algorithmic Formulation ... 97
7.4 Intelligent Hotspot Prediction & Proactive Throttling 98
7.4.1 Proactive ANN Hardware Pipeline 99
7.4.2 Closed-Loop Proportional-Integral (PI) Thermal Throttling 100
7.5 Silicon Fault Tolerance: The ZoneDefense Framework 100
7.6 Hardware Microarchitecture: Fault-Tolerant Router Bypass Logic 101
7.7 Simulation Model & Experimental Verification 103
7.8 Microchannel Liquid Cooling and Conjugate Heat Transfer 104
7.8.1 Fluid Dynamics and Convective Heat Transfer in Microchannels 104
7.8.2 Pressure Drop and Pumping Power Overhead 104
7.9 Thermal-Aware Multi-Core Floorplanning Optimization 105
7.9.1 Mathematical Problem Formulation 105
7.9.2 Simulated Annealing Optimization Framework 105
7.10 Summary 106
7.11 Interactive Lab & Hands-on Challenges 106
7.12 Exercises and Review Questions 107

Part IV Advanced Packaging, 3D NoC & Multi-Chiplets

8 Three-Dimensional Networks-on-Chip (3D-NoC) & TSVs ... 111
8.1 The Three-Dimensional Scaling Paradigm 111
8.1.1 Architectural Advantages of 3D Interconnects 111
8.2 Through-Silicon Via (TSV) Physical Characteristics & Parasitics 112
8.2.1 TSV Fabrication Methodologies: Via-First, Via-Middle, and Hybrid Direct Bonding 113
8.2.2 High-Frequency Transmission and Parasitic Substrate Modeling 114
8.2.3 The Keep-Out Zone (KOZ) and Thermo-Mechanical Lamé Formulations 114
8.3 3D Router Microarchitectures & Vertical Arbitration 115
8.3.1 Crossbar Area Expansion in 7-Port 3D Routers 115
8.3.2 Dimensionally Partitioned Routers and 3D Crossbar Factorization 117
8.3.3 Vertical Arbitration Contention vs. Arbitration-Free Fast-Paths 117

8.4 Thermal Stacking & Heat Trapping in 3D Silicon 117
8.5 TSV Defect Tolerance & Redundancy 118
8.6 Hardware Microarchitecture: 7-Port Vertical Link Arbiter 118
8.7 Simulation Model & Numerical Verification 120
8.8 Pre-Bond and Post-Bond 3D TSV Testing and BIST 120
8.8.1 Pre-Bond TSV Test Architectures 121
8.8.2 Post-Bond IEEE 1500 Boundary-Scan Wrappers 121
8.9 Multi-Tier Power Distribution Network (PDN) Antiresonance . . 122
8.9.1 Inter-Tier Loop Inductance and Antiresonance Peaking . 122
8.9.2 Target Impedance and Decap Budget Allocation 122
8.10 Summary 123
8.11 Interactive Lab & Hands-on Challenges 123
8.12 Exercises and Review Questions 123

9 Multi-Chiplet Interconnects, 2.5D Systems & UCIe 125
9.1 The Chiplet Revolution: Breaking the Monolithic Barrier 125
9.2 2.5D Packaging Taxonomies: How Chiplets Connect 127
9.3 The Universal Chiplet Interconnect Express (UCIe) Standard . . 129
9.4 Inter-Chiplet Latency & Credit Buffer Sizing 130
9.4.1 Latency Decomposition: Where Does Time Go? 130
9.4.2 Little's Law: The Baseball Catching Problem 131
9.5 Deadlock in Multi-Chiplet Networks & In/Outbound Swap Recovery 132
9.5.1 Why Do Multi-Chiplet Networks Suffer Deadlocks? 132
9.5.2 Deadlock Avoidance vs. In/Outbound Swap Recovery . . . 133
9.6 Hardware Implementation: Synthesizable D2D Bridge 134
9.7 Industrial Multi-Chiplet Interconnect Case Studies 137
9.7.1 AMD Zen 4/5: CCD and IOD Infinity Fabric Architecture 137
9.7.2 Intel Sapphire Rapids and Granite Rapids: EMIB Multi-Tile Meshes 137
9.7.3 NVIDIA Blackwell B200 / GB200: NVLink-C2C and CoWoS-L 138
9.7.4 CEA-Leti IntAct 96-Core Active Interposer Platform 138
9.8 Simulation Model & Quantitative Evaluation 138
9.9 Summary: What You Learned 139
9.10 Interactive Lab & Hands-on Challenges 139
9.11 Exercises and Review Questions 140

Part V Emerging Frontiers: Optical, Wireless, Security & AI

10 Optical & Wireless Networks-on-Chip (ONoC / WNoC) 143
10.1 The Non-Traditional Interconnect Paradigm 143
10.2 Silicon Nanophotonics: Guiding Light Inside a Chip 144
10.2.1 Micro-Ring Resonators (MRRs): The Optical Filter 145
10.2.2 Switching Element Architectures: CSE vs. PSE 146
10.3 Optical Loss Modeling & Laser Power Sizing 146

10.3.1 Where Does the Light Go? . . . 148
10.3.2 Sizing the Laser: Why Decibels Matter . . . 148
10.4 Thermal Sensitivity: The Out-of-Tune Guitar . . . 149
10.5 Wireless Networks-on-Chip (WNoC): Radio Inside the CPU . . . 150
10.5.1 How Does an On-Chip Radio Work? . . . 150
10.5.2 Electromagnetic Wave Propagation in Lossy Silicon Substrates . . . 151
10.5.3 Energy-per-Bit and Transceiver Optimization . . . 152
10.5.4 Finding the Sweet Spot: Why 60 GHz? . . . 152
10.6 All-to-All Optical Crossbar Architectures: Corona and ATAC . . . 152
10.6.1 SWMR vs. MWSR Waveguide Topologies . . . 154
10.6.2 The Corona and ATAC Architectures . . . 154
10.7 Wireless Medium Access Control and Broadcast Acceleration . . . 154
10.7.1 Hardware MAC Protocols for On-Chip RF Channels . . . 155
10.7.2 Broadcast Acceleration in Cache Coherence Protocols (W-MESI) . . . 155
10.8 Hardware Implementation: Synthesizable Photonic MRR Switch 155
10.9 Simulation Model & Numerical Verification . . . 157
10.10 Summary: What You Learned . . . 158
10.11 Interactive Lab & Hands-on Challenges . . . 158
10.12 Exercises and Review Questions . . . 159

A Open-Source NoC Frameworks, RTL Repositories & PDKs . 161
A.1 Architectural Taxonomy of Open-Source NoC Frameworks . . . 161
A.1.1 Deep Architectural Profiles . . . 161
A.2 Synthesizable SystemVerilog Reference Router . . . 162
A.2.1 Lookahead Dimension-Order Routing Unit . . . 163
A.2.2 Single-Cycle Round-Robin Matrix Arbiter . . . 163
A.2.3 Parameterizable 5×5 Crossbar Datapath . . . 164
A.2.4 Top-Level 5-Port Virtual Channel Router . . . 165
A.3 Physical ASIC Implementation Flow with OpenROAD & OpenLane . . . 166
A.3.1 Target Process Design Kits . . . 166
A.3.2 Physical Synthesis Script for Yosys and OpenROAD . . . 167
A.3.3 Standard Cell Area and Power Breakdown . . . 167

Solutions to Checkpoint Questions & Exercises . . . 169
References . . . 176

Index . . . 181

List of Acronyms

2D/3D	Two-/Three-Dimensional
AMBA	Advanced Microcontroller Bus Architecture
AXI4	Advanced eXtensible Interface 4
BEP	Bit Error Probability
BGA	Ball Grid Array
CDG	Channel Dependency Graph
CMP	Chip Multiprocessor
CRC	Cyclic Redundancy Check
D2D	Die-to-Die
DMA	Direct Memory Access
DVFS	Dynamic Voltage and Frequency Scaling
Flit	Flow Control Unit
GALS	Globally Asynchronous Locally Synchronous
HOL	Head-of-Line (Blocking)
HPCA	High-Performance Computer Architecture
IP	Intellectual Property
IRDS	International Roadmap for Devices and Systems
ISCA	International Symposium on Computer Architecture
ITRS	International Technology Roadmap for Semiconductors
JSSC	IEEE Journal of Solid-State Circuits
LLM	Large Language Model
MPSoC	Multi-Processor System-on-Chip
NI	Network Interface
NoC	Network-on-Chip
OCSSI	On-Chip Source-Synchronous Interconnect
ONoC	Optical Network-on-Chip
PE	Processing Element
Phit	Physical Unit

PMA	Physical Medium Attachment
PPA	Power, Performance, and Area
PUF	Physically Unclonable Function
QoS	Quality of Service
RC	Route Computation
RoB	Reorder Buffer
RTT	Round-Trip Time
SA	Switch Allocation
SAF	Store-and-Forward (Switching)
SCC	Single-Chip Cloud Computer
SNN	Spiking Neural Network
SoC	System-on-Chip
ST	Switch Traversal
TCAD	IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems
TCAS	IEEE Transactions on Circuits and Systems
TSV	Through-Silicon Via
TVLSI	IEEE Transactions on Very Large Scale Integration Systems
UCIe	Universal Chiplet Interconnect Express
VA	Virtual Channel Allocation
VC	Virtual Channel
VCT	Virtual Cut-Through (Switching)
WH	Wormhole (Switching)
WNoC	Wireless Network-on-Chip
WSC	Wafer-Scale Computing

Part I
Foundations & Router Microarchitecture

Chapter 1
The Multi-Core Communication Crisis & Interconnect Physics

Abstract For over three decades, semiconductor compute performance accelerated through Dennard scaling and monolithic bus interconnects. Around 2005, the physical collapse of Dennard scaling halted clock frequency scaling, forcing a historic pivot toward horizontal core parallelism. However, scaling processing core counts immediately exposed the interconnect crisis: shared multi-drop buses suffer from cubic Elmore delay scaling and severe serialization bottlenecks; centralized crossbars explode with quadratic silicon area and routing congestion; and deep-submicron copper lines undergo severe electron scattering under the Mayadas-Shatzkes effect. This chapter establishes the physical foundations of on-chip wire scaling, mathematically analyzes bus and crossbar limitations, and introduces the Network-on-Chip (NoC) paradigm as the scalable packet-switched backbone of modern manycore processors, highlighted through silicon case studies including the 2,048-core HammerBlade SoC.

1.1 The Breakdown of Dennard Scaling and the Multi-Core Pivot

For over three decades, the semiconductor industry advanced along the dual trajectories of **Moore's Law** (transistor density doubling approximately every 18–24 months) and **Dennard Scaling** [12] (constant power density through proportional scaling of supply voltage V_{dd}, threshold voltage V_{th}, and physical oxide thickness t_{ox}). Under classical Dennard scaling:

$$\text{Power Density} = \frac{P_{active}}{\text{Area}} \propto \frac{C \cdot V_{dd}^2 \cdot f}{1/s^2} \approx \text{Constant} \tag{1.1}$$

where $s \approx 1.4$ is the per-generation linear scaling factor. Because operating frequency could be increased by $\approx 40\%$ each generation without increasing power dissipation per unit silicon area, single-threaded performance grew exponentially without encountering thermal catastrophic failure.

Around 2004–2005, at the 90 nm and 65 nm process nodes, Dennard scaling suffered an irreversible collapse. As gate oxide thickness reached atomic dimensions ($t_{ox} \leq 1.2$ nm, merely 4–5 silicon dioxide monolayers), **quantum mechanical direct tunneling** through the gate dielectric surged exponentially:

$$J_{tunnel} \approx A \cdot \left(\frac{V_{ox}}{t_{ox}}\right)^2 \exp\left(-\frac{B \cdot [1-(1-V_{ox}/\Phi_B)^{3/2}]}{V_{ox}/t_{ox}}\right) \tag{1.2}$$

At the same time, sub-threshold leakage current I_{sub} prevented further aggressive reductions of the supply voltage V_{dd}, as the threshold voltage V_{th} could not be reduced without causing unacceptable static standby power dissipation:

$$I_{sub} = \mu_0 C_{ox} \frac{W}{L}(m-1)V_T^2 \exp\left(\frac{V_{gs}-V_{th}}{mV_T}\right)\left(1-\exp\left(-\frac{V_{ds}}{V_T}\right)\right) \tag{1.3}$$

where $V_T = \frac{k_B T}{q} \approx 26$ mV at 300 K, and m is the sub-threshold swing factor ($m \approx 1.1$–1.4). Consequently, chip architects struck the unyielding **Power Wall** ($\approx$ 100–130 W/cm^2 for air-cooled packaging). Increasing the clock frequency beyond 3.5–4.0 GHz became thermally untenable.

To maintain compute throughput growth, processor design executed a historic pivot from *frequency scaling* to *spatial concurrency*: multi-core processors, Chip Multiprocessors (CMPs), and Multi-Processor Systems-on-Chip (MPSoCs). However, scaling the number of processing elements from dual-core to tens, hundreds, and thousands of cores immediately triggered an existential crisis in **on-chip interconnect physical design**.

1.2 The Failure of Shared Multi-Drop Buses

In early Systems-on-Chip, intellectual property (IP) blocks (CPUs, DSPs, DMA engines, SRAM banks) communicated through shared multi-drop buses such as ARM AMBA (AHB/APB), IBM CoreConnect, or Wishbone. While shared buses are simple and cost-effective for small SoCs ($N \leq 4$–8 masters), they collapse under severe physical and architectural bottlenecks as core counts grow.

1.2.1 Capacitive Loading & Elmore Wire Delay

A shared bus consists of long physical metal traces spanning across the silicon die. Every connected master and slave attaches an input/output transceiver tap to the shared wire. According to the **Elmore Delay Model** [15], the propagation delay τ_{bus} of a distributed RC line loaded by N capacitive taps is:

Table 1.1 Physical Interconnect Scaling Across Core Counts (Die Edge: 12 mm, Process: 12nm FinFET, $r_{wire} = 800\ \Omega$/mm, $c_{wire} = 0.18$ pF/mm, $c_{tap} = 8$ fF)

Cores (N)	Bus Delay (ps)	Max Bus Freq	Aggregate Bus BW	Crossbar Switches	2D Mesh Bisection BW
2	9,044 ps	55.3 MHz	0.88 GB/s	4	128 GB/s
4	9,129 ps	54.8 MHz	0.88 GB/s	16	128 GB/s
8	9,298 ps	53.8 MHz	0.86 GB/s	64	192 GB/s
16	9,638 ps	51.9 MHz	0.83 GB/s	256	256 GB/s
32	10,316 ps	48.5 MHz	0.78 GB/s	1,024	384 GB/s
64	11,673 ps	42.8 MHz	0.69 GB/s	4,096	512 GB/s
128	14,386 ps	34.8 MHz	0.56 GB/s	16,384	768 GB/s
256	19,814 ps	25.2 MHz	0.40 GB/s	65,536	1,024 GB/s

$$\tau_{bus} \approx 0.38 \cdot R_{wire} \cdot C_{wire} + R_{driver} \cdot \left(C_{wire} + \sum_{i=1}^{N} C_{tap,i} \right) + \sum_{i=1}^{N} R_{wire,i} \cdot C_{tap,i} \tag{1.4}$$

where:

- $R_{wire} = r_{wire} \cdot L$ is the total resistance of the global metal wire across the die ($L \approx 10$–20 mm).
- $C_{wire} = c_{wire} \cdot L$ is the distributed wire capacitance (area capacitance to substrate plus fringe capacitance).
- $C_{tap,i}$ is the input/output capacitance of the i-th master/slave transceiver tap.
- R_{driver} is the equivalent output resistance of the driving transistor.

Because every attached core adds capacitive load ($C_{taps} = N \cdot C_{tap}$), the total capacitance scales linearly with core count ($C_{total} \propto N$). Furthermore, because the wire length L scales with the die perimeter ($L \propto \sqrt{N}$), the intrinsic wire RC product grows with the square of the die dimension:

$$\tau_{bus} \propto R_{wire} \cdot C_{total} \propto (\sqrt{N}) \cdot (N + \sqrt{N}) \propto N^{3/2} \tag{1.5}$$

Consequently, the maximum achievable bus clock frequency ($f_{max} \approx \frac{1}{2\tau_{bus}}$) plunges precipitously. While local processor logic easily ran at 2.0–3.0 GHz, global shared buses were throttled to 200–400 MHz.

1.2.2 The Serialization Contention Bottleneck

A shared bus is a **single shared transmission medium**. At any given clock cycle, the centralized bus arbiter can grant bus access to **exactly one master**:

- **Fixed Aggregate Bandwidth**: $BW_{agg} = W \times f_{bus}$, where W is bus width (e.g., 64 or 128 bits) and f_{bus} is bus clock frequency.

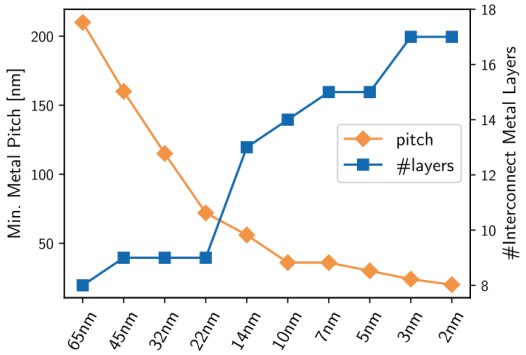


Fig. 1. Technology scaling of on-chip wire resources based on IDRS reports [8] for 2–14 nm and [9] for 22–65 nm.

Fig. 1.1 Technology scaling of on-chip wire resources based on ITRS/IRDS reports for 2–14 nm and 22–65 nm nodes. While transistor delay decreases exponentially, global wire delay surges. [Reproduced from: T. Fischer, M. Rogenmoser, T. Benz, F. K. Gürkaynak, and L. Benini, "FlooNoC: A 645-Gb/s/link 0.15-pJ/B/hop Open-Source NoC With Wide Physical Links and End-to-End AXI4 Parallel Multistream Support," *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 33, no. 4, pp. 1094–1107, Apr. 2025 [17]. © 2025 IEEE]

- **Per-Core Bandwidth Starvation**: $BW_{per_core} \leq \frac{W \times f_{bus}}{N}$. As core count N increases, per-core bandwidth asymptotically approaches zero.
- **Arbitration Scalability**: Centralized round-robin or matrix arbiters require $O(N)$ logic depth, transforming grant generation into a critical-path timing bottleneck.

1.3 Sub-10nm Interconnect Physics & Wire Scaling

The physical crisis of global wiring is exacerbated by the physics of deep-submicron copper interconnects [23]. In technology nodes below 14 nm, 7 nm, and 3 nm, interconnect resistance does not scale inversely with cross-sectional area.

1.3.1 The Mayadas-Shatzkes & Fuchs-Sondheimer Scattering Models

In bulk copper at room temperature, the electron mean free path is $\lambda_0 \approx 39.0$ nm. When the physical width w of an on-chip metal wire shrinks below $2\lambda_0 \approx 80$ nm, two scattering mechanisms cause the effective resistivity ρ_{eff} to surge dramatically above the bulk value ($\rho_{bulk} \approx 1.68 \times 10^{-8}\ \Omega \cdot \text{m}$):

1. **Surface Scattering (Fuchs-Sondheimer Model)**: Electrons reflect diffusely off the outer interfaces of the wire:

$$\frac{\rho_{FS}}{\rho_{bulk}} = 1 + \frac{3}{8\kappa}(1-p) \tag{1.6}$$

where $\kappa = \frac{w_{eff}}{\lambda_0}$ is the normalized wire dimension and $p \in [0,1]$ is the specularity parameter ($p \approx 0.2$ for realistic barrier interfaces).

2. **Grain Boundary Scattering (Mayadas-Shatzkes Model)** [38]: Nanoscale copper recrystallizes into small grains with average grain diameter $d \approx w_{eff}$. Electrons scatter at grain boundaries:

$$\frac{\rho_{MS}}{\rho_{bulk}} = \frac{1}{3\left[\frac{1}{3} - \frac{\alpha}{2} + \alpha^2 - \alpha^3 \ln\left(1 + \frac{1}{\alpha}\right)\right]} \tag{1.7}$$

where $\alpha = \frac{\lambda_0}{d}\frac{R}{1-R}$, with $R \approx 0.3$–0.5 representing the grain boundary reflection coefficient.

3. **Diffusion Barrier Liner Encroachment**: Copper electromigration requires high-resistivity diffusion liners (Ta/TaN, Co, Ru) of thickness $t_{barrier} \approx$ 1.0–1.5 nm. Because the liner does not scale down proportionally with the wire trench, the conductive copper cross-sectional area A_{Cu} shrinks much faster than total wire trench area:

$$A_{Cu} = (w - 2t_{barrier}) \cdot (h - t_{barrier}), \qquad \rho_{total} = \rho_{eff} \cdot \left(\frac{w \cdot h}{A_{Cu}}\right) \tag{1.8}$$

At the 3 nm node ($w \approx 10$ nm), the effective copper resistivity exceeds 15× bulk copper, turning long global wires into unbearable thermal, resistance, and latency bottlenecks.

1.3.2 Sub-3nm Transistor Architectures: GAAFET, RibbonFET, and CFET

As silicon dimensions scale into the sub-3nm regime, conventional FinFETs encounter severe short-channel degradation and sub-threshold leakage due to quantum confinement limits and corner parasitic leakage. Semiconductor foundries have transitioned to **Gate-All-Around Field-Effect Transistors (GAAFETs)** (termed *RibbonFET* by Intel and *MBCFET* by Samsung):

1. **Nanosheet Architecture**: Rather than a vertical 3D silicon fin gated on three sides, a GAAFET comprises multiple horizontally stacked silicon nanosheet ribbons (typically 3 to 4 sheets of thickness $t_{sheet} \approx$ 5–6 nm and width $w_{sheet} \approx$ 15–40 nm) completely wrapped by the high-κ metal gate dielectric (HfO_2/TiN).

2. **Sub-Threshold Swing and Electrostatic Control**: Enclosing the channel on all four sides maximizes gate capacitive coupling, driving the sub-threshold swing S_S down to ≈ 64–$67\,\text{mV/decade}$ (near the theoretical thermodynamic limit of $60\,\text{mV/decade}$ at $300\,\text{K}$) and suppressing Drain-Induced Barrier Lowering (DIBL $\leq 30\,\text{mV/V}$).
3. **Complementary FET (CFET)**: The ultimate scaling evolution of GAAFETs is the **CFET**, in which an NMOS nanosheet stack is monolithically fabricated *vertically on top* of a PMOS nanosheet stack. By stacking opposite-polarity transistors vertically, standard cell track height shrinks from 5.0T down to 3.0T or 4.0T, doubling transistor density without requiring physical gate pitch reduction.

1.3.3 Backside Power Delivery Networks (BSPDN / PowerVia)

In classical Front-End-of-Line / Back-End-of-Line (FEOL/BEOL) manufacturing, both low-voltage logic signal wires and high-current power distribution rails (V_{dd}, V_{ss}) compete for the same frontside metal tracks (M0 through M15). Power routing consumes up to 30% of intermediate metal tracks, exacerbating routing congestion and inducing severe resistive voltage drop (IR-drop).

To overcome this interconnect bottleneck, modern sub-2nm nodes introduce the **Backside Power Delivery Network (BSPDN)** (such as Intel PowerVia and TSMC A16):

- The wafer is thinned from the backside down to a thickness of $\leq 500\,\text{nm}$.
- Power delivery networks (V_{dd} and ground grids) are fabricated entirely on the wafer **backside** using thick, ultra-low-resistance copper layers.
- Sub-micron **Nano-Through-Silicon Vias (nTSVs)** connect backside power rails directly to transistor source/drain terminals.
- Frontside metal layers (M0–M14) are freed exclusively for high-density signal interconnects, slashing IR-drop by over 30% and increasing signal routing density by up to 20%.

1.3.4 Alternative Interconnect Metallization: Ruthenium, Cobalt, and Molybdenum

Because copper requires high-resistivity diffusion barrier liners (Ta/TaN) that do not scale below 1.0–1.5 nm, the effective copper conductor area in sub-12nm trenches collapses, triggering catastrophic resistivity surges (Eq. 1.8).

To solve this physical scaling wall, semiconductor manufacturers are transitioning the lowest metal levels (M0, M1, M2) from copper to **linerless refractory and platinum-group metals**:

- **Ruthenium (Ru)**: Ruthenium is a semi-noble metal with high cohesive energy that resists electromigration without requiring any diffusion barrier liner. Crucially, the bulk electron mean free path of ruthenium is only $\lambda_0(Ru) \approx 4.9\,\text{nm}$—nearly $8\times$ **shorter than copper** ($\lambda_0(Cu) \approx 39.0\,\text{nm}$). Consequently, size-effect grain boundary scattering in ruthenium is drastically attenuated at sub-10nm line widths.
- **Molybdenum (Mo) and Cobalt (Co)**: Molybdenum exhibits an electron mean free path of $\lambda_0(Mo) \approx 10.0\,\text{nm}$ and low bulk resistivity ($5.3\,\mu\Omega \cdot \text{cm}$), providing superior aspect-ratio fill characteristics without voiding.

The Metal Crossover Threshold: At wire trench widths $w > 15\,\text{nm}$, copper remains superior due to its low bulk resistivity ($1.68\,\mu\Omega \cdot \text{cm}$ vs. $7.1\,\mu\Omega \cdot \text{cm}$ for Ru). However, when wire width shrinks below the **crossover width** ($w \leq 12\,\text{nm}$), the linerless nature and short mean free path of Ruthenium cause its total line resistance to become **lower than linered copper**, establishing Ruthenium as the future metallization standard for sub-2nm NoC link routers.

1.3.5 Historical Technology Scaling Across Semiconductor Generations

Table 1.2 chronicles the physical parameters of interconnect and transistor scaling from the 180nm planar CMOS node to the 2nm GAAFET generation.

Table 1.2 Comprehensive Interconnect and Transistor Metrics Across Semiconductor Technology Generations

Node	Year	Device Type	V_{dd} (V)	Min Pitch	Gate Delay	Wire Delay (1 mm)
180 nm	1999	Planar Bulk	1.80	500 nm	35.0 ps	2.5 ps
90 nm	2003	Strained Si	1.20	220 nm	15.0 ps	8.0 ps
45 nm	2007	High-κ Metal	1.00	160 nm	7.0 ps	25.0 ps
22 nm	2012	3D FinFET	0.85	80 nm	3.5 ps	90.0 ps
7 nm	2018	EUV FinFET	0.75	40 nm	1.8 ps	350.0 ps
3 nm	2023	GAAFET	0.70	24 nm	1.1 ps	1,200.0 ps
2 nm	2025	RibbonFET/BSPDN	0.65	18 nm	0.8 ps	2,800.0 ps

As documented in Table 1.2, while gate delay dropped by over $43\times$ from 180nm to 2nm, unrepeated wire delay across a 1 mm global span escalated by more than $1{,}100\times$. This catastrophic divergence provides the definitive physical justification for packet-switched Networks-on-Chip.

1.4 High-Frequency Transmission Line Effects and Inductive Peaking

As clock frequencies exceed 2–4 GHz and signal rise times t_r compress below 30–50 ps, on-chip global wires cease to behave as purely diffusive distributed RC circuits and transition into distributed RLC transmission lines. When the signal transition time t_r is comparable to the electromagnetic time of flight $t_f = L\sqrt{\mathcal{L} \cdot \mathcal{C}}$, inductive reactance $\omega\mathcal{L}$ rivals series resistance $\mathcal{R}$, inducing reflection ringing, inductive crosstalk, and signal overshoot.

1.4.1 The Telegrapher's Equations and the Wave Equation

A transmission line is governed by the classical Telegrapher's equations:

$$-\frac{\partial v(x,t)}{\partial x} = \mathcal{R} \cdot i(x,t) + \mathcal{L} \cdot \frac{\partial i(x,t)}{\partial t} \tag{1.9}$$

$$-\frac{\partial i(x,t)}{\partial x} = \mathcal{G} \cdot v(x,t) + \mathcal{C} \cdot \frac{\partial v(x,t)}{\partial t} \tag{1.10}$$

where $\mathcal{R}$ (Ω/m), $\mathcal{L}$ (H/m), $\mathcal{G}$ (S/m), and $\mathcal{C}$ (F/m) denote per-unit-length series resistance, loop inductance, dielectric conductance, and shunt capacitance, respectively. Differentiating with respect to spatial coordinate x yields the lossy wave equation:

$$\frac{\partial^2 v(x,t)}{\partial x^2} = \mathcal{LC}\frac{\partial^2 v(x,t)}{\partial t^2} + (\mathcal{RC} + \mathcal{LG})\frac{\partial v(x,t)}{\partial t} + \mathcal{RG}v(x,t) \tag{1.11}$$

The transition criterion from the diffusive RC regime to the wave-propagation RLC regime is established by comparing the wire length L to the attenuation attenuation length and signal rise time:

$$\text{RLC Effects Significant} \iff t_r \leq 2.5 \cdot t_f = 2.5 \cdot L\sqrt{\mathcal{LC}} \quad \text{and} \quad \mathcal{R} \cdot L \leq 2Z_0 \tag{1.12}$$

where $Z_0 = \sqrt{\mathcal{L}/\mathcal{C}}$ is the characteristic impedance of the lossless transmission line. In intermediate metal layers (M_4–M_7), high series resistance ($\mathcal{R} \gg \omega\mathcal{L}$) damps inductive waves, keeping lines in the RC regime. However, in top-level global redistribution layers (thick metal M_{top}, $h \geq 1.5\,\mu$m, $w \geq 1.0\,\mu$m), series resistance drops to $\mathcal{R} \leq 20\,\Omega$/mm while loop inductance reaches $\mathcal{L} \approx 0.5$–1.0 nH/mm. Under these conditions, wire propagation speed approaches the speed of light in silicon dioxide ($v = c/\sqrt{\epsilon_r} \approx 150\,\mu$m/ps), but un-terminated lines suffer severe reflections with reflection coefficient $\Gamma_L = \frac{R_L - Z_0}{R_L + Z_0}$.

1.4.2 High-Frequency Skin Effect and Dielectric Loss

At multi-gigahertz frequencies, alternating magnetic fields force current toward the outer conductor perimeter, confining conduction to a characteristic skin depth δ:

$$\delta = \sqrt{\frac{\rho}{\pi \cdot f \cdot \mu_0 \cdot \mu_r}} \tag{1.13}$$

For copper ($\rho = 1.68 \times 10^{-8}\,\Omega \cdot \mathrm{m}$, $\mu_r = 1$), the skin depth is $\delta \approx 0.95\,\mu\mathrm{m}$ at 5 GHz and shrinks to $\delta \approx 0.67\,\mu\mathrm{m}$ at 10 GHz. When conductor thickness exceeds 2δ, current crowding raises the effective AC resistance:

$$\mathcal{R}_{AC}(f) \approx \mathcal{R}_{DC} \cdot \left[1 + \left(\frac{f}{f_{skin}}\right)^2\right]^{1/4} \propto \sqrt{f} \tag{1.14}$$

Concurrently, dielectric polarization lag in the low-k inter-layer dielectric introduces a loss tangent $\tan\delta_{diel}$:

$$\mathcal{G}(f) = 2\pi f \cdot \mathcal{C} \cdot \tan\delta_{diel} \tag{1.15}$$

Together, skin-effect attenuation and dielectric absorption severely attenuate the high-frequency harmonics of digital clock and data edges, causing eye-diagram closure on long unbuffered buses.

1.5 The Monolithic Crossbar and Its Quadratic Area Explosion

To avoid the serialized contention of a single shared bus, architects initially considered a full centralized crossbar switch connecting N masters directly to M slaves. While a crossbar allows non-interfering parallel communication between disjoint pairs of masters and slaves, its hardware complexity scales quadratically:

$$\text{Crosspoints} = N \times M \approx O(N^2), \qquad \text{Wiring Tracks} = (N+M) \times W_{data_width} \approx O(N) \tag{1.16}$$

For an MPSoC with $N = 64$ cores, a full crossbar requires $64 \times 64 = 4,096$ crosspoint multiplexer switches and $128 \times 128 = 16,384$ wires crossing a central silicon layout area, causing severe layout congestion, routing track starvation, and parasitic capacitive coupling crosstalk. A full crossbar switch is physically viable only for small clusters ($N \leq 4$–8 cores).

1.6 Optimal Repeater Insertion Theory (Bakoglu's Formulation)

To counteract the quadratic delay growth of long RC lines ($\tau \propto L^2$), designers segment wires by inserting CMOS inverters or non-inverting buffers (repeaters) at regular intervals. By partitioning a line of length L into k equal segments of length $l = L/k$, each driven by a buffer sized h times the minimum-sized inverter, the delay transforms from quadratic to linear in total length:

$$\tau_{total} = k \cdot \tau_{segment} = \frac{L}{l} \cdot \left[R_0 \cdot \left(\frac{\mathcal{C} \cdot l}{h} + C_0 \right) + \frac{\mathcal{R} \cdot l}{h} \cdot C_0 + \frac{1}{2} \mathcal{R}\mathcal{C} \cdot l^2 \right] \quad (1.17)$$

where R_0 and C_0 are the output drive resistance and input gate capacitance of a minimum-sized inverter ($h = 1$), and $\mathcal{R}, \mathcal{C}$ are the per-unit-length wire resistance and capacitance.

By setting the partial derivatives $\frac{\partial \tau_{total}}{\partial l} = 0$ and $\frac{\partial \tau_{total}}{\partial h} = 0$, Bakoglu derived the closed-form optimal repeater spacing l_{opt} and buffer sizing factor h_{opt}:

$$l_{opt} = \sqrt{\frac{2R_0C_0}{\mathcal{R} \cdot \mathcal{C}}} \quad (1.18)$$

$$h_{opt} = \sqrt{\frac{R_0 \cdot \mathcal{C}}{\mathcal{R} \cdot C_0}} \quad (1.19)$$

Substituting Eq. (1.18) and Eq. (1.19) into Eq. (1.17) yields the minimum achievable interconnect propagation delay per unit length:

$$\left(\frac{\tau}{L}\right)_{min} = 2\sqrt{R_0C_0 \cdot \mathcal{R}\mathcal{C}} \cdot \left(1 + \sqrt{\frac{1}{2}}\right) \approx 2.41 \cdot \sqrt{R_0C_0 \cdot \mathcal{R}\mathcal{C}} \quad (1.20)$$

The Repeater Insertion Dilemma: While optimal repeater insertion linearizes wire delay, it imposes severe silicon area and power penalties:

- **Active Silicon Tax**: For a 1024-bit bus traversing a 20 mm die, optimal repeaters consume over 12% of active silicon area in intermediate tiers.
- **Dynamic Switching Power**: $P_{dyn} = k \cdot (h \cdot C_0 + \mathcal{C} \cdot l_{opt}) \cdot V_{dd}^2 \cdot f$, dissipating gigawatts of unrecoverable charging energy across global links.
- **Bidirectional Routing Obstacles**: Repeaters are unidirectional CMOS amplifiers; bidirectional buses require complex tri-state transceivers and direction-control routing.

1.7 The Network-on-Chip Paradigm: Packet-Switched Modular Interconnects

In 2001, pioneering papers by Dally & Towles (*"Route Packets, Not Wires"* [8]) and Benini & De Micheli (*"Networks on Chips: A New SoC Paradigm"* [2]) introduced the definitive architectural solution:

The Network-on-Chip (NoC) Principle: Replace long, dedicated, unstructured global metal wires with a structured, packet-switched micronetwork composed of modular routers and short, pipelined point-to-point links.

Key advantages of the NoC architecture include:

1. **Bounded Wire Lengths**: Every link connects strictly adjacent routers. The link wire length is fixed to the physical pitch of one compute tile ($L_{link} \approx 0.5$–1.5 mm), regardless of whether the chip contains 16 or 2,048 cores.
2. **Predictable Timing & Pipelining**: Because link wire delay is short and localized, routers and links can be clocked at high frequencies ($f_{clk} \approx 1.5$–2.5 GHz) without repeater insertion crises.
3. **Concurrent Multi-Hop Communication**: Multiple packets traverse non-overlapping segments of the network simultaneously, yielding aggregate throughput that scales proportionally with the network size:

$$BW_{bisection} \propto \sqrt{N} \times W_{link} \times f_{clk} \tag{1.21}$$

4. **Modularity & Silicon Reuse**: Adding processing tiles does not require redesigning the global wiring tree.

1.8 Industrial Case Studies: From 48 Cores to Wafer-Scale Engines

To examine the real-world execution of NoC architecture in modern silicon, we analyze landmark implementations across academic and industrial milestones:

1.8.1 The Intel TeraFLOPS 80-Core Processor (65nm CMOS)

Fabricated by Intel in 2007 (Vangal et al., ISSCC), the TeraFLOPS processor represented the world's first microprocessor to achieve 1.0 TFLOPS double-precision floating-point performance on a single chip. It featured an 8×10 2D mesh of compute tiles running at 4.0 GHz at 1.2 V (62 W). Each tile contained

a 5-port 5-stage pipelined packet-switched router with crossbar switch bandwidth of 80 GB/s per node, providing an unprecedented 1.6 Tb/s of bisection bandwidth across short, regular 1.2 mm point-to-point links.

1.8.2 The Tilera TILE64 Manycore Processor (90nm CMOS)

The Tilera TILE64 architecture departed from traditional single-mesh NoCs by implementing **five independent, concurrent physical 2D mesh networks** on the same silicon die:

1. **UDI (User Dynamic Interconnect)**: Zero-overhead streaming user-level message passing between processor cores.
2. **IDN (I/O Dynamic Interconnect)**: Dedicated DMA transfers from PCIe and 10GbE network interfaces to on-chip tile SRAM caches.
3. **MDN (Memory Dynamic Interconnect)**: Cache-to-DDR memory controller request and response channels.
4. **TDN (Tile Dynamic Interconnect)**: Coherent L2 cache-to-cache intervention and invalidation traffic.
5. **STN (Static Interconnect)**: Circuit-switched deterministic real-time communication.

By separating traffic into independent physical networks, the TILE64 completely prevented cache coherence transactions from being delayed by bulk I/O streaming, delivering over 30 Tb/s of aggregate bisection throughput.

1.8.3 The Cerebras Wafer-Scale Engine (WSE-2 & WSE-3)

The Cerebras Wafer-Scale Engine represents the extreme physical limit of monolithic Network-on-Chip scaling. Rather than cutting wafers into individual dies, Cerebras fabricates an entire 300 mm silicon wafer as a single contiguous compute system:

- **Core Count**: $850,000$ AI-optimized tensor processing cores (WSE-2) and $900,000$ cores in 5nm (WSE-3).
- **Wafer-Scale Interconnect**: The chip-cutting boundary (scribe lines) is bridged using dedicated high-density lithographic cross-reticle metal lines, creating a seamless 2D mesh fabric spanning $46,225\,\text{mm}^2$.
- **On-Chip Fabric Bandwidth**: The contiguous 2D mesh fabric delivers over **220** Pb/s (petabits per second) of internal fabric bandwidth with single-cycle router latency per hop.
- **Hardware Redundancy**: A built-in fault-tolerant bypass fabric automatically detects and routes around defective fabrication sites, ensuring 100% operational yield on a wafer with thousands of native silicon defect points.

Table 1.3 Evolution of Silicon Landmark Network-on-Chip Processors

Architecture	Year / Venue	Process Node	Tile Array / Cores	Peak Bisection BW	Interconnect Topology
Intel TeraFLOPS	2007 JSSC [12]	65nm Bulk	8×10 (80)	1.6 Tb/s	2D Mesh (5-port)
Intel 48-Core SCC	2011 JSSC [24]	45nm Bulk	6×4 (48)	256 Gb/s	2D Mesh (e-cube)
Theseus WSC	2025 TCAD [56]	Wafer-Scale	Wafer Grid	100+ Tb/s	2D Torus Wafer NoC
FlooNoC (Benini)	2025 TVLSI [17]	12nm Fin-FET	8×4 (288)	103 Tb/s	Wide AXI4 2D Mesh
HammerBlade (Gao)	2026 TVLSI [21]	12nm Fin-FET	64×32 (2,048)	53.2 Tb/s	Ruche Multi-Hop NoC

To examine the real-world execution of NoC architecture in modern silicon, we analyze landmark implementations from our literature corpus:

1.8.4 The 2,048-Core HammerBlade Manycore SoC (12nm FinFET)

A landmark demonstration of extreme monolithic scaling is the **HammerBlade (HB)** processor developed by Gao et al. [21] and Davidson et al. [10].

Key architectural and physical metrics of the HammerBlade processor include:

- **Silicon Area**: 99 mm^2 in TSMC 12nm FinFET.
- **Compute Density**: 2,048 custom RISC-V cores organized into an array of 64×32 tiles.
- **Interconnect Topology**: **Ruche Network**, a wire-maximal topology overlaying nearest-neighbor 2D mesh links with long-distance multi-hop bypass express wires.
- **Clocking & Throughput**: Operating at **1.49 GHz at 0.8 V**, the NoC delivers a peak aggregate bandwidth of **2,572.6 Tb/s** and a bisection bandwidth of **53.2 Tb/s**.
- **Routing Density**: Achieves **4,289 bit/mm** across higher metal layers.
- **Integration**: Leverages **On-Chip Source-Synchronous Interconnect (OCSSI)** under a Globally Asynchronous Locally Synchronous (GALS) paradigm to eliminate chip-wide clock tree skew.

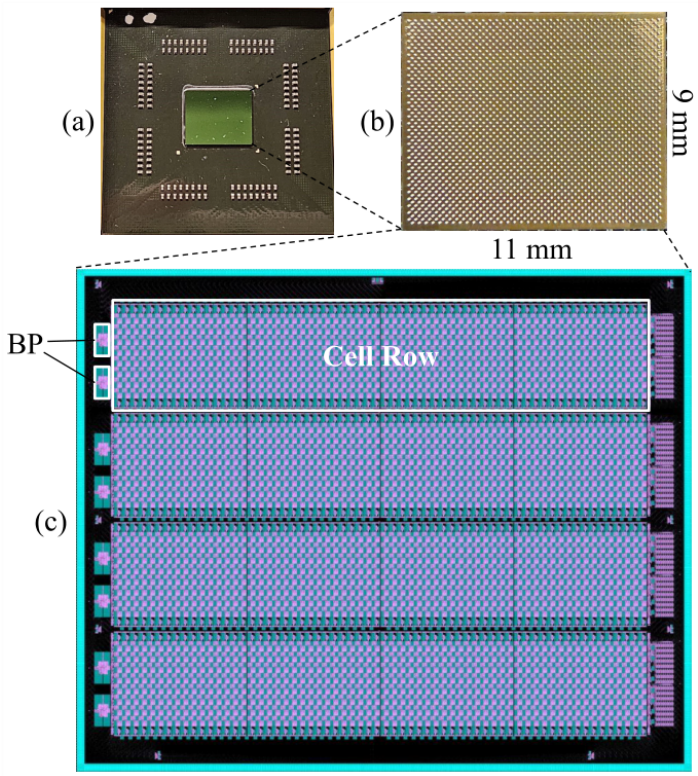


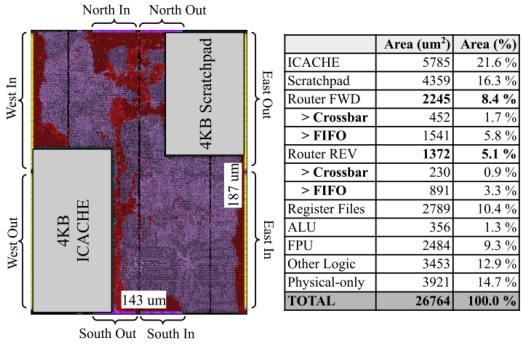
Fig. 1. (a) BGA 2116 flip-chip packaging, (b) 11 × 9-mm die photograph, and (c) chip floorplan. BP = BlackParrot Processor [7].

Fig. 1.2 Modern monolithic manycore scaling in 12nm silicon: (a) BGA 2116 flip-chip packaging, (b) 11 × 9-mm die photograph, and (c) chip floorplan of the 2,048-core HammerBlade SoC. [Reproduced from: P. Gao et al., "HammerBlade in Silicon: A 12-nm 2048-Core RISC-V Manycore SoC With Ruche Networks," *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 34, no. 7, pp. 2306–2310, July 2026 [21]. © 2026 IEEE]

1.9 Interactive Lab & Hands-on Challenges

To build an experimental understanding of interconnect trade-offs, explore the companion simulators on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **2D Mesh Routing Visualizer** (https://network-on-chip.github.io/learning/interactive/visualizer/): Inject uniform random and hotspot traffic into a 4×4 and 8×8 mesh; observe how localized point-to-point links sustain continuous throughput without centralized bus stalls.



	Area (um^2)	Area (%)
ICACHE	5785	21.6 %
Scratchpad	4359	16.3 %
Router FWD	**2245**	**8.4 %**
> **Crossbar**	452	1.7 %
> **FIFO**	1541	5.8 %
Router REV	**1372**	**5.1 %**
> **Crossbar**	230	0.9 %
> **FIFO**	891	3.3 %
Register Files	2789	10.4 %
ALU	356	1.3 %
FPU	2484	9.3 %
Other Logic	3453	12.9 %
Physical-only	3921	14.7 %
TOTAL	**26764**	**100.0 %**

Fig. 2. HB compute tile floorplan and its area breakdown. Router cells (red) and the core logic (pink).

Fig. 1.3 HammerBlade compute tile floorplan and physical area breakdown showing router cells (red) and core logic (pink). [Reproduced from: P. Gao et al., "HammerBlade in Silicon: A 12-nm 2048-Core RISC-V Manycore SoC With Ruche Networks," *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 34, no. 7, pp. 2306–2310, July 2026 [21]. © 2026 IEEE]

- **Topology Titans: Graph Architect** (https://network-on-chip.github.io/learning/interactive/topology/): Synthesize 2D Mesh, 2D Torus, and Hypercube topologies within strict wire length and pin degree limits.
- **Arbiter Arena: Crossbar Clash** (https://network-on-chip.github.io/learning/interactive/arbiter/): Benchmark matrix, round-robin, and wavefront arbiters under heavy contention.

1.10 Checkpoint Assessment

1. **Question 1.1 (Tier 1: Fundamental)**: Why does the propagation delay of a shared multi-drop bus grow as $O(N^{3/2})$ to $O(N^2)$ as core count N scales, rather than remaining constant?
2. **Question 1.2 (Tier 2: Architectural Analysis)**: What is the Mayadas-Shatzkes effect, and why does it represent a critical barrier for on-chip global interconnects below 14 nm?
3. **Question 1.3 (Tier 3: Design Trade-Off)**: Why is a centralized $N \times N$ crossbar switch unacceptable for a 128-core processor, despite offering non-blocking zero-hop communication between disjoint pairs?

Detailed answers and derivations for these questions are provided in the Solutions chapter at the end of this volume.

Chapter 2
Switching Techniques, Flow Control & Buffer Management

Abstract The efficiency, latency, and area footprint of an on-chip network are fundamentally dictated by its switching and flow control architecture. This chapter explores how raw application messages are decomposed into packets, flits, and phits, and details how silicon Network Interfaces bridge industry-standard protocols such as ARM AMBA AXI4 to internal NoC channels. We develop rigorous mathematical latency formulations comparing Store-and-Forward (SAF), Virtual Cut-Through (VCT), and Wormhole switching. We analyze lossless flow control mechanisms, proving the fundamental relationship between credit turnaround loop round-trip time (RTT) and minimum buffer sizing to sustain wire-rate throughput. Finally, we examine low-power alternative flow-control architectures, featuring the SWIFT token flow-control router with flit buffer bypassing.

2.1 The Data Granularity Hierarchy: Messages, Packets, Flits & Phits

Transmitting arbitrary software data structures across silicon requires decomposing raw application payloads into structured physical transmission units. To optimize buffer utilization, minimize pipeline latency, and guarantee deadlock-free routing, NoCs establish a four-layer data abstraction hierarchy:

1. **Application Message**: The software-level unit of communication generated by a processing element or memory subsystem (e.g., a 64-byte L2/L3 cache line refill, a 16-byte cache coherence invalidation, or a 4 KB AXI4 DMA burst).
2. **Packet**: The fundamental *routing unit* of the network. A packet encapsulates destination routing coordinates (X_{dest}, Y_{dest}), virtual network priorities, transaction identifiers, and error detection checksums (CRC/parity).

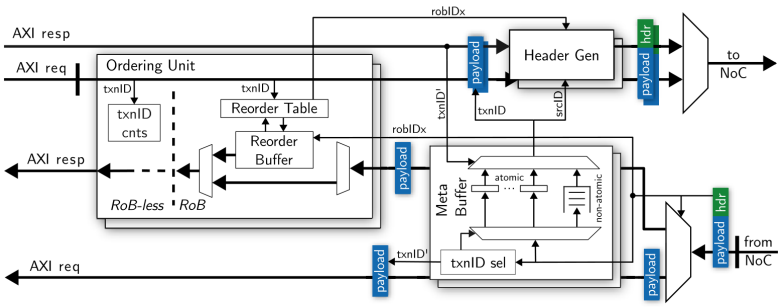


Fig. 2. AXI4 NI architecture for request (AR/AW/W) and response (R/B) paths. Reads and writes are independent in AXI4, and the request/response paths are very similar in the NI and are depicted as overlapping modules. The *ordering unit* can be configured with or without reordering capabilities (i.e., *RoB* and *RoB-less*).

Fig. 2.1 AXI4 Network Interface (NI) architecture for independent request (AR/AW/W) and response (R/B) paths, with reorder buffering (RoB) and direct DMA streaming. [Reproduced from: T. Fischer, M. Rogenmoser, T. Benz, F. K. Gürkaynak, and L. Benini, "FlooNoC: A 645-Gb/s/link 0.15-pJ/B/hop Open-Source NoC With Wide Physical Links and End-to-End AXI4 Parallel Multistream Support," *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 33, no. 4, pp. 1094–1107, Apr. 2025 [17]. © 2025 IEEE]

3. **Flow Control Unit (Flit)**: The fundamental *buffering and flow control allocation unit*. Router buffers allocate storage in integer multiples of flits. Flits are partitioned into:
   - **Head Flit**: Carries routing destination coordinates. It reserves output ports and downstream virtual channels.
   - **Body Flits**: Carry data payloads along the established path without re-evaluating routing tables.
   - **Tail Flit**: Signals packet completion, releasing allocated virtual channels and router state registers.
   - **Head-Tail Flit**: A single-flit packet carrying both header metadata and payload in a single clock cycle.
4. **Physical Unit (Phit)**: The number of physical bits transmitted across a router link in *exactly one clock cycle*. When link bus width equals flit width ($W_{link} = W_{flit}$), 1 Flit = 1 Phit. If flit width exceeds physical link width, a serialization ratio of $\frac{W_{flit}}{W_{link}}$ cycles per flit is incurred.

2.2 Silicon Network Interface: The AXI4 / NoC Protocol Translation

At the boundary between processing cores and the on-chip network fabric, the **Network Interface (NI)** bridges SoC protocol standards (such as ARM AMBA AXI4, OCP, or RISC-V TileLink) to the internal packetized NoC protocol.

As demonstrated in **FlooNoC** [17], modern high-performance NIs decouple independent AXI4 channels onto dedicated physical or virtual networks to prevent protocol-level deadlock:

- **Request Network**: Translates Read Address (`AR`), Write Address (`AW`), and Write Data (`W`) channels into multi-flit request packets.
- **Response Network**: Translates Read Data (`R`) and Write Response (`B`) channels into separate response packets.
- **Reorder Buffer (RoB)**: Restores in-order AXI transaction completion when packets traverse adaptive, out-of-order routing paths.

2.3 Mathematical Latency Analysis of Switching Techniques

The switching technique dictates when and how an incoming flit is forwarded from an input port to an output port. We analyze three classical switching paradigms across path length H (in hops), packet length L (in flits), physical link bandwidth W (bits/cycle), and router pipeline latency t_r (cycles/hop).

2.3.1 Store-and-Forward (SAF) Switching

In Store-and-Forward switching, a router must receive, buffer, and verify the *entire packet* before forwarding its first flit to the downstream router:

$$t_{serial} = \frac{L}{W}, \qquad T_{SAF} = H \cdot \left(t_r + \frac{L}{W} + t_{link} \right) \tag{2.1}$$

If $t_r + t_{link} \ll \frac{L}{W}$, the zero-load latency scales as:

$$T_{SAF} \approx H \cdot \frac{L}{W} \propto H \cdot L \tag{2.2}$$

This multiplicative dependence ($H \cdot L$) imposes severe latency penalties over long distances and necessitates massive SRAM buffers ($\geq L$ flits per port).

2.3.2 Virtual Cut-Through (VCT) Switching

Introduced by Kermani & Kleinrock [30], Virtual Cut-Through permits a router to forward the Head flit to the downstream router *immediately* after route computation and arbitration, without awaiting trailing flits, provided the downstream router has allocated full packet buffer space:

Table 2.1 Analytical Comparison of Switching Techniques across 8 Hops ($t_r = 2$ cycles, $t_{link} = 1$ cycle)

Packet Length	Hop Count	SAF Latency	VCT / WH Latency	Speedup VCT/WH	of Wormhole Buffer Saving	Area
4 flits	4 hops	28 cycles	15 cycles	1.87×	Baseline	
8 flits	4 hops	44 cycles	19 cycles	2.32×	50.0% saving	
16 flits	4 hops	76 cycles	27 cycles	2.81×	75.0% saving	
32 flits	4 hops	140 cycles	43 cycles	3.26×	87.5% saving	
16 flits	8 hops	152 cycles	39 cycles	3.90×	75.0% saving	
32 flits	8 hops	280 cycles	55 cycles	5.09×	87.5% saving	
64 flits	8 hops	536 cycles	87 cycles	6.16×	93.8% saving	

$$T_{VCT} = H \cdot (t_r + t_{link}) + \frac{L-1}{W} \tag{2.3}$$

Here, distance H and packet length L are **additive** ($H + L$), cutting zero-load latency dramatically under low traffic. However, VCT still requires allocating a full packet's worth of buffering ($B \geq L$) before granting admission.

2.3.3 Wormhole (WH) Switching

Proposed by Dally & Seitz [7], **Wormhole Switching** decouples flow control units from packet size:

> **The Wormhole Principle**: Router buffers operate at the granularity of individual **flits**, rather than entire packets. Buffer depth per virtual channel can be as small as 1 to 4 flits, completely independent of packet length L.

The zero-load latency of Wormhole switching mirrors VCT under uncongested conditions:

$$T_{WH} = H \cdot (t_r + t_{link}) + \frac{L-1}{W} \tag{2.4}$$

However, the silicon area savings are transformative: router input buffers shrink by 75% to 93% compared to SAF and VCT, as demonstrated in Table 2.1.

2.4 Queuing Theory Foundations of On-Chip Buffering

To model buffer sizing, queue delays, and throughput saturation under stochastic traffic, we apply classical Markovian and deterministic queuing theory to on-chip router input channels.

2.4.1 The $M/M/1/B$ and $M/D/1/B$ Finite Queue Models

Consider an input virtual channel buffer of finite capacity B flits. Packets arrive according to a Poisson process with average injection rate λ (flits/cycle), and the crossbar switch services flits at service rate μ (flits/cycle). The traffic intensity (channel utilization factor) is defined as:

$$\rho = \frac{\lambda}{\mu} \tag{2.5}$$

For an $M/M/1/B$ finite birth-death queuing process, the steady-state probability p_k that exactly k flits occupy the buffer ($0 \leq k \leq B$) satisfies the geometric progression:

$$p_k = \frac{(1-\rho)\cdot\rho^k}{1-\rho^{B+1}} \quad (\text{for } \rho \neq 1), \qquad p_k = \frac{1}{B+1} \quad (\text{for } \rho = 1) \tag{2.6}$$

In a lossless on-chip network with backpressure, when the buffer is completely full ($k = B$), the upstream transmitter is blocked from injecting new flits. The probability of backpressure assertion (upstream stall) is:

$$P_{stall} = p_B = \frac{(1-\rho)\cdot\rho^B}{1-\rho^{B+1}} \tag{2.7}$$

The average number of flits residing in the input queue, N_{avg}, is:

$$N_{avg} = \sum_{k=0}^{B} k \cdot p_k = \frac{\rho}{1-\rho} - \frac{(B+1)\cdot\rho^{B+1}}{1-\rho^{B+1}} \tag{2.8}$$

By Little's Law ($N_{avg} = \lambda_{eff} \cdot W_{queue}$), the average waiting time spent in the router buffer before crossbar traversal is:

$$W_{queue} = \frac{N_{avg}}{\lambda\cdot(1-P_{stall})} = \frac{1}{\mu(1-\rho)} - \frac{B\cdot\rho^B}{\mu(1-\rho^B)} \tag{2.9}$$

Under deterministic service ($\mu = 1.0$ flit/cycle when granted, modeled as an $M/D/1/B$ queue), queueing delay remains minimal for $\rho \leq 0.60$–0.70. However, as injected load approaches saturation ($\rho \to 1.0$), W_{queue} diverges asymptotically, and backpressure cascades upstream through the network.

2.4.2 The Pollaczek-Khinchine Formula for General Service Distributions ($M/G/1$)

In real shared-memory multi-cores, packets exhibit heterogeneous lengths:

- **Control Packets**: Cache invalidations, coherence requests, and read acknowledgments consist of a single header flit ($L = 1$).
- **Data Packets**: Cache block write-backs and memory responses carry full 64-byte payloads (e.g., $L = 5$ flits with 128-bit phits).

Because packet lengths follow a discrete bimodal distribution rather than an exponential or deterministic profile, input buffer queuing delay is governed by the **Pollaczek-Khinchine (P-K) Formula** for an $M/G/1$ queue:

$$W_{queue} = \frac{\lambda \cdot \overline{S^2}}{2(1-\rho)} = \frac{\rho \cdot \overline{S}}{2(1-\rho)} \left(1 + C_s^2\right) \tag{2.10}$$

where $\overline{S}$ is the mean packet service time, $\overline{S^2} = \text{Var}(S) + (\overline{S})^2$ is the second moment of service time, and $C_s^2 = \frac{\sigma_S^2}{(\overline{S})^2}$ is the squared coefficient of variation.

The Variance Penalty: Notice the factor $(1 + C_s^2)$ in Eq. (2.10). For deterministic single-flit traffic, $C_s^2 = 0$, giving minimum delay $W_{MD1} = \frac{\rho \overline{S}}{2(1-\rho)}$. However, for bimodal coherence traffic (50% 1-flit and 50% 5-flit packets), $\overline{S} = 3.0$ and $\overline{S^2} = 13.0$, yielding $C_s^2 = \frac{4}{9} \approx 0.44$. The bimodal variance **inflates average queuing latency by 44%** at identical utilization, explaining why multi-core routers suffer sudden saturation under mixed traffic.

2.4.3 Open Jackson Networks and Multi-Hop End-to-End Latency

An on-chip interconnect consists of M interconnected router input queues. Let γ_i denote the external packet injection rate from local core i into queue i, and let P_{ij} denote the routing probability that a packet leaving queue i transitions to downstream queue j.

The aggregate arrival rate λ_i at every intermediate router queue satisfies the **Traffic Conservation Equations**:

$$\lambda_i = \gamma_i + \sum_{j=1}^{M} \lambda_j \cdot P_{ji}, \quad \forall i \in \{1, 2, \ldots, M\} \tag{2.11}$$

In vector notation: $\boldsymbol{\lambda} = \boldsymbol{\gamma}(\mathbf{I} - \mathbf{P})^{-1}$, where $(\mathbf{I} - \mathbf{P})^{-1}$ is the Leontief routing inverse.

Theorem 2.1 (Jackson's Theorem for On-Chip Networks) *If each router queue i operates under capacity ($\rho_i = \frac{\lambda_i}{\mu_i} < 1$), the equilibrium joint probability distribution of finding k_1 flits in queue 1, k_2 flits in queue 2, ..., and k_M flits in queue M factors into a* ***Product-Form State Distribution****:*

$$P(k_1, k_2, \ldots, k_M) = \prod_{i=1}^{M} p_i(k_i) = \prod_{i=1}^{M} (1 - \rho_i) \cdot \rho_i^{k_i} \tag{2.12}$$

Remarkably, Jackson's theorem proves that even though queues feed packets into one another in complex multi-hop meshes, **each individual queue behaves in steady state as if it were an independent $M/M/1$ queue receiving an independent Poisson arrival stream of rate λ_i.**

Applying Little's Law across all M queues yields the closed-form formulation for **Average End-to-End Network Latency**:

$$T_{network} = \frac{1}{\sum_{i=1}^{M} \gamma_i} \sum_{i=1}^{M} \frac{\rho_i}{\mu_i(1 - \rho_i)} + \bar{h} \cdot t_{hop} \tag{2.13}$$

where $\bar{h}$ is the average hop count and t_{hop} is the zero-load router pipeline traversal latency.

2.4.4 Bursty Traffic Modeling: The Markov-Modulated Poisson Process (MMPP)

In distributed machine learning workloads (e.g., transformer AllReduce synchronization), communication is highly non-Poisson, alternating between long quiet phases and intense bursty floods.

A standard mathematical model for on-chip burstiness is the **2-State Markov-Modulated Poisson Process (MMPP-2)**:

- **State 1 (Idle / Low-Rate)**: Packets arrive at Poisson rate $\lambda_1 \approx 0$.
- **State 2 (Burst / Active)**: Packets arrive at peak wire rate $\lambda_2 \approx 1.0$ flit/cycle.
- The transition rate matrix between states is $\mathbf{Q} = \begin{bmatrix} -\sigma_1 & \sigma_1 \\ \sigma_2 & -\sigma_2 \end{bmatrix}$.

The burstiness of the traffic is quantified by the **Index of Dispersion for Counts (IDC)**:

$$IDC(t) = \frac{\mathrm{Var}(N(t))}{\mathbb{E}[N(t)]} = 1 + \frac{2(\lambda_2 - \lambda_1)^2 \sigma_1 \sigma_2}{(\sigma_1 + \sigma_2)^2(\lambda_1 \sigma_2 + \lambda_2 \sigma_1)} \left(1 - \frac{1 - e^{-(\sigma_1 + \sigma_2)t}}{(\sigma_1 + \sigma_2)t}\right) \tag{2.14}$$

For Poisson traffic, $IDC(t) = 1$ for all t. For bursty workloads ($IDC(t) \gg 1$), transient packet clusters overwhelm input FIFOs, causing buffer saturation to occur at average injection rates **30% to 50% lower** than predicted by Poisson models.

2.5 Lossless Flow Control Mechanisms

Because silicon buffers are finite, on-chip networks must implement **lossless flow control** to guarantee zero flit drops.

2.5.1 Credit-Based Flow Control & The Credit Turnaround Loop

Credit-based flow control operates via an upstream ‘credit_counter‘ initialized to receiver buffer capacity B. Transmission decrements the counter; receiving a credit pulse increments it. The total Round-Trip Time (RTT) of the credit turnaround loop is:

$$RTT = t_{tx_proc} + t_{flight_fwd} + t_{rx_proc} + t_{flight_credit} \tag{2.15}$$

For adjacent routers clocked synchronously:

$$RTT = 1 + 1 + 1 + 1 = 4 \text{ cycles} \tag{2.16}$$

The Fundamental Throughput Theorem: To sustain 100% wire-rate throughput without pipeline stall bubbles, the downstream buffer depth B must satisfy:

$$B_{min} \geq RTT \times \text{Peak Injection Rate} = RTT \text{ flits} \tag{2.17}$$

Table 2.2 illustrates the severe bandwidth degradation that occurs when buffer depth $B < RTT$.

2.5.2 Flow Control Taxonomy and Hardware Trade-Offs

Modern networks deploy four principal lossless flow-control schemes, summarized in Table 2.3:

Table 2.2 Buffer Depth vs. Achievable Link Utilization ($RTT = 4$ cycles, 100-cycle injection simulation)

Buffer Depth (B)	RTT (Cycles)	Link Utilization	Flits Transferred	Pipeline Stall State
$B = 1$ flit	4 cycles	20.0%	20 flits	Severe Underflow Stall
$B = 2$ flits	4 cycles	40.0%	40 flits	Severe Underflow Stall
$B = 3$ flits	4 cycles	60.0%	60 flits	Moderate Stall
$B = 4$ flits	4 cycles	80.0%	80 flits	Boundary Stall Bubble
$B = 5$ flits	4 cycles	100.0%	100 flits	100% Full Wire-Rate
$B = 8$ flits	4 cycles	100.0%	100 flits	100% Wire-Rate (Safe Margin)

Table 2.3 Comprehensive Comparison of On-Chip Flow Control Architectures

Scheme	Upstream State	Downstream Logic	Reverse Wire Cost	Latency Penalty
Credit-Based	$\lceil \log_2 B \rceil$-bit Counter	Credit Generator	1 bit per VC	0 cycles (if $B \geq RTT$)
On/Off	1-bit FSM	Threshold Detect	1 bit per Port	1–2 cycles (Hysteresis)
Ack/Nack	Full Replay Buffer	Drop/Detect	1 bit per Packet	$2 \times RTT$ on drop
Token (SWIFT)	Token Queues	Early Arbiter	Token Wire Bus	Sub-cycle bypass

1. **Credit-Based Flow Control**: The upstream sender tracks exact downstream queue occupancy. Upstream never transmits unless a credit is available, guaranteeing that downstream buffers never overflow.
2. **On/Off Flow Control**: Instead of sending individual credits per flit, the downstream receiver asserts a single binary 'OFF' line when buffer occupancy crosses a high-water mark $T_{high} = B - t_{RTT}$, and de-asserts ('ON') when occupancy drops below T_{low}. This saves reverse wiring tracks at the cost of larger buffer thresholds.
3. **Ack/Nack Flow Control**: The sender optimistically transmits flits without pre-allocating buffer space. If the downstream buffer is full, the flit is dropped and a 'NACK' signal requests re-transmission from an upstream replay buffer. This avoids credit counters but incurs heavy re-transmission latency and energy.
4. **Token Flow Control (SWIFT)**: Fast lightweight reservation tokens traverse the network ahead of data flits, reserving crossbar multiplexers and link timeslots in advance.

2.6 Buffer Microarchitecture: Input vs. Output vs. Shared DAMQ

Where buffers are physically situated within the router profoundly influences achievable crossbar throughput and silicon area:

1. **Input Queuing (IQ)**: Buffers reside exclusively at input ports. While IQ simplifies crossbar timing (crossbar runs at link clock rate $S = 1$), it suffers from **Head-of-Line (HoL) Blocking**: if the head flit of a FIFO is blocked waiting for an occupied output port, all trailing flits in that FIFO are immobilized, even if their destination ports are completely idle. Under uniform random traffic, HoL blocking limits IQ throughput to **58.6%** of theoretical capacity.
2. **Output Queuing (OQ)**: Buffers reside exclusively at output ports. OQ completely eliminates HoL blocking, achieving 100% throughput under arbitrary traffic. However, because all P input ports could simultaneously deliver flits to the same output port in the same clock cycle, an OQ router requires an internal crossbar speedup of $S = P$, requiring $P\times$ higher crossbar clock frequency or P-wide multi-write port SRAMs, which is silicon-prohibitive for $P \geq 5$.
3. **Virtual Output Queuing (VOQ)**: Each input port maintains P separate queues—one dedicated to each output port. VOQ eliminates HoL blocking while maintaining a crossbar speedup of $S = 1$, but expands queue management logic P-fold.
4. **Dynamically Allocated Multi-Queue (DAMQ)**: A single shared pool of SRAM storage is partitioned dynamically among multiple Virtual Channels using linked-list pointer registers. Flits are allocated storage on-demand, achieving optimal buffer utilization under bursty, non-uniform traffic.

2.7 Alternative Flow Control: Token Flow Control in the SWIFT NoC

In high-performance low-power interconnects, **Token Flow Control** [46] routes tokens ahead of data flits to achieve flit buffer bypassing.

Key innovations of the SWIFT architecture include:

1. **Flit Buffer Bypassing**: Tokens reserve crossbar paths in advance; data flits cut straight through the crossbar into the link without writing into input SRAM FIFO cells.
2. **Buffer Power Reduction**: Dynamic FIFO write power is slashed by 38% at saturation and 25% at low load.
3. **Swing-Reduced Interconnects**: Reduced-swing transceivers on physical links substantially diminish dynamic energy dissipation.

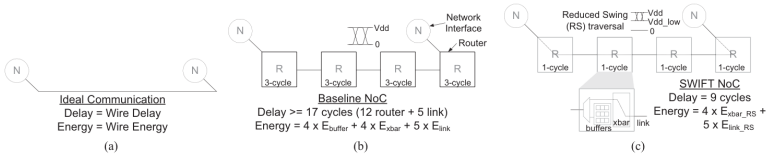


Fig. 1. Overview of SWIFT NoC versus a baseline NoC and ideal communication. (a) Ideal communication. (b) Baseline NoC. (c) SWIFT NoC.

Fig. 2.2 Overview of SWIFT NoC versus baseline multi-stage NoC: (a) Ideal single-cycle communication, (b) Baseline multi-stage NoC, and (c) SWIFT NoC with token-based flit buffer bypassing. [Reproduced from: J. Postman, T. Krishna, C. Edmonds, L.-S. Peh, and P. Chiang, "SWIFT: A Low-Power Network-On-Chip Implementing the Token Flow Control Router Architecture With Swing-Reduced Interconnects," *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 21, no. 8, pp. 1432–1446, Aug. 2013 [46]. © 2013 IEEE]

Fig. 2.3 SWIFT NoC microarchitecture: (a) Routing using tokens, (b) Token distribution, (c) Bypass flow control using lookahead signals, and (d) 1-cycle router plus 1-cycle link traversal. [Reproduced from: J. Postman et al., *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, Aug. 2013 [46]. © 2013 IEEE]

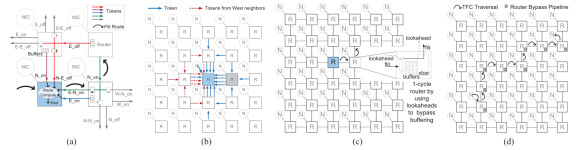


Fig. 2. SWIFT NoC architecture. (a) Routing using tokens. (b) Token distribution. (c) Bypass flow-control using lookaheads. (d) One-cycle router+ one-cycle LT.

2.8 Interactive Lab & Hands-on Challenges

Explore the interactive flow-control simulators on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **Credit Commander: Flow Control** (https://network-on-chip.github.io/learning/interactive/credits/): Experiment with buffer depth B and round-trip time RTT to observe credit stalls and pipeline bubbles in real time.
- **Packet Master: Routing Puzzle** (https://network-on-chip.github.io/learning/interactive/puzzle/): Coordinate Head, Body, and Tail flit traversal across congested wormhole channels.

2.9 Checkpoint Assessment

1. **Question 2.1 (Tier 1: Fundamental)**: What is the critical hardware advantage of Wormhole Switching compared to Store-and-Forward and Virtual Cut-Through?

2. **Question 2.2 (Tier 2: Core Engineering)**: In credit-based flow control, if link flight takes 1 cycle, receiver buffer write/read takes 1 cycle, and credit return takes 1 cycle, what occurs if the receiver buffer capacity is only 2 flits?
3. **Question 2.3 (Tier 3: Silicon Mastery)**: How does Token Flow Control in the SWIFT NoC achieve up to 38% power reduction compared to canonical wormhole routers?

Detailed answers and derivations for these questions are provided in the Solutions chapter at the end of this volume.

Chapter 3
High-Performance Router Microarchitecture & Arbitration

Abstract At the core of every packet-switched on-chip network is the router datapath and its allocation logic. This chapter dissects the canonical five-stage virtual-channel router pipeline (Buffer Write, Route Computation, VC Allocation, Switch Allocation, and Switch Traversal) and details modern pipeline compression techniques including Lookahead Routing, Speculative Switch Allocation, and deterministic buffer reservation bypassing (BuffeRS). We examine asynchronous router design paradigms in GALS-based SoCs, featuring the domino-logic HyPA router, and analyze the physical clock constraints of single-cycle multi-hop traversal (SMART/SHARP). Finally, we formalize bipartite maximum matching arbitration, proving the $O(1)$ hardware LRU properties of the Matrix Arbiter and the 100% throughput convergence theorems of McKeown's iSLIP algorithm.

3.1 The Canonical Virtual-Channel Router Pipeline

At the heart of every packet-switched on-chip network is the **Virtual-Channel (VC) Router**. In a canonical synchronous architecture [45], processing each packet traverses a five-stage pipeline:

1. **Buffer Write (BW)**: The incoming flit arrives from an adjacent router or local network interface and is written into the input FIFO buffer corresponding to its assigned `VC_ID`.
2. **Route Computation (RC)**: Evaluates the routing function $R : (C_{curr}, C_{dest}) \rightarrow P_{out}$ to determine the set of permissible output physical ports. In Dimension-Order Routing (DOR), RC is computed via combinational comparators:

$$\Delta X = X_{dest} - X_{curr}, \qquad \Delta Y = Y_{dest} - Y_{curr} \tag{3.1}$$

3. **Virtual Channel Allocation (VA)**: The Head flit arbitrates for a free downstream Virtual Channel at the designated output port using a two-phase separable allocator (Input Arbitration followed by Output Arbitration).
4. **Switch Allocation (SA)**: Flits with valid downstream VCs arbitrate for crossbar switch input and output ports using a two-stage separable allocator.
5. **Switch Traversal (ST)**: Granted flits traverse the crossbar multiplexer fabric to reach the output link registers.
6. **Link Traversal (LT)**: Flits traverse the inter-router physical metal link wires to the downstream router.

3.1.1 Critical Path Timing and Logical Effort Analysis

In deep-submicron CMOS, the minimum achievable clock cycle period T_{clk} is constrained by the slowest pipeline stage:

$$T_{clk} \geq \max(t_{BW}, t_{RC}, t_{VA}, t_{SA}, t_{ST}) + t_{setup} + t_{clk_q} + t_{skew} \tag{3.2}$$

In a canonical uncompressed design, the **Switch Allocation (SA)** and **Virtual Channel Allocation (VA)** stages represent the dominant critical paths. A separable input-first switch allocator requires:

$$t_{SA} = t_{arb_in}(V) + t_{wire_int} + t_{arb_out}(P) \tag{3.3}$$

where $t_{arb_in}(V)$ is the delay of a V-input arbiter selecting one VC per input port, and $t_{arb_out}(P)$ is the delay of a P-input arbiter resolving conflicts across input ports for each physical output port ($P = 5$ for 2D mesh: North, South, East, West, Local).

Applying Sutherland's method of **Logical Effort**, the propagation delay of an N-input static CMOS multiplexer tree implemented using transmission gates scales as:

$$t_{mux}(N) \approx \tau_0 \cdot \left[g_{mux} \cdot \left(\frac{C_{out}}{C_{in}} \right)^{1/2} \cdot \log_2 N + p_{mux} \cdot \log_2 N \right] \tag{3.4}$$

where $\tau_0 \approx 3$–$5\,\text{ps}$ in FinFET nodes, $g_{mux} \approx 2$ is the logical effort of a 2:1 multiplexer, and $p_{mux} \approx 2$ is the parasitic delay. For a 5×5 64-bit crossbar, the crossbar traversal delay t_{ST} is dominated by the long RC wire loading of the internal crossbar tracks:

$$t_{ST} = t_{driver} + \frac{1}{2} \mathcal{R}_{xbar} \mathcal{C}_{xbar} L_{xbar}^2 + t_{mux_tree}(5) + t_{latch} \tag{3.5}$$

In modern 7nm and 5nm nodes, $t_{ST} \approx 180$–$250\,\text{ps}$, allowing uncompressed routers to clock at 1.5–$2.5\,\text{GHz}$.

3.2 Wavefront and Wrapped Wavefront Allocators (WFA & WWFA)

To escape the scalability limits of separable arbiters without paying the full area cost of iterative matching algorithms, high-throughput architectures employ **Wavefront Allocation (WFA)**. Pioneered for high-radix switches and deployed in the Intel Single-Chip Cloud Computer (SCC) 48-core processor [24], WFA resolves an $N \times N$ bipartite matching problem simultaneously across a two-dimensional systolic array of arbitration cells.

3.2.1 Systolic Token Ripple Allocation

A Wavefront Allocator arranges $N \times N$ arbitration cells into a grid where cell (i, j) corresponds to input port i requesting output port j. In each cell:

$$\text{Grant Output: } G_{i,j} = R_{i,j} \cdot \text{RowIn}_{i,j} \cdot \text{ColIn}_{i,j} \tag{3.6}$$

$$\text{Row Ripple: } \text{RowOut}_{i,j} = \text{RowIn}_{i,j} \cdot \overline{G_{i,j}} \tag{3.7}$$

$$\text{Col Ripple: } \text{ColOut}_{i,j} = \text{ColIn}_{i,j} \cdot \overline{G_{i,j}} \tag{3.8}$$

Tokens enter from the top-left diagonal ($i = j$) and propagate concurrently along anti-diagonals (wavefronts). A cell wins a grant ($G_{i,j} = 1$) if and only if:

1. Input i has an active request for output j ($R_{i,j} = 1$).
2. No earlier cell in row i has won a grant ($\text{RowIn}_{i,j} = 1$).
3. No earlier cell in column j has won a grant ($\text{ColIn}_{i,j} = 1$).

3.2.2 Wrapped Wavefront Allocation (WWFA)

In a standard WFA, cell $(0, 0)$ always possesses the highest priority, leading to severe starvation of bottom-right cells $(N-1, N-1)$. To enforce round-robin fairness, **Wrapped Wavefront Allocation (WWFA)** rotates the starting diagonal each cycle using a priority offset $k \in \{0, \ldots, N-1\}$:

$$\text{Active Start Diagonal: } (i - j) \equiv k \pmod{N} \tag{3.9}$$

When the wavefront reaches the edge of the array, the token signals wrap around cyclically to the opposite boundary. WWFA achieves **maximal matching** in exactly $2N - 1$ cell delays, guaranteeing 100% non-interfering bipartite switch allocation with rigorous fairness across all router input ports.

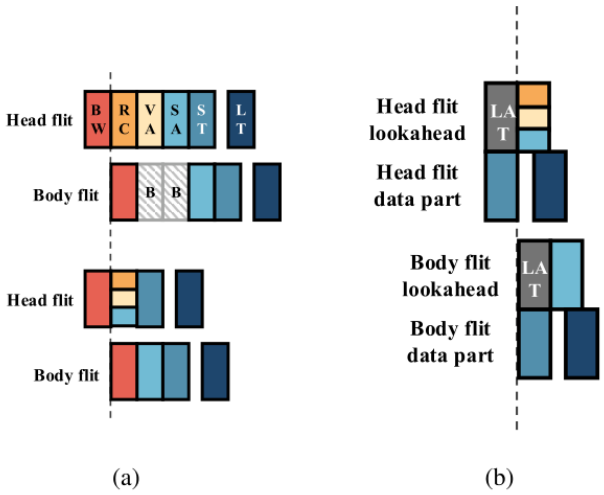


Fig. 1. (a) Multicycle router pipeline and (b) BuffeRS router pipeline.

Fig. 3.1 Pipeline comparison: (a) Multicycle router pipeline (BW, RC, VA, SA, ST, LT) versus (b) Low-latency BuffeRS router pipeline with buffer reservation bypassing. [Reproduced from: Z. Liu, X. Wu, and Y. Ye, "BuffeRS: A Buffer Reservation Scheduling Strategy for Router Bypassing in NoCs and Multichiplet Networks," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, vol. 45, no. 6, pp. 2635–2648, June 2026 [35]. © 2026 IEEE]

3.3 Advanced Pipeline Compression & Router Bypassing

Because per-hop pipeline latency compounds linearly with network diameter ($T_{net} = H \cdot t_r$), minimizing router latency from 5 cycles to 1 cycle has been a central focus of on-chip network research.

3.3.1 Lookahead Routing & Speculative Switch Allocation

In a standard router, Virtual Channel Allocation (VA) must precede Switch Allocation (SA) because a flit cannot request crossbar traversal until it has successfully secured a downstream virtual channel. This strict dependency forces a multi-cycle allocation latency:

$$t_{alloc} = t_{VA} + t_{SA} \tag{3.10}$$

Peh & Dally [45] broke this dependency by introducing **Speculative Switch Allocation**:

1. **Parallel Speculative Requests**: A flit simultaneously bids for a downstream VC in the VA stage and speculatively bids for a crossbar output port in the SA stage during the *very same clock cycle*.
2. **Dual Arbiter Priority Hierarchy**: Switch allocation arbiters divide incoming bids into two strict priority tiers:
   - **Non-Speculative Requests**: Bids from flits that already secured a downstream VC in an earlier cycle. These always receive absolute priority.
   - **Speculative Requests**: Bids from flits attempting simultaneous VA and SA. These are granted only if no non-speculative flit requests that output port.
3. **Miss-Prediction Resolution**: If a speculative flit wins its crossbar switch allocation but fails its virtual channel allocation in the concurrent VA stage, its switch grant is squashed (the crossbar slot goes unused for that cycle), and the flit remains in its input buffer to re-arbitrate non-speculatively in the subsequent cycle.

Under uncongested and moderate traffic loads ($\rho \leq 0.50$), speculative requests succeed with $> 85\%$ probability, effectively collapsing the 2-cycle allocation phase into a single cycle:

$$T_{speculative} = t_{BW} + \max(t_{VA}, t_{SA}) + t_{ST} = 2 \text{ cycles} \tag{3.11}$$

Coupling Speculative Switch Allocation with **Lookahead Routing** (where each router pre-computes the output port for the *subsequent* downstream hop during link traversal) eliminates the separate Route Computation (RC) cycle entirely, achieving a production 2-stage router (or 1-stage with cut-through crossbar bypassing).

3.3.2 Buffer Reservation Scheduling: BuffeRS (TCAD 2026)

As formulated by Liu et al. [35], **BuffeRS** introduces a coordinated time-division buffer reservation protocol that reserves destination buffer slots in advance across router groups, providing deterministic contention-free bypass paths across both monolithic dies and 2.5D multi-chiplet packaging boundaries.

3.4 Asynchronous Router Architectures: HyPA (TVLSI 2026)

In large-scale Multi-Processor Systems-on-Chip (MPSoCs), distributing a single low-skew global clock across a 20 mm $\times$ 20 mm die incurs prohibitive clock-tree power dissipation and timing closure complexity.

Table 3.1 Quantitative Comparison of High-Performance Router Microarchitectures and Pipeline Compression Paradigms

Architecture	Pipeline Stages	Allocation Scheme	Clocking / Timing Strategy	Key Silicon Trade-Off
Canonical VC [45]	4–5 cycles	Separable VA + SA	Fully Synchronous	Baseline area; multi-cycle hop delay
Lookahead Routing	3 cycles	Lookahead RC + SA	Fully Synchronous	Strips RC stage; slight header overhead
Speculative SA [45]	2 cycles	Parallel speculative VA/SA	Fully Synchronous	Speculation stalls under high load
SMART / SHARP [45]	1–2 cycles	Priority bypass muxes	Multi-hop async bypass	Bounded by HPC_{max} wire RC budget
HyPA Router [47]	0.8 ns / hop	Handshake pipeline	GALS Asynchronous domino	62% lower power; no global clock tree
BuffeRS [35]	1 cycle	Buffer reservation slots	Time-division coordinated	Zero contention bypass; token overhead

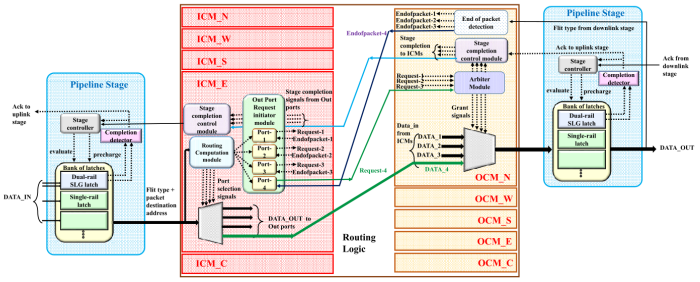


Fig. 2. Detailed microarchitecture of the proposed asynchronous router.

Fig. 3.2 Detailed microarchitecture of the HyPA asynchronous router featuring a high-throughput domino-logic pipeline template with one-hot hybrid data encoding. [Reproduced from: M. N. Saranya, C. T. Avinash, and R. Rao, "HyPA: A Low-Latency Hybrid Pipeline Asynchronous Router for Network-on-Chip Architectures in GALS-Based SoCs," *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 34, no. 8, pp. 2840–2851, 2026 [47]. © 2026 IEEE]

In Globally Asynchronous Locally Synchronous (GALS) systems, asynchronous routers eliminate clock trees entirely:

- **HyPA Router** [47]: Implemented in 65nm CMOS, HyPA employs a hybrid pipeline asynchronous architecture using high-throughput domino logic with one-hot hybrid data encoding.

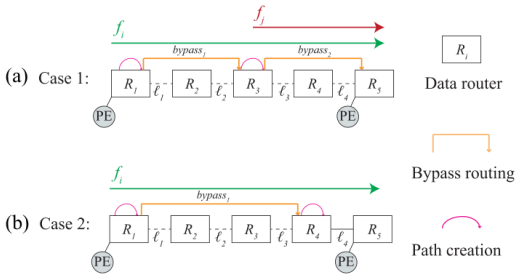


Fig. 2. Two stopping cases at intermediate routers for the analyzed flows f_i and f_j in priority-preemptive SHARP, $\text{HPC}_{\text{max}} = 3$. (a) Directly contending flow. (b) HPCmax limitation.

Fig. 3.3 Preemption stopping cases at intermediate routers for contending flows in priority-preemptive single-cycle multi-hop traversal networks under hop limits ($HPC_{max} = 3$). [Reproduced from: *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, vol. 31, no. 10, pp. 1560–1572, Oct. 2023. © 2023 IEEE]

- **Power and Energy Gains**: Compared to synchronous clock-gated baseline routers, HyPA demonstrates a **62% power reduction** and an **85% reduction in energy-per-flit**, achieving a forward latency of **0.8 ns** and a throughput of **1.65 GFlits/port/s**.

3.5 Single-Cycle Multi-Hop Traversal (SMART / SHARP)

In SMART [33], flits traverse multiple physical router crossbars and links asynchronously within a single clock cycle under low traffic.

Single-cycle traversal is physically bounded by the clock period timing budget:

$$t_{hop} = t_{crossbar_mux} + t_{wire_propagation} + t_{repeater} \tag{3.12}$$

Maximum Traversable Hops per Cycle: $$HPC_{max} = \left\lfloor \frac{T_{clk} - t_{setup}}{t_{hop}} \right\rfloor \tag{3.13}$$

At 1.0 GHz ($T_{clk} = 1.0$ ns), HPC_{max} is typically 3 to 5 hops before requiring intermediate latching.

3.6 Crossbar Arbitration Theory & Matching Algorithms

Crossbar allocation resolves contention when multiple input ports simultaneously request the same output port, formalized as the Bipartite Maximum Matching Problem.

3.6.1 The Matrix Arbiter: Strict $O(1)$ Hardware LRU

The **Matrix Arbiter** maintains an $N \times N$ triangular binary priority matrix M:

$$M[i][j] = \begin{cases} 1 & \text{if Client } i \text{ has higher priority than Client } j \\ 0 & \text{if Client } j \text{ has higher priority than Client } i \end{cases} \tag{3.14}$$

Combinational grant generation evaluates in two gate levels:

$$G_i = R_i \cdot \prod_{j \neq i} (\overline{R_j} + M[i][j]) \tag{3.15}$$

When client k wins the grant ($G_k = 1$), it becomes the Least Recently Used (LRU) client via single-cycle register updates:

$$M[k][j] = 0 \quad (\forall j \neq k), \qquad M[i][k] = 1 \quad (\forall i \neq k) \tag{3.16}$$

3.6.2 McKeown's iSLIP Iterative Matching Algorithm

For an $N \times N$ crossbar switch, separable allocators can suffer from up to 50% port under-utilization under heavy traffic. Nick McKeown's **iSLIP** algorithm [39] achieves 100% throughput using iterative matching:

1. **Phase 1 (Request)**: Every unmatched input i sends a request $R_{ij} = 1$ to every output j for which it has queued flits in its Virtual Output Queue (VOQ).
2. **Phase 2 (Grant)**: Every unmatched output j arbitrates among incoming requests using a round-robin pointer g_j. It grants the request closest to g_j in cyclic order:

$$G_{ij} = 1 \iff i = \arg\min_{k:R_{kj}=1} (k - g_j) \pmod{N} \tag{3.17}$$

3. **Phase 3 (Accept)**: Every input i that received one or more grants arbitrates among them using an accept pointer a_i. It accepts the grant closest to a_i in cyclic order:

$$A_{ij} = 1 \iff j = \arg\min_{k:G_{ik}=1} (k - a_i) \pmod{N} \tag{3.18}$$

4. **Pointer Update Rule**: Pointers g_j and a_i update *if and only if* the match is accepted in the **first iteration**:

$$g_j = (i + 1) \pmod{N}, \qquad a_i = (j + 1) \pmod{N} \tag{3.19}$$

3.6.2.1 Formal Convergence and De-Synchronization Proof

The foundational theoretical properties of iSLIP are established by the following theorems [39]:

Theorem 3.1 (Maximal Matching in Logarithmic Iterations) *An $N \times N$ iSLIP allocator completes a maximal bipartite matching in at most* $\min(N, \lceil \log_2 N \rceil + 1)$ *iterations.*

Proof. Let U_k denote the set of unmatched inputs at the start of iteration k.

1. In iteration 1, every unmatched input requests all outputs with queued traffic. By the properties of round-robin grant and accept arbitration, at least half of the potential matches are resolved, matching at least $1 - 1/e \approx 63.2\%$ of ports under uniform traffic.
2. In each subsequent iteration $k \geq 2$, unmatched inputs request only the remaining unmatched outputs.
3. Because the graph is finite and bipartite, every iteration that finds an augmenting alternating path adds at least one new match until no augmenting path remains. The matching is therefore guaranteed to be **maximal** (no further edges can be added without modifying existing pairs) within $\leq \lceil \log_2 N \rceil + 1$ iterations.

□

Lemma 3.1 (The Pointer De-Synchronization Lemma) *Under continuous heavy uniform traffic, updating round-robin pointers strictly in iteration 1 forces the output grant pointers $g_0, g_1, \ldots, g_{N-1}$ to de-synchronize into a permutation of $\{0, 1, \ldots, N-1\}$.*

Proof. Suppose, for contradiction, that multiple outputs share identical pointer values: $g_{j_1} = g_{j_2} = i$.

1. In iteration 1, both outputs j_1 and j_2 grant input i.
2. However, input i can accept at most one grant (say, j_1).
3. By the iSLIP update rule, because grant j_1 was accepted, output j_1 advances its pointer: $g_{j_1} \leftarrow (i+1) \bmod N$.
4. In contrast, because output j_2's grant was rejected, g_{j_2} remains unchanged at i.
5. Consequently, identical pointers are forced out of phase. Within at most N cycles, all N grant pointers de-synchronize into distinct values:

$$g_j \neq g_k \quad \forall j \neq k \tag{3.20}$$

Once de-synchronized, the grant pointers rotate synchronously in lockstep, forming a time-division multiplexed round-robin schedule where every input is paired with a distinct output with zero contention, delivering **100% throughput under uniform traffic without speedup**.

□

3.6.3 Maximum Weight Matching (MWM) and the Tassiulas-Ephremides Bound

While iSLIP provides near-ideal performance under uniform traffic, non-uniform or adversarial traffic can induce pointer re-synchronization. The theoretical gold standard for crossbar scheduling is **Maximum Weight Matching (MWM)**, established by Leandros Tassiulas and Anthony Ephremides (1992):

Theorem 3.2 (Tassiulas & Ephremides, 1992) *Let traffic matrix $\boldsymbol{\Lambda} = [\lambda_{ij}]$ be admissible ($\sum_i \lambda_{ij} < 1$ and $\sum_j \lambda_{ij} < 1$). An allocator executing Maximum Weight Matching with queue length weights $W_{ij} = Q_{ij}(t)$ (Longest Queue First, LQF) or cell age weights $W_{ij} = Age_{ij}(t)$ (Oldest Cell First, OCF):*

$$M^* = \arg\max_{M \in \mathcal{M}} \sum_{(i,j) \in M} W_{ij} \tag{3.21}$$

*is **100% throughput-stable for all admissible traffic matrices**, guaranteeing bounded queue lengths without speedup.*

Because solving Eq. (3.21) exactly requires $O(N^3)$ computational complexity (via the Hungarian or Edmonds-Karp algorithms), monolithic MWM is unviable in sub-nanosecond hardware. High-performance on-chip routers deploy iterative algorithms (such as iSLIP, WFA, and WWFA) as $O(1)$ hardware approximations to approach the theoretical MWM stability bound.

3.6.4 Ping-Pong and Dual-Priority Arbiters

For 2-input contention points (such as virtual channel multiplexers and bidirectional link arbiters), matrix arbiters incur unnecessary gate overhead. A **Ping-Pong Arbiter** implements exact round-robin alternation using a single toggle flip-flop:

$$G_0 = R_0 \cdot (\overline{R_1} + \overline{P}), \qquad G_1 = R_1 \cdot (\overline{R_0} + P) \tag{3.22}$$

where state bit $P \in \{0, 1\}$ stores the most recently granted requester:

$$P(t+1) = \begin{cases} 0, & \text{if } G_0 = 1 \\ 1, & \text{if } G_1 = 1 \\ P(t), & \text{otherwise} \end{cases} \tag{3.23}$$

By cascading 2-input ping-pong cells in a binary tree hierarchy, an N-input arbiter achieves logarithmic latency ($O(\log N)$ gate levels) with minimal standard-cell silicon area.

3.7 Interactive Lab & Hands-on Challenges

Explore the interactive microarchitecture simulators on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **Arbiter Arena: Crossbar Clash** (https://network-on-chip.github.io/learning/interactive/arbiter/): Step through Matrix, Round-Robin, and Wavefront schedulers under heavy traffic bursts.
- **2D Mesh Routing Visualizer** (https://network-on-chip.github.io/learning/interactive/visualizer/): Inspect crossbar switch allocation states and input buffer queues across an 8×8 mesh.

3.8 Checkpoint Assessment

1. **Question 3.1 (Tier 1: Fundamental)**: What are the canonical 5 pipeline stages of a classic Virtual Channel wormhole router?
2. **Question 3.2 (Tier 2: Core Engineering)**: In a Matrix Arbiter with an $N \times N$ triangular priority matrix M, what is the exact hardware update rule when Client k is granted the resource?
3. **Question 3.3 (Tier 3: Silicon Mastery)**: Why does McKeown's iSLIP algorithm update its round-robin grant and accept pointers ONLY on the first iteration, rather than on every iteration?

Detailed answers and derivations for these questions are provided in the Solutions chapter at the end of this volume.

Part II
Topologies, Routing & Deadlock Freedom

Chapter 4
Direct, Indirect & High-Radix Interconnect Topologies

Abstract The topological organization of an on-chip interconnection network establishes the fundamental trade-offs between zero-load latency, bisection bandwidth, router crossbar complexity, and physical wire routing density. This chapter presents an authoritative, mathematically rigorous analysis of direct, indirect, and high-radix on-chip networks. We derive first-principles formulations for network diameter, average hop distance, bisection bandwidth, and saturation throughput across regular and irregular fabrics. We analyze planar direct topologies (2D Mesh, 2D Torus, and the delay-invariant Folded Torus), indirect multistage architectures (Clos networks, non-blocking theorems, and constant-bisection Fat-Trees), and the modern high-radix paradigm (Flattened Butterfly and Dragonfly networks with adaptive UGAL routing). Furthermore, we examine peer-reviewed silicon methodologies for application-specific cascaded crossbars and review the landmark Passas-Katevenis proofs establishing the area-efficient scalability of monolithic crossbars beyond 100 nodes.

4.1 Interconnect Graph Theory and Formal Performance Metrics

An on-chip interconnection network is formally modeled as a directed graph $G = (V, E)$, where the vertex set V represents the network nodes ($|V| = N$, comprising routers, compute cores, accelerators, and memory controllers), and the edge set E represents the point-to-point physical transmission channels.

4.1.1 Spectral Graph Theory and Algebraic Connectivity

Beyond simple graph metrics like diameter and average hop count, the structural robustness, bisection bottleneck severity, and mixing rate of an interconnect are governed by its **Laplacian Spectrum**.

Let $A \in \{0,1\}^{N\times N}$ denote the adjacency matrix of graph G, and let $D = \text{diag}(d_1, d_2, \ldots, d_N)$ denote the diagonal degree matrix. The unnormalized **Graph Laplacian** is defined as:

$$\mathcal{L} = D - A \tag{4.1}$$

The eigenvalues of $\mathcal{L}$ are real and non-negative, ordered as:

$$0 = \lambda_1 \leq \lambda_2 \leq \lambda_3 \leq \cdots \leq \lambda_N \leq 2d_{max} \tag{4.2}$$

The multiplicity of eigenvalue 0 equals the number of connected components in the graph (G is connected $\iff \lambda_2 > 0$).

Fiedler's Algebraic Connectivity (λ_2): The second-smallest eigenvalue $\lambda_2(\mathcal{L})$ quantifies the intrinsic communication bottleneck of the network. By **Cheeger's Inequality**:

$$\frac{\lambda_2}{2} \leq h(G) \leq \sqrt{2d_{max} \cdot \lambda_2} \tag{4.3}$$

where $h(G) = \min_{S\subset V, |S|\leq N/2} \frac{|E(S,\bar{S})|}{|S|}$ is the Cheeger conductance (edge expansion).

A small λ_2 indicates a severe bisection bottleneck (such as a 1D ring or large 2D mesh, where $\lambda_2 = 2-2\cos(\pi/\sqrt{N}) \approx O(1/N)$). In contrast, expander graphs and high-radix topologies achieve $\lambda_2 \geq \Theta(1)$, ensuring that no sparse bisection cut can choke global traffic.

4.1.2 Classical Topological Performance Metrics

The choice of topology defines the structural limits of communication performance:

1. **Router Radix (k)**: The number of physical input/output ports per router:

$$k = \deg_{in}(u) = \deg_{out}(u) \tag{4.4}$$

Because internal crossbar switch area and power scale quadratically ($O(k^2)$), low-radix designs minimize local switch overhead, whereas high-radix switches reduce network traversal hops at the expense of larger crossbars.

2. **Network Diameter (D)**: The maximum shortest-path distance between any pair of nodes in the network:

$$D = \max_{u,v \in V} d(u, v) \tag{4.5}$$

where $d(u, v)$ is the minimum hop count between u and v. Diameter dictates worst-case communication latency.

3. **Average Hop Distance ($\bar{h}$)**: The expected shortest-path distance under a spatial traffic probability distribution $p(u, v)$:

$$\bar{h} = \sum_{u \in V} \sum_{v \neq u} p(u, v) \cdot d(u, v) \tag{4.6}$$

Under Uniform Random (UR) traffic ($p(u, v) = \frac{1}{N(N-1)}$), this reduces to $\bar{h}_{UR} = \frac{1}{N(N-1)} \sum_{u \neq v} d(u, v)$.

4. **Bisection Bandwidth (B_C)**: The minimum bandwidth cut partitioning the vertex set V into two equal disjoint subsets V_1, V_2 ($|V_1| = |V_2| = N/2$):

$$B_C = \min_{V_1, V_2} \sum_{u \in V_1, v \in V_2, (u,v) \in E} C(u, v) \tag{4.7}$$

where $C(u, v)$ denotes link bandwidth. In networks with uniform channel width W bits and clock frequency f_{clk}, bisection bandwidth is expressed as channel count N_C: $B_C = N_C \cdot W \cdot f_{clk}$.

5. **Maximum Channel Load (γ_{max}) and Saturation Throughput (Θ_{sat})**: Let γ_e denote the expected traffic volume traversing physical edge $e \in E$ when each node injects 1 flit per cycle. The maximum channel load is:

$$\gamma_{max} = \max_{e \in E} \gamma_e \tag{4.8}$$

Because each channel can transport at most 1 flit per cycle, the theoretical maximum saturation throughput per node is bounded by:

$$\Theta_{sat} \leq \frac{1}{\gamma_{max}} \leq \frac{2B_C}{N} \quad \text{[flits/node/cycle]} \tag{4.9}$$

6. **Zero-Load Latency (T_0)**: The average latency of an L-flit packet across uncontended routers and links:

$$T_0 = (\bar{h} + 1) \cdot t_{router} + \bar{h} \cdot t_{wire} + \left(\frac{L - 1}{\text{Rate}} \right) \tag{4.10}$$

4.2 Direct Topologies: Mesh, Torus, and the Folded Torus

Direct topologies couple every router directly to a compute node, forming regular planar grids well-suited for standard CMOS silicon tiling.

4.2.1 2D Mesh Networks: First-Principles Hop Derivations

In a $k \times k$ 2D Mesh ($N = k^2$):

- **Diameter**: Opposing corners require $2(k-1)$ hops: $D_{Mesh} = 2(k-1)$.
- **Combinatorial Average Hop Derivation**: Along a single 1D dimension of k nodes, the sum of pairwise Manhattan distances across all k^2 source-destination pairs (including self-traffic of distance 0) is:

$$\sum_{i=0}^{k-1}\sum_{j=0}^{k-1}|i-j| = 2\sum_{i=0}^{k-1}\sum_{j=0}^{i}(i-j) = 2\sum_{i=0}^{k-1}\frac{i(i+1)}{2} = \frac{(k-1)k(k+1)}{3} \tag{4.11}$$

  Dividing by k^2 gives the 1D average distance: $\bar{h}_{1D} = \frac{k^2-1}{3k} \approx \frac{k}{3}$. Because Manhattan distance is strictly additive across independent orthogonal dimensions, the 2D mesh average hop distance is:

$$\bar{h}_{Mesh} = 2 \cdot \bar{h}_{1D} = \frac{2(k^2-1)}{3k} \approx \frac{2}{3}k \tag{4.12}$$

  For an 8×8 mesh ($k = 8$), Eq. (4.12) yields $\bar{h} = \frac{2(63)}{24} = 5.25$ hops (or 5.33 hops when strictly excluding self-communication $u \neq v$).
- **Bisection Bandwidth**: Bisecting the mesh along the center cut slices k bidirectional links ($2k$ directional channels):

$$B_{C,Mesh} = 2k \cdot W \implies \Theta_{Mesh} \leq \frac{4W}{k} \tag{4.13}$$

Under uniform random traffic, peripheral nodes route through the center, generating severe center-channel congestion ($\gamma_{max} \approx k/4$).

4.2.2 Channel Load Under Adversarial Traffic Permutations

Under deterministic Dimension-Order (XY) routing, non-uniform traffic patterns induce extreme load imbalances:

- **Bit-Complement Traffic**: Node (x, y) transmits to $(k-1-x, k-1-y)$. Every packet must cross the center bisection in both dimensions, generating

maximum channel load $\gamma_{\max} = \frac{k}{2} C_{\text{link}}$ and causing throughput collapse to $\Theta_{\max} = \frac{2}{k}$ (50% of uniform capacity).

- **Matrix Transpose Traffic**: Node (x, y) transmits to (y, x). Packets above the diagonal route along $+X$ and then $-Y$, concentrating heavy traffic onto specific intermediate turning channels while opposite-direction channels remain completely idle.
- **Hotspot Traffic**: A fraction α_h of all injected packets from all cores target a specific memory controller tile (x_h, y_h). The local ingress channels of the hotspot router saturate at an injection rate of only:

$$\Theta_{\max,\text{hotspot}} \leq \frac{k_{\text{local}}}{\alpha_h \cdot N} \tag{4.14}$$

throttling the entire chip even if all other 99% of network channels are completely empty.

4.2.3 2D Torus and the Wire Delay Dilemma

A $k \times k$ 2D Torus (k-ary 2-cube) adds periodic wrap-around links. The topology becomes completely edge-symmetric:

$$D_{Torus} = 2 \left\lfloor \frac{k}{2} \right\rfloor, \quad \bar{h}_{Torus} \approx \frac{k}{2}, \quad B_{C,Torus} = 4k \cdot W \tag{4.15}$$

A 2D Torus delivers exactly twice the bisection bandwidth and a 25% reduction in average hop count compared to a 2D Mesh.

However, in a naive planar layout, the perimeter wrap-around wire must span the entire die:

$$L_{wrap} = (k-1) \cdot L_0 \tag{4.16}$$

where L_0 is the tile pitch. By the Elmore delay model ($\tau \propto L^2$), an 8×8 torus suffers 49× higher propagation delay on wrap-around wires than on local links, destroying clock frequency unless aggressively pipelined.

4.2.4 The Folded Torus Invariance

The **Folded Torus** topology solves the wrap-around wire crisis by physically interleaving even and odd nodes along each dimension:

$$\text{Physical Placement} : 0, \quad 2, \quad 4, \quad \ldots, \quad k-2, \quad k-1, \quad \ldots, \quad 3, \quad 1 \tag{4.17}$$

Every link in the Folded Torus—including all wrap-around connections—has an identical physical length:

$$L_{link,Folded} = 2 \cdot L_0 \tag{4.18}$$

By trading a constant wire length of $2L_0$ for the elimination of catastrophic $(k-1)L_0$ wires, the Folded Torus achieves high clock frequencies while retaining torus bisection bandwidth and low diameter.

4.3 Indirect Multistage Interconnects: Clos and Fat-Trees

Indirect topologies place processing elements exclusively at the network periphery, interconnecting them through intermediate switching stages.

4.3.1 Clos Networks and Non-Blocking Theorems

A 3-stage Clos network $C(n, m, r)$ comprises r input switches of size $n \times m$, m middle switches of size $r \times r$, and r output switches of size $m \times n$ [6].

- **Strictly Non-Blocking (SNB)**: Any idle input can connect to any idle output without rerouting existing calls if and only if:

$$m \geq 2n - 1 \tag{4.19}$$

- **Rearrangeably Non-Blocking (RNB)**: By the Slepian-Duguid theorem, any input can connect to any output with potential rerouting if:

$$m \geq n \tag{4.20}$$

4.3.2 Fat-Trees and Constant Bisection Bandwidth

Leiserson's Fat-Tree [34] avoids root bottlenecks by progressively increasing channel capacity toward the tree root. At each hierarchical level, the number of parallel channels doubles, preserving constant bisection bandwidth across all horizontal network partitions.

4.4 High-Radix Interconnects: Flattened Butterfly, Dragonfly, and Extremal Graphs

Modern deep-submicron and FinFET CMOS technologies provide vast routing gate density, prompting a fundamental architectural shift from low-radix/high-diameter networks (such as 2D meshes and tori) toward high-radix/low-diameter

architectures ($k \geq 16$–$64+$). In a high-radix network, router crossbar complexity increases quadratically ($O(k^2)$), but the network diameter D and average hop count $\bar{h}$ drop precipitously, drastically curtailing serialization delay, packet buffering overhead, and per-hop energy dissipation.

4.4.1 Flattened Butterfly Topologies

Kim et al. [31] introduced the **Flattened Butterfly** ($FBfly$) topology by projecting the multi-stage switches of a conventional k-ary n-fly butterfly into a single direct router per network tile. Formally, a d-dimensional Flattened Butterfly $FBfly(d, k)$ organizes $N = k^d$ nodes into a d-dimensional grid of radix-k routers.

In each dimension $i \in \{1, 2, \ldots, d\}$, a router at coordinate $(x_1, \ldots, x_i, \ldots, x_d)$ maintains dedicated point-to-point channels to all other $k-1$ routers along that coordinate axis:

$$(x_1, \ldots, x_i, \ldots, x_d) \longleftrightarrow (x_1, \ldots, x_i', \ldots, x_d) \quad \forall x_i' \neq x_i \tag{4.21}$$

Consequently, the network diameter and router radix are governed by:

$$D_{FBfly} = d, \quad k_{radix} = d(k-1) + p \tag{4.22}$$

where p denotes the local injection terminal density (number of compute cores or IP blocks attached per router). For a standard two-dimensional Flattened Butterfly ($d = 2$), any node in the network is reachable within at most $D = 2$ network hops (one row hop and one column hop), slashing zero-load router traversal cycles.

The bisection bandwidth of a balanced d-dimensional Flattened Butterfly bisecting along any single dimension cut is:

$$B_C = \frac{k^d \cdot W}{4} = \frac{N \cdot W}{4} \tag{4.23}$$

where W is the channel bit-width. Comparing Eq. (4.23) to a 2D mesh bisection bandwidth $B_{C,mesh} = 2\sqrt{N}W$, the Flattened Butterfly increases global bisection bandwidth by a factor of $\frac{\sqrt{N}}{8}$, completely eliminating the classical center-cut saturation bottleneck under adversarial traffic patterns.

However, dimension-order routing (DOR) on a Flattened Butterfly is vulnerable to worst-case adversarial traffic (e.g., bit-complement or transpose traffic), which concentrates flows onto a single intermediate dimension and saturates bisection channels. To overcome this, Kim et al. employ **Valiant's Randomized Routing (VAL)**: every packet is first forwarded to a uniformly random intermediate router (r_{rand}), and subsequently forwarded from (r_{rand}) to the destination router. While VAL doubles the average hop count from d to $2d$, it

converts any arbitrary worst-case traffic pattern into two independent uniform random distributions, guaranteeing 50% of peak theoretical throughput.

4.4.2 The Dragonfly Topology and Hierarchical Grouping

While the Flattened Butterfly scales efficiently across on-chip distances where dense intermediate metal layers are available, scaling beyond hundreds of nodes with direct point-to-point cables encounters severe physical packaging and cabling cost limits. Kim, Dally, Scott, and Abts [32] introduced the **Dragonfly** topology, which couples high router radix with hierarchical grouping to minimize expensive global channels.

A Dragonfly network is formally parameterized by four integer quantities:

- p: Number of terminal compute cores (or injection ports) attached to each router.
- a: Number of routers within each local group.
- h: Number of global inter-group channels hosted on each router.
- g: Number of identical groups constituting the entire system.

4.4.2.1 Hierarchical Connectivity and Radix Balance

Within each local group, the a routers are interconnected via a high-density, low-cost intra-group local network, typically organized as an intra-group clique (complete graph K_a) or a 2D flattened butterfly. In an intra-group clique, each router allocates $a - 1$ ports to local neighbors.

Globally, each group exposes a total of $a \cdot h$ global channels to the rest of the network. In the canonical balanced Dragonfly topology, every pair of groups is interconnected by exactly one global link:

$$g = a \cdot h + 1 \tag{4.24}$$

The total number of routers N_r and terminal processing nodes N in the balanced system are:

$$N_r = a \cdot g = a(ah + 1), \quad N = p \cdot N_r = a \cdot p \cdot (ah + 1) \tag{4.25}$$

The radix of each individual router is partitioned into terminal, local, and global ports:

$$k_{radix} = p + (a - 1) + h \tag{4.26}$$

To achieve cost optimality and balance global versus local throughput under uniform random traffic, the router port allocations satisfy the radix balance constraint:

$$a = 2p = 2h \tag{4.27}$$

Under this canonical balance, the total network diameter is bounded strictly by $D = 3$ hops:

1. **Local Hop 1**: From the source router R_s within group G_s to the local gateway router $R_g \in G_s$ hosting the global link to destination group G_d.
2. **Global Hop 2**: Traversal across the long global link from $R_g \in G_s$ to the corresponding ingress router $R_i \in G_d$.
3. **Local Hop 3**: From the ingress router R_i within group G_d to the ultimate destination router $R_d \in G_d$.

4.4.2.2 Universal Globally-Adaptive Routing (UGAL)

Because global channels in a Dragonfly represent scarce bisection resources, static minimal routing (MIN) suffers catastrophic throughput collapse (dropping to $\Theta \leq 1/a$ of peak capacity) under adversarial traffic that targets a specific global link.

To maintain near-ideal throughput across both uniform and adversarial traffic distributions, Kim et al. developed **Universal Globally-Adaptive Routing (UGAL)** [32]. Under UGAL, an injecting router dynamically selects between minimal routing (P_{MIN}, path length ≤ 3) and non-minimal Valiant routing (P_{VAL}, path length ≤ 6, routed through a randomly selected intermediate group G_r):

$$\text{Path Decision} = \begin{cases} P_{MIN}, & \text{if } w_{MIN} \cdot q_{MIN} \leq w_{VAL} \cdot q_{VAL} + \Delta_{bias} \\ P_{VAL}, & \text{otherwise} \end{cases} \tag{4.28}$$

where q_{MIN} denotes the instantaneous packet queue occupancy on the minimal global egress port, q_{VAL} is the queue occupancy toward the candidate intermediate group, and $w_{MIN} = 1$, $w_{VAL} = 2$ represent the path hop weights (since non-minimal routing consumes twice as much global link bandwidth). The threshold parameter $\Delta_{bias} \geq 0$ prevents premature non-minimal routing during transient queue spikes.

4.4.2.3 The Dragonfly+ Variant

In massive-scale systems where intra-group router count a exceeds 32, an intra-group clique becomes prohibitive due to $O(a^2)$ local link scaling. The **Dragonfly+** topology replaces the intra-group clique with a two-tier bipartite (spine-and-leaf) structure inside each group. Leaf routers connect to terminal nodes and intra-group spine switches; spine switches host the long global channels to remote groups. Dragonfly+ eliminates intra-group multi-hop routing and provides full cross-sectional bisection within each group while preserving the 3-hop global diameter.

4.4.3 Moore Bound, Slim Fly, and PolarFly Extremal Topologies

A fundamental open question in interconnect theory is: *Given a fixed router radix k, what is the theoretical maximum number of compute nodes N that can be interconnected with a maximum diameter D?*

4.4.3.1 The Theoretical Moore Bound

By considering a breadth-first search tree originating from an arbitrary vertex in a regular graph of degree k, the number of reachable vertices at hop distance i is at most $k(k-1)^{i-1}$. Summing over all distances up to diameter D yields the **Moore Bound**:

$$N(k, D) \leq 1 + k \sum_{i=0}^{D-1} (k-1)^i = 1 + k \frac{(k-1)^D - 1}{k-2} \tag{4.29}$$

For diameter $D = 2$, Eq. (4.29) simplifies to the classical Hoffman-Singleton bound:

$$N(k, 2) \leq k^2 + 1 \tag{4.30}$$

Graphs achieving the exact Moore bound for $D = 2$ exist only for degrees $k \in \{2, 3, 7\}$ (and potentially $k = 57$, which remains an unsolved mathematical conjecture). Consequently, real-world interconnection networks seek **extremal graph constructions** that approach the Moore bound as closely as possible.

4.4.3.2 The Slim Fly Topology

Besta and Hoefler [3] introduced **Slim Fly**, a high-radix network architecture based on **McKay-Miller-Siran (MMS) graphs**. An MMS graph is constructed over the Galois field $\mathbb{F}_q$, where q is a prime power congruent to $4m+1$ ($m \in \mathbb{N}$). The vertices are partitioned into two disjoint subsets of size q^2:

$$V = \{(0, x, y) \mid x, y \in \mathbb{F}_q\} \cup \{(1, x', y') \mid x', y' \in \mathbb{F}_q\}, \quad |V| = 2q^2 \tag{4.31}$$

Links are established according to quadratic residues within $\mathbb{F}_q$:

$$(0, x, y) \sim (0, x, y') \iff y - y' \in R_\xi \tag{4.32}$$

$$(1, x', y') \sim (1, x', y'') \iff y' - y'' \in \xi R_\xi \tag{4.33}$$

$$(0, x, y) \sim (1, x', y') \iff y - y' = x \cdot x' \tag{4.34}$$

where R_ξ is the multiplicative subgroup of quadratic residues.

The resulting Slim Fly topology achieves:

$$D_{SF} = 2, \quad k_{radix} = \frac{3q - \delta}{2} \approx \frac{3}{2}\sqrt{\frac{N}{2}} \tag{4.35}$$

where $\delta \in \{1, 3\}$. Slim Fly reaches **nearly 50% of the theoretical Moore bound for diameter 2**, outperforming both the Flattened Butterfly and Dragonfly by reducing total router count by 25% and network cabling costs by over 30% at identical bisection capacity.

4.4.3.3 PolarFly: Non-Bipartite Polarity Graphs

To eliminate the bipartite constraint of Slim Fly and enable arbitrary node cardinalities, modern research has synthesized **PolarFly**, built upon **Erdős-Rényi polarity graphs** from projective geometry $PG(2, q)$. In PolarFly, every router corresponds to a 1-dimensional subspace (point) in $\mathbb{F}_q^3$, and links correspond to orthogonal points with respect to a non-degenerate symmetric bilinear form (quadric). PolarFly achieves diameter 2 with zero quadrangle cycles, yielding exceptional resilience against multi-link failures and unprecedented algebraic connectivity λ_2.

4.5 Physical Packaging Hierarchy and Multi-Die Interconnects

The topology of an interconnect cannot be designed in isolation from the physical packaging hierarchy. Modern high-performance SoCs span multiple packaging domains, each exhibiting orders-of-magnitude differences in wire pitch, capacitance, signal attenuation, and pin count.

Table 4.1 Physical Interconnect Metrics Across Packaging Hierarchies

Hierarchy Level	**Wire Pitch (μm)**	**Capacitance (fF/μm)**	**Latency**	**Energy (pJ/bit)**
On-Chip Metal (M1–M6)	0.04 – 0.15	0.15 – 0.25	$<$ 100 ps/mm	0.01 – 0.05
On-Chip Thick Metal (M7–M12)	0.50 – 2.00	0.10 – 0.18	60 ps/mm	0.05 – 0.15
Silicon Interposer (2.5D CoWoS)	0.80 – 1.50	0.35 – 0.60	120 ps/mm	0.10 – 0.30
High-Density Bridge (EMIB)	1.00 – 2.00	0.30 – 0.50	110 ps/mm	0.15 – 0.40
Organic Package Substrate	25.0 – 50.0	1.20 – 2.00	180 ps/mm	0.80 – 2.50
PCB Traces / Backplane	100.0 – 250.0	2.50 – 4.50	220 ps/mm	3.00 – 10.00
Active Optical Cables (AOC)	N/A (fiber)	N/A (photons)	5000 ps/m	10.00 – 25.00

As shown in Table 4.1, crossing an organic package substrate boundary incurs nearly 20× higher energy per bit than routing across local thick-metal on-chip layers. This physical discontinuity mandates **hierarchical topology mapping**:

- **Intra-Die Domain**: High-radix, dense meshes or cascaded crossbars exploiting dense on-chip wiring pitches ($\leq 0.5\,\mu$m).
- **Die-to-Die (D2D) Domain**: Ultra-short reach (USR) parallel buses routed across silicon bridges (e.g., UCIe standard, EMIB, CoWoS), preserving terabyte-per-second beachhead bandwidth without serialization latency.
- **Inter-Socket / System Domain**: Serialized SerDes links driving differential copper traces or optical transceivers organized in low-diameter Dragonfly or Slim Fly topologies.

4.6 Industrial High-Radix Interconnect Case Studies

To bridge theoretical topology formulations with industrial silicon realization, we examine two modern high-radix switching platforms: the NVIDIA NVSwitch fabric and the Cray Slingshot interconnect.

4.6.1 Case Study 1: NVIDIA NVSwitch Fabric

In massive distributed deep learning workloads, multi-GPU communication patterns are dominated by collective primitives such as `AllReduce`, `AllGather`, and `ReduceScatter`. Traditional PCI Express topologies create severe bisection bottlenecks, limiting scaling efficiency.

To eliminate this bottleneck, NVIDIA developed the **NVSwitch** architecture. An HGX baseboard integrates multiple dedicated high-radix crossbar ASICs:

- **NVSwitch 1 (Volta)**: 18-port crossbar router delivering 900 GB/s bidirectional bandwidth per switch across 16 V100 GPUs.
- **NVSwitch 2 (Ampere)**: 36-port crossbar router in 7 nm TSMC CMOS, providing 7.2 Tb/s aggregate switching bandwidth with hardware-accelerated SHARP (Scalable Hierarchical Aggregation and Reduction Protocol) reduction engines embedded directly inside the crossbar datapaths.
- **NVSwitch 3 (Hopper / Blackwell)**: 64-port physical crossbar delivering 14.4 Tb/s (Hopper) and 28.8 Tb/s (Blackwell) aggregate non-blocking bandwidth per chip.

By deploying a multi-tier fat-tree of NVSwitch chips, an NVLink network creates a unified, single-address-space distributed crossbar connecting up to 72 GPUs in a single NVL72 rack without intermediate PCIe or Ethernet hops, providing 1.8 TB/s of all-to-all bidirectional bandwidth per GPU.

4.6.2 Case Study 2: The Cray Slingshot Interconnect

The **Cray Slingshot** architecture [11] represents a modern commercial implementation of a high-radix Dragonfly network, deployed in leadership supercomputers including Frontier and Aurora.

The Slingshot router is an ASIC fabricated in 16 nm FinFET comprising **64 ports of 200 Gb/s** each (delivering 12.8 Tb/s of bidirectional switching bandwidth per router):

1. **Dragonfly Group Organization**: The 64 ports of each router are partitioned into 16 compute-node injection ports, 32 local intra-group ports, and 16 global optical inter-group ports. A single Slingshot Dragonfly group contains up to 32 routers (interconnecting 512 compute nodes), and up to 256 groups are interconnected into an exascale fabric supporting 131,072 endpoints with a maximum diameter of $D = 3$.
2. **Hardware-Enforced Fine-Grained Adaptive Routing**: Rather than relying on static hashing (such as ECMP in datacenter Ethernet), Slingshot performs per-packet dynamic route selection on every router. When global channels become congested, packets are automatically rerouted across non-minimal intermediate groups.
3. **Credit-Driven Congestion Isolation**: Slingshot monitors queue residency across all virtual channels in real time. When a congested flow accumulates excessive buffer occupancy, the router throttles ingress injection credits specifically for that offending flow without head-of-line blocking innocent, non-congested flows.

4.7 Application-Specific Synthesis & Cascaded Crossbars

Heterogeneous SoCs exhibit highly non-uniform communication demands. Rigid standard topologies introduce severe silicon area and power waste.

4.7.1 MIRO Synthesis of Cascaded Crossbar Networks

Jun, Yoo, and Chung [27] formulated the synthesis of cascaded crossbar switches to replace power-hungry centralized bus matrices.

Jun et al. proved that while Mixed Integer Linear Programming (MILP) discovers exact optimal topologies, its computational complexity explodes for $|V| > 16$. The MIRO heuristic clusters masters based on communication affinities, achieving **up to 69.5% area reduction** over conventional architectures while satisfying all clock frequency requirements.

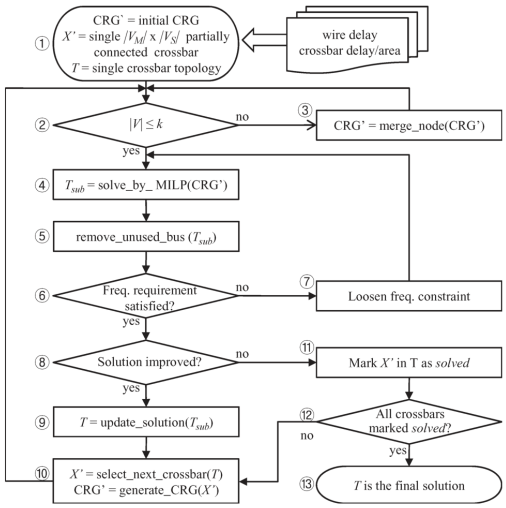


Fig. 1. MIRO Process.

Fig. 4.1 The MIRO (Multi-stage Interconnect Reduction and Optimization) process for synthesizing cascaded crossbar networks under bandwidth and latency constraints. [Reproduced from: M. Jun, S. Yoo, and E.-Y. Chung, "Topology Synthesis of Cascaded Crossbar Switches," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, Vol. 28, No. 6, pp. 926–930, June 2009, DOI: 10.1109/TCAD.2009.2017079. © 2009 IEEE]

4.7.2 Switch Merger and Partitioning Algorithms

Dumitriu and Khan [14] introduced transaction-oriented topology generation using iterative switch merging and partitioning:

By using predictive contention modeling during synthesis, the resulting topologies matched or exceeded regular 2D mesh throughput while utilizing **50% fewer network resources**.

4.8 Silicon Floorplanning Limits: The Passas Radix-128 Proof

A persistent misconception in early literature asserted that full crossbar networks are unviable beyond 32 nodes. Passas, Katevenis, and Pnevmatikatos [43] refuted this by modeling and fabricating a **radix-128 crossbar in 90 nm CMOS**.

Passas et al. proved that:

1. With 8+ metal layers, crossbar wiring tracks fit directly **over standard cells** on intermediate metal layers without expanding standard-cell silicon area:

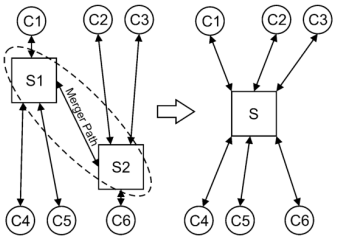


Fig. 3. Example of a switch merger.

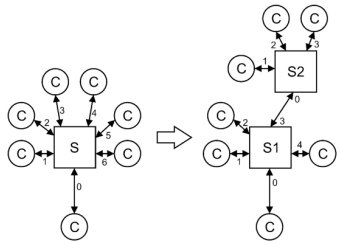


Fig. 4. Example of the switch partitioning process.

Fig. 4.2 Application-specific topology optimization: (left) switch merger process combining low-traffic routers; (right) switch partitioning process dividing congested routers to meet physical pin and timing constraints. [Reproduced from: V. Dumitriu and G. N. Khan, "Throughput-Oriented NoC Topology Generation and Analysis for High Performance SoCs," *IEEE Trans. Very Large Scale Integr. (VLSI) Syst.*, Vol. 17, No. 10, pp. 1433–1446, October 2009, DOI: 10.1109/TVLSI.2008.2006798. © 2009 IEEE]

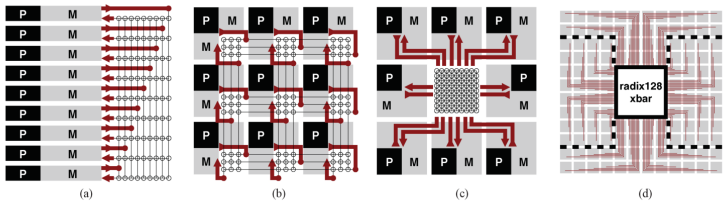


Fig. 1. Die floorplans. (a) Stack of user tiles using side crossbar. (b) Matrix of user tiles using distributed crossbar. (c) Matrix of user tiles using centralized crossbar. (d) Balanced link density of (c) in a "12×12 minus 4×4" matrix of user tiles; *P*s denote processors and *M*s memories.

Fig. 4.3 Die floorplans for crossbar-based on-chip networks. (a) Stack of user tiles with side crossbar. (b) Matrix of user tiles with distributed crossbar. (c) Centralized crossbar. (d) Balanced link density layout in a 12×12 minus 4×4 matrix of tiles. [Reproduced from: G. Passas, M. Katevenis, and D. Pnevmatikatos, "Crossbar NoCs Are Scalable Beyond 100 Nodes," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, Vol. 31, No. 4, pp. 573–585, April 2012, DOI: 10.1109/TCAD.2011.2176730. © 2012 IEEE]

$$\text{Area}_{crossbar} \approx O(N^2 \cdot W), \quad \text{Delay}_{crossbar} \approx O(\sqrt{N \cdot W}) \tag{4.36}$$

2. In a 150 mm^2 die containing 128 1 mm^2 processor tiles, the radix-128 crossbar occupied **only 11% of total die area** (including repeaters and pipeline flip-flops).
3. Under global and stressed traffic, the monolithic crossbar delivered **up to 3× higher throughput and 2.5× lower latency** than a 2D mesh by eliminating multi-hop queueing stalls.

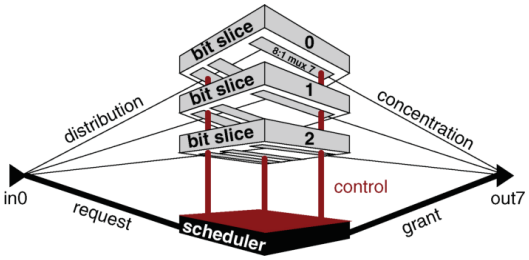


Fig. 3. 3-bit radix-8 crossbar.

Fig. 4.4 Standard-cell layout of a 3-bit radix-8 crossbar bit-slice. Multiplexer circuits are routed entirely beneath upper-level metal tracks. [Reproduced from: G. Passas, M. Katevenis, and D. Pnevmatikatos, "Crossbar NoCs Are Scalable Beyond 100 Nodes," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, Vol. 31, No. 4, pp. 573–585, April 2012, DOI: 10.1109/TCAD.2011.2176730. © 2012 IEEE]

4.9 Quantitative Topology Evaluation Matrix

Table 4.2 summarizes the mathematical and microarchitectural metrics across major topology classes for an $N = 64$ core system.

Table 4.2 Analytical and Empirical Metric Matrix across Interconnect Topologies ($N = 64$ nodes)

Topology	Radix (k)	Diam (D)	Avg Hops ($\bar{h}$)	Bisect (B_C/W)	T_0 (cyc)	Sat. Thru.
Ring (1D)	2	32	16.00	4	28.00	0.125
2D Mesh (8×8)	4 (avg 3.5)	14	5.33	16	12.00	0.500
2D Torus (8×8)	4	8	4.06	32	10.10	1.000
Folded Torus (8×8)	4	8	4.06	32	10.10	1.000
Hypercube (6-cube)	6	6	3.05	64	8.57	2.000
Fat-Tree (3-level)	4	6	3.80	64	9.70	2.000
Flattened Butterfly (8×8)	14	2	1.78	256	6.67	8.000
Full Crossbar (64×64)	64	1	1.00	2048	5.00	64.000

4.10 Interactive Lab & Hands-on Challenges

Explore the interactive topology visualizers on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **Topology Titans: Graph Architect** (https://network-on-chip.github.io/learning/interactive/topology/): Synthesize 2D Mesh, 2D Torus,

and Hypercube topologies within strict wire length and pin degree constraints.

- **2D Mesh Routing Visualizer** (https://network-on-chip.github.io/learning/interactive/visualizer/): Inject uniform random, bit-complement, and hotspot traffic into a 4×4 and 8×8 mesh to observe bisection link saturation in real time.
- **Arbiter Arena: Crossbar Clash** (https://network-on-chip.github.io/learning/interactive/arbiter/): Benchmark matrix, round-robin, and wavefront arbiters under high-radix switch allocation.

4.11 Exercises and Review Questions

1. **Bisection Bandwidth Bound**: Prove that under uniform random traffic, the maximum throughput per node is bounded by $\Theta \leq \frac{2B_C}{N}$. What fraction of injected traffic must traverse the bisection cut?
2. **Folded Torus Delay**: For an 8×8 tiled chip with tile pitch $L_0 = 1.2$ mm, compare the wire delay of a naive torus wrap-around link versus a folded torus link under the Elmore delay model.
3. **Clos Non-Blocking Sizing**: For a 3-stage Clos network $C(n, m, r)$ with $N = 64$ and $n = 8$, calculate the minimum m required for strictly non-blocking operation versus rearrangeably non-blocking operation.
4. **High-Radix Trade-Offs**: Explain why high-radix networks become more energy-efficient than low-radix 2D meshes as core counts scale beyond 100 nodes.

Chapter 5
Deadlock Theory, Turn Models & Adaptive Routing

Abstract In packet-switched Networks-on-Chip (NoCs), deadlock represents an unrecoverable failure state in which a set of packets are mutually blocked indefinitely, freezing communication across the silicon die. While Store-and-Forward networks localize buffer dependencies to adjacent nodes, Wormhole Switching dramatically inflates the resource dependency space because packets span multiple physical links simultaneously. This chapter provides a rigorous theoretical and microarchitectural treatment of routing deadlock freedom. We analyze the Dally-Seitz Channel Dependency Graph (CDG) theorem, establish the necessary and sufficient conditions for cycle-free routing, and examine the Glass-Ni Turn Models that deliver partially adaptive routing without virtual channels. We formulate Duato's theorem proving that adaptive subnetworks may contain cycles if an acyclic escape subnetwork is present, and examine the Odd-Even turn model. Finally, we explore state-of-the-art IEEE TCAD and TCAS-II 2026 breakthroughs, addressing protocol-level deadlock in cache-coherent manycores via the PCMT framework and 3D multi-chiplet deadlock avoidance via ScaDA-Rou.

5.1 The Deadlock Problem in Wormhole-Switched Fabrics

In a packet-switched interconnection network, **deadlock** occurs when a set of concurrent packets cannot make forward progress because each packet holds one or more buffer resources while awaiting buffers held by other packets in the set.

In **Wormhole Switching**, buffer allocation operates at the flit granularity while routing and channel acquisition occur at the packet level. Consequently, a single packet of length L flits spans across multiple physical links and intermediate routers simultaneously:

- The **Head flit** arbitrates for downstream virtual channels and crossbar ports.
- The **Body flits** trail behind in intermediate router FIFO buffers.
- The **Tail flit** releases channel reservations upon traversing the link.

When a Head flit encounters buffer backpressure at a downstream router, the body flits stall in place, holding buffers across the upstream routers. This simultaneous *holding* and *requesting* of distributed resources creates cyclical resource dependency loops.

Architects distinguish deadlock from two related communication anomalies:

1. **Deadlock**: Permanent cessation of forward progress caused by circular resource dependencies.
2. **Livelock**: Packets remain in continuous motion across the network but never arrive at their destinations due to unconstrained non-minimal deflection.
3. **Starvation**: A packet is permanently denied switch arbitration access due to unfair allocation policies, despite the network remaining active.

5.2 The Dally-Seitz Channel Dependency Graph Theorem

The foundational mathematical framework for proving deadlock freedom in packet networks was developed by William J. Dally and Charles L. Seitz [7].

5.2.1 Formal Definition of the CDG

Let $I = (V, C)$ be an interconnection network, where V is the set of router nodes and C is the set of directed physical and virtual channels. Let $R : C \times V \to \mathcal{P}(C)$ be the routing function determining permissible egress channels from an ingress channel $c_i \in C$ toward destination $v \in V$.

The **Channel Dependency Graph (CDG)** $D = (C, E_{CDG})$ is a directed graph where:

- Vertices correspond to channels $c \in C$.
- A directed edge $(c_i, c_j) \in E_{CDG}$ exists if and only if a packet can transition directly from channel c_i to channel c_j under routing function R:

$$(c_i, c_j) \in E_{CDG} \iff \exists v \in V \quad \text{such that} \quad c_j \in R(c_i, v) \tag{5.1}$$

5.2.2 The Theorem

Theorem 5.1 (Dally & Seitz, 1987) *A deterministic routing function R for an interconnection network I is deadlock-free if and only if its Channel Dependency Graph $D = (C, E_{CDG})$ has no directed cycles.*

Proof. **Sufficiency ($\Leftarrow$)**: If D has no directed cycles, it is a finite Directed Acyclic Graph (DAG). There exists at least one channel $c_{sink} \in C$ with zero

out-degree (delivering to local core memory). A packet occupying c_{sink} cannot be blocked by any other packet in the fabric and will clear in finite time. By topological sorting of the DAG, every channel can be assigned an integer rank $\text{rank}(c_i)$ such that $(c_i, c_j) \in E_{CDG} \implies \text{rank}(c_i) < \text{rank}(c_j)$. Flits on channels of rank K are guaranteed forward progress once channels of rank $> K$ clear. Hence, deadlock is impossible.

Necessity ($\Rightarrow$): If D contains a directed cycle $c_0 \to c_1 \to \cdots \to c_{M-1} \to c_0$, one can construct a traffic state with M packets $P_0, \ldots, P_{M-1}$ such that P_i holds channel c_i and requests channel $c_{(i+1) \bmod M}$. If all channels have finite buffer capacity and are full, no packet can advance, producing an unrecoverable deadlock. □

5.2.3 Application: Dimension-Order Routing (XY Routing)

In a 2D Mesh, deterministic XY routing requires packets to exhaust all displacement in the X-dimension before turning onto the Y-dimension.

- **Allowed Turns**: $(X \to X), (X \to Y), (Y \to Y)$.
- **Prohibited Turns**: $(Y \to X)$ (no turns from Y-channels onto X-channels).

Because X-channels never depend on Y-channels, all edges in the CDG flow strictly from X to Y. The CDG is strictly acyclic, proving that **XY routing is provably deadlock-free on any 2D mesh without requiring virtual channels**.

5.2.4 Dateline Routing in k-ary n-cube Torus Networks

While dimension-order routing guarantees cycle freedom in a mesh, it fails catastrophically in a **Torus** (k-ary n-cube). In a 1D ring of k nodes, packets moving in the positive X direction form a circular dependency loop of length k:

$$c_0 \longrightarrow c_1 \longrightarrow c_2 \longrightarrow \cdots \longrightarrow c_{k-1} \longrightarrow c_0 \tag{5.2}$$

Because the physical topology contains cycles, no purely topological turn prohibition can break the circular dependency without partitioning the network into a line (sacrificing the wrap-around link).

To solve this without sacrificing wrap-around bandwidth, Dally and Seitz [7] introduced **Virtual Channel Dateline Routing**:

1. Each physical channel along dimension X is split into at least two Virtual Channels: VC_0 (low channel) and VC_1 (high channel).
2. An arbitrary physical link is designated as the **Dateline** (typically the wrap-around link connecting router $k-1$ to router 0).
3. **Virtual Channel Transition Invariant**:

- Packets inject into the network on VC_0.
- As long as a packet traverses links prior to the dateline, it remains in VC_0.
- Upon crossing the dateline, the packet is forced to transition into VC_1.
- A packet in VC_1 is strictly forbidden from ever returning to VC_0.

Theorem 5.2 (Torus Deadlock Freedom via Dateline Routing) *Dateline routing on a k-ary 1-cube (ring) using 2 virtual channels per physical link yields a strictly acyclic Channel Dependency Graph.*

Proof. Let the channels be ordered lexicographically by pair (VC, x), where $VC \in \{0, 1\}$ and $x \in \{0, 1, \ldots, k-1\}$ represents the channel originating at node x:

$$(0,0) < (0,1) < \cdots < (0,k-2) < (1,0) < (1,1) < \cdots < (1,k-1) \tag{5.3}$$

Traversals within VC_0 step from $(0, x)$ to $(0, x+1)$, strictly increasing channel rank. Crossing the dateline transitions from $(0, k-1)$ to $(1, 0)$, strictly increasing channel rank. Traversals within VC_1 step from $(1, x)$ to $(1, x+1)$, strictly increasing channel rank. Because every routing step strictly increases channel rank, the CDG is a directed acyclic graph (DAG), guaranteeing deadlock freedom. □

For a d-dimensional torus (k-ary d-cube), applying dimension-order routing (e.g., $X \to Y \to Z$) combined with independent datelines per dimension guarantees deadlock freedom with 2^d (or $2 \times d$) virtual channels.

5.2.5 Matrix Formulation and Algorithmic Cycle Detection in CDGs

For complex, irregular, or dynamically reconfigurable NoCs, manual cycle analysis is infeasible. The CDG can be formalized algebraically via an **Adjacency Dependency Matrix** $M_{CDG} \in \{0, 1\}^{|C| \times |C|}$, where:

$$M_{CDG}[i, j] = \begin{cases} 1, & \text{if } (c_i, c_j) \in E_{CDG} \\ 0, & \text{otherwise} \end{cases} \tag{5.4}$$

The (i, j)-th entry of the m-th matrix power, $(M^m_{CDG})[i, j]$, represents the number of directed dependency paths of exact length m from channel c_i to channel c_j. Consequently:

$$\text{CDG is Acyclic} \iff \text{Tr}(M^m_{CDG}) = 0 \quad \forall m \in \{1, 2, \ldots, |C|\} \tag{5.5}$$

which is equivalent to stating that M_{CDG} is a nilpotent matrix whose eigenvalues are all identically zero ($\lambda_i(M_{CDG}) = 0 \; \forall i$).

In automated CAD routing synthesizers, deadlock verification is executed in $O(|C| + |E_{CDG}|)$ time using **Tarjan's Strongly Connected Components (SCC) algorithm**. If every strongly connected component of the CDG consists of a single channel with no self-loop, the routing algorithm is provably deadlock-free.

5.3 Glass & Ni Turn Models for Adaptive Routing

While XY routing is deadlock-free, it is strictly non-adaptive: packets cannot route around congested or faulty routers. Christopher J. Glass and Lionel M. Ni [22] introduced the **Turn Model**, enabling partially adaptive deadlock-free routing without adding virtual channels.

5.3.1 Turn Analysis in a 2D Mesh

In a 2D Mesh, flits travel in 4 directions: East, West, North, South. There are exactly $4 \times 2 = 8$ possible 90-degree turns:

$$\{E \to N, \quad E \to S, \quad W \to N, \quad W \to S, \quad N \to E, \quad N \to W, \quad S \to E, \quad S \to W\}$$

These 8 turns form two simple directed cycles:

1. **Clockwise Cycle**: $E \to S \to W \to N \to E$ (uses $E \to S, S \to W, W \to N, N \to E$).
2. **Counter-Clockwise Cycle**: $E \to N \to W \to S \to E$ (uses $E \to N, N \to W, W \to S, S \to E$).

Theorem 5.3 (Glass & Ni, 1994) *To prevent deadlock in a mesh network, it is necessary and sufficient to prohibit at least one turn from every simple cycle. To maximize routing adaptiveness, exactly one turn from each directed cycle should be prohibited.*

By prohibiting exactly **2 out of the 8 turns** (one clockwise, one counter-clockwise), the network achieves **75% (6/8) turn freedom** while completely breaking all dependency cycles!

5.3.2 The Canonical 2D Turn Models

1. **West-First**:
   - Prohibited turns: $N \to W$ and $S \to W$.

- Rule: If a packet must travel West ($\Delta X < 0$), it must execute all Westward hops first before making any North or South turns.
- Adaptiveness: East-bound packets ($\Delta X > 0$) enjoy 100% full adaptiveness between East, North, and South channels.

2. **North-Last**:
   - Prohibited turns: $N \to E$ and $N \to W$.
   - Rule: Northward hops must be postponed until all East, West, and South hops are complete. Once a packet turns North, it can never turn again.
3. **Negative-First**:
   - Prohibited turns: Transitions from positive directions ($+X, +Y$: East, North) to negative directions ($-X, -Y$: West, South). Specifically, $E \to W$ (180°, illegal anyway), $E \to S$, $N \to W$, and $N \to S$.
   - Rule: Flits must complete all negative dimension hops before acquiring any positive channels.

5.3.3 3D Turn Models and Combinatorial Complexity

As semiconductor packaging transitions to true 3D integration with Through-Silicon Vias (TSVs), routing must extend to 3D meshes (X, Y, Z). In a 3D mesh, there are 6 cardinal directions ($\pm X, \pm Y, \pm Z$).

At each router, a flit arriving from any of the 6 directions can turn onto any of the 4 orthogonal directions, resulting in:

$$N_{turns,3D} = 6 \times 4 = 24 \text{ possible 90-degree turns} \tag{5.6}$$

These 24 turns project onto the three orthogonal coordinate planes:

- The XY plane: 8 turns, forming 2 planar cycles (Clockwise and Counter-Clockwise).
- The XZ plane: 8 turns, forming 2 planar cycles.
- The YZ plane: 8 turns, forming 2 planar cycles.

In total, there are $2 \times 3 = 6$ simple planar cycles, in addition to non-planar 3D cycles (e.g., $+X \to +Y \to +Z \to -X \to -Y \to -Z$).

To guarantee cycle freedom across all 3 dimensions without virtual channels:

- Under **3D Negative-First Routing**, all turns from $\{+X, +Y, +Z\}$ to $\{-X, -Y, -Z\}$ are prohibited (a total of $3 \times 3 = 9$ prohibited turns out of 24, leaving 15 permitted turns or 62.5% adaptiveness).
- Packets must route all $-X, -Y, -Z$ displacements first, and only then route along $+X, +Y, +Z$. This completely prevents both planar and non-planar circular dependencies.

5.4 The Odd-Even Turn Model

To eliminate the directional asymmetry inherent in Glass-Ni turn models (where West-bound packets suffer 0% adaptiveness while East-bound enjoy 100%), Ge-Ming Chiu [5] introduced the **Odd-Even Turn Model**. Rather than enforcing uniform turn rules globally, Odd-Even routing modulates turn restrictions according to the **column parity** (x mod 2) of the current router:

Definition 5.1 (Odd-Even Turn Rules)

- **Rule 1 (Even Columns, x is Even)**: A packet is not permitted to take an $E \to N$ or $E \to S$ turn at any node located in an even column. Furthermore, $N \to W$ and $S \to W$ turns are prohibited at even columns.
- **Rule 2 (Odd Columns, x is Odd)**: A packet is not permitted to take an $N \to E$ or $S \to E$ turn at any node located in an odd column. Furthermore, $W \to N$ and $W \to S$ turns are prohibited at odd columns.

Theorem 5.4 (Deadlock Freedom of Odd-Even Routing) *The Odd-Even turn model prevents all circular dependencies in a 2D mesh without requiring virtual channels.*

Proof. A routing cycle in a 2D mesh requires at least one right-to-left (Westward) path segment and at least one left-to-right (Eastward) path segment connected by vertical (North/South) transitions.

1. In any potential Clockwise cycle ($E \to S \to W \to N \to E$), the turn from E to vertical must occur at an odd column (since $E \to S$ is banned at even columns). The subsequent turn from W to N must occur at an even column (since $W \to N$ is banned at odd columns).
2. However, for a packet to transition from an odd column to an even column while traveling West, it must step from x_{odd} to $x_{odd} - 1 = x_{even}$. To return Eastward, it must step from x_{even} to $x_{even} + 1 = x_{odd}$.
3. Chiu proved that the alternating sequence of permitted turns forces any hypothetical closed path to violate the column parity constraints, making the formation of a directed cycle topologically impossible.

□

Odd-Even routing distributes adaptiveness evenly across the fabric: both East-bound and West-bound traffic retain substantial routing flexibility, avoiding the severe thermal and congestion hotspots produced by West-First routing on the western die periphery.

5.5 Duato's Theorem for Fully Adaptive Routing

In 1993, José Duato [13] established a foundational breakthrough: **An adaptive routing function does NOT require an acyclic Channel Dependency Graph to guarantee deadlock freedom**.

5.5.1 Formal Mathematical Statement

Let $I = (V, C)$ be an interconnection network and $R : C \times V \to \mathcal{P}(C)$ be an adaptive routing function.

Theorem 5.5 (Duato, 1993) *An adaptive routing function R is deadlock-free if the channel set C can be partitioned into an adaptive channel subset C_A and an escape channel subset C_E ($C = C_A \cup C_E$, $C_A \cap C_E = \emptyset$) satisfying three conditions:*

1. ***Escape Subnetwork Connectedness**: The subnetwork $I_E = (V, C_E)$ provides at least one connected path between every source $s \in V$ and destination $d \in V$ under an escape routing function $R_E \subseteq R$.*
2. ***Escape Acyclicity**: The Channel Dependency Graph of the escape subnetwork, $D_E = (C_E, E_{CDG,E})$, contains no directed cycles.*
3. ***Unconditional Escape Transition**: A packet occupying any channel in C_A is always permitted to transition into an eligible channel in C_E.*

5.5.2 Proof of Duato's Theorem

Proof. Suppose, for contradiction, that a deadlock configuration exists. A deadlocked configuration consists of a non-empty set of packets $\mathcal{S}_{dead}$ where every packet $P \in \mathcal{S}_{dead}$ is permanently blocked awaiting buffer space on channels held by other packets in $\mathcal{S}_{dead}$.

1. No packet in $\mathcal{S}_{dead}$ can be occupying solely channels in C_E with its requested channels residing in C_E, because by Condition 2, D_E is strictly acyclic (a finite DAG). The packet occupying the topologically highest channel in C_E has an unobstructed path to its destination and must advance, contradicting the deadlock assumption.
2. Therefore, any circular wait must involve packets occupying channels in C_A.
3. However, by Condition 3, every packet in C_A can request escape channels in C_E. By Condition 1, C_E connects all nodes.
4. Since C_E is acyclic and possesses finite buffer capacity, packets entering C_E are guaranteed to drain in finite time. Once an escape channel buffer becomes vacant, a stalled packet in C_A acquires it, breaking the stall.
5. Hence, an unrecoverable cyclic dependency cannot persist, and the system is provably deadlock-free.

□

5.5.3 Microarchitectural Realization: Adaptive and Escape VCs

In modern industrial routers, Duato's theorem is realized using Virtual Channels:

- **Virtual Channel 0 (VC_0)**: Dedicated exclusively as the **Escape Subnetwork (C_E)**. Routing on VC_0 is restricted to deterministic Dimension-Order (XY) routing.
- **Virtual Channels 1 through $V-1$ ($VC_1 \dots VC_{V-1}$)**: Constitute the **Adaptive Subnetwork (C_A)**. Routing in C_A is fully adaptive (minimal or non-minimal), permitting arbitrary turns and cycles in the CDG.

During switch allocation, incoming packets prioritize available buffers in C_A to route around congested links. If all downstream C_A virtual channels are congested or stalled, the packet requests allocation on the escape channel VC_0. Once on VC_0, the flit advances deterministically, guaranteeing drainage and completely preventing deadlock collapse.

5.6 Protocol-Level Deadlock and the PCMT Architecture

In shared-memory multi-core processors, routing-level deadlock freedom is insufficient to guarantee system forward progress. Cache coherence protocols introduce high-level message dependencies that can cause **Protocol-Level Deadlock**.

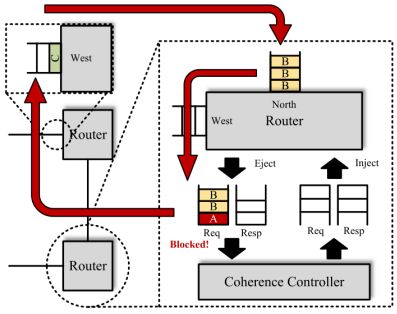


Fig. 1. Possible cyclic dependency of protocol-level deadlock in NoC.

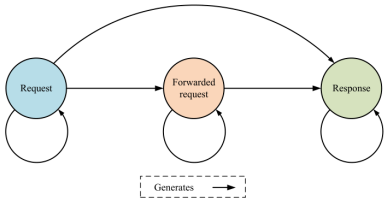


Fig. 4. Coherence message types in directory-based systems.

Fig. 5.1 Protocol-level deadlock in cache-coherent multi-cores: (left) circular dependency loop between coherence Request and Response message buffers; (right) hierarchical message classification in directory-based systems. [Reproduced from: Y. Jia, Z. Zhang, C. Fu, L. Wan, and J. Han, "PCMT: Prioritizing Coherence Message Types for NoC Protocol-Level Deadlock Freedom," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, Vol. 45, No. 4, pp. 1949–1962, April 2026, DOI: 10.1109/TCAD.2025.3524944. © 2026 IEEE]

5.6.1 Coherence Dependency Chains and Classical Virtual Networks

In directory-based protocols (such as MESI or MOESI), cache transactions generate a multi-tier dependency chain:

$$\text{Request } (M_{req}) \longrightarrow \text{Forwarded Invalidation } (M_{fwd}) \longrightarrow \text{Response / Data } (M_{resp}) \tag{5.7}$$

If all three message classes share a common buffer pool in the router, a catastrophic protocol deadlock occurs:

1. Core A issues a memory read request (M_{req}) toward Directory D.
2. Directory D accepts the request and attempts to inject an invalidation probe (M_{fwd}) to Core B, but the network is full.
3. Core B cannot accept incoming messages because its local input buffer is occupied by an unserviced request waiting for an outgoing response buffer to clear.
4. A circular wait is established across the memory controllers and network buffers, freezing the entire multi-core chip even though all network routers use deadlock-free XY routing.

To break this protocol-level dependency, classical systems mandate **Virtual Networks (VNs)**:

$$VN_{req} \cap VN_{fwd} = \emptyset, \quad VN_{fwd} \cap VN_{resp} = \emptyset \tag{5.8}$$

Each Virtual Network is assigned dedicated, physically isolated virtual channels. However, Jia et al. [26] proved that static VN partitioning incurs severe **Head-of-Line (HoL) blocking and buffer fragmentation**:

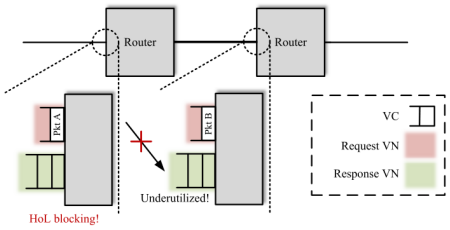


Fig. 2. Possible HoL blocking occurs with underutilized VCs.

Fig. 5.2 Underutilization and Head-of-Line blocking in virtual networks. Dedicated partition walls force flits to stall while response VCs remain empty. [Reproduced from: Y. Jia et al., "PCMT," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, Vol. 45, No. 4, April 2026, DOI: 10.1109/TCAD.2025.3524944. © 2026 IEEE]

Under write-intensive benchmarks, VN_{req} buffers overflow and throttle injection, while adjacent VN_{resp} buffers sit completely idle, wasting up to 65% of on-chip SRAM capacity.

5.6.2 PCMT: Prioritizing Coherence Message Types

To eliminate rigid VN partition walls, Jia et al. [26] introduced **PCMT (Prioritizing Coherence Message Types)**, uniting all virtual channels into a **single shared buffer pool**:

1. **Dynamic Priority Schedulers**: Every flit header is stamped with its protocol message category. The switch allocator enforces strict non-preemptive priority:

$$\text{Priority}(M_{resp}) > \text{Priority}(M_{fwd}) > \text{Priority}(M_{req}) \tag{5.9}$$

   Because Responses sit at the sink end of the coherence dependency chain and require no further outbound network messages to be consumed by a core, granting them absolute arbitration priority ensures they drain immediately.
2. **Preemptive Buffer Eviction and Anti-Starvation Counters**: To prevent lower-priority Requests from starving during prolonged Response bursts, PCMT implements an age threshold counter τ_{starve}. Flits waiting longer than τ_{starve} cycles temporarily promote to highest priority, guaranteeing bounded latency.

In 7 nm FinFET standard-cell synthesis, PCMT achieves protocol-level deadlock freedom while **slashing router input buffer silicon area by 63%** and improving saturated throughput by 34% over rigid virtual networks.

5.6.3 Detailed Coherence Transaction Sequences and Race Deadlocks in MOESI

To visualize how protocol deadlock emerges in real silicon, consider a directory-based MOESI cache coherence protocol with distributed memory controllers across an 8×8 mesh. When Core C_0 attempts to write to memory address A currently cached in Shared (S) state by Core C_1 and Owned (O) state by Core C_2:

1. **Phase 1 (Request)**: Core C_0 injects a `GetM` (Get Modified) packet on Virtual Network VN_{req} addressed to the Home Directory $D(A)$.
2. **Phase 2 (Forward & Invalidate)**: Directory $D(A)$ receives `GetM`, transitions its directory state for line A to the transient state M^D (Modified-Pending), and injects:

- A forwarded request `FwdGetM` on VN_{fwd} to Owner Core C_2;
- An invalidation probe `Inv` on VN_{fwd} to Sharer Core C_1.

3. **Phase 3 (Response & Invalidation Ack)**:
   - Owner Core C_2 receives `FwdGetM`, downgrades its local cache line from O to I (Invalid), and transmits the 64-byte data payload `Data` on VN_{resp} directly to requestor C_0.
   - Sharer Core C_1 receives `Inv`, invalidates its local copy, and transmits an `InvAck` acknowledgment on VN_{resp} to requestor C_0.
4. **Phase 4 (Completion)**: Once Core C_0 collects both the `Data` payload and all pending `InvAck` tokens, it transitions its cache line to Modified (M), completes the processor store instruction, and transmits a final `CompAck` on VN_{resp} to Directory $D(A)$ to clear the transient lock.

The Unbuffered Circular Wait Disaster: Suppose an architect mistakenly pools all message types into a shared physical buffer without priority preemption. If multiple cores concurrently launch store instructions, VN_{req} flits saturate all input buffers. When Directory $D(A)$ attempts to forward `FwdGetM`, its egress buffer is blocked by incoming `GetM` flits. Owner C_2 cannot consume incoming requests because its internal outbound queue is blocked waiting to inject `Data`. Because C_0 cannot receive `Data`, it cannot retire its store, establishing an unbreakable **circular wait across the memory hierarchy**.

5.7 Topology-Agnostic Deadlock-Free Routing: Up*/Down* and Segment-Based

While dimension-order and turn model routing operate effectively on regular orthogonal topologies (meshes and tori), modern SoCs frequently contain irregular geometries caused by heterogeneous block sizes, memory controller placement, and permanent silicon defects.

5.7.1 The Up*/Down* Routing Algorithm

Schroeder et al. [48] established **Up*/Down* Routing**, a provably deadlock-free algorithm applicable to **any arbitrary connected graph topology** $G = (V, E)$.

The algorithm proceeds in three phases:

1. **Spanning Tree Construction**: A single node is selected as the root node $r \in V$. A breadth-first search (BFS) constructs a spanning tree T_{BFS}, assigning each node $u \in V$ an integer depth $d(u)$ representing its shortest hop distance from root r.
2. **Edge Direction Assignment**: Every bidirectional physical channel $(u, v) \in E$ is partitioned into two directed channels categorized as **Up** or **Down**:

$$\text{Channel } (u \to v) \text{ is } \begin{cases} \textbf{Up}, & \text{if } d(u) > d(v) \text{ or } (d(u) = d(v) \text{ and } \mathrm{ID}(u) > \mathrm{ID}(v)) \\ \textbf{Down}, & \text{if } d(u) < d(v) \text{ or } (d(u) = d(v) \text{ and } \mathrm{ID}(u) < \mathrm{ID}(v)) \end{cases} \tag{5.10}$$

3. **The Up*/Down* Routing Rule**: A packet path is legal if and only if it consists of zero or more **Up** transitions followed by zero or more **Down** transitions:

$$\text{Legal Path Pattern}: \quad [\textbf{Up}]^* \cdot [\textbf{Down}]^* \tag{5.11}$$

   Once a packet traverses a **Down** link, it is strictly forbidden from ever traversing another **Up** link!

Theorem 5.6 (Deadlock Freedom of Up*/Down* Routing) *Under the Up*/Down* routing rule, the Channel Dependency Graph contains no directed cycles on any connected graph.*

Proof. Suppose, for contradiction, that a directed cycle exists in the CDG: $c_0 \to c_1 \to \cdots \to c_{k-1} \to c_0$.

1. A cycle must contain at least one transition from an incoming channel c_i to an outgoing channel c_{i+1} where the node depth increases and subsequently decreases, or vice-versa.
2. However, because Down-to-Up transitions ([**Down**] $\to$ [**Up**]) are strictly prohibited, no path can transition from a higher depth back to a lower depth after descending.
3. Channels can only transition [**Up**] $\to$ [**Up**], [**Up**] $\to$ [**Down**], or [**Down**] $\to$ [**Down**].
4. Assigning channel rank $\mathrm{rank}(c_{u\to v}) = -d(v)$ for Up channels and $\mathrm{rank}(c_{u\to v}) = +d(v)$ for Down channels guarantees that every legal transition strictly increases topological rank.
5. Because channel rank is strictly monotonic, circular wait is topologically impossible.

□

5.7.2 The Root Bottleneck and Segment-Based Routing (SR)

While Up*/Down* routing guarantees deadlock freedom on irregular fabrics, it exhibits a severe performance limitation: **The Root Bottleneck**. Because packets between distant nodes must climb the spanning tree before descending,

links near the root node r experience severe congestion, saturating long before periphery channels.

To eliminate the root bottleneck without adding virtual channels, Mejia et al. [40] introduced **Segment-Based Routing (SR)**. SR partitions the irregular graph into small, locally independent sub-graphs called *segments*. By prohibiting turns locally within each segment rather than enforcing a single global root hierarchy, SR distributes traffic uniformly across all links, achieving up to **2.5× higher saturation throughput** than standard Up*/Down* routing.

5.8 3D Multi-Chiplet Deadlock Freedom: ScaDA-Rou

In 3D multi-chiplet systems, vertical Through-Silicon Vias (TSVs) introduce cross-plane dependencies that interact catastrophically with planar turn models.

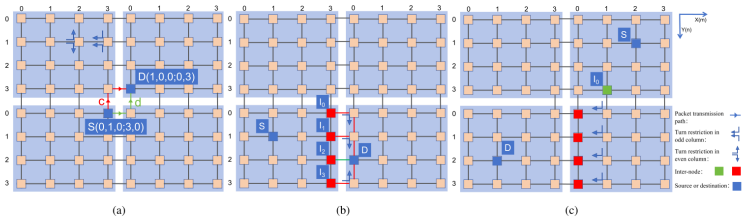


Fig. 2. Special consideration in 2D routing: (a) selection of inter-Chiplet routing order, (b) corner cases in X-dimension routing, (c) corner cases in Y-dimension routing.

Fig. 5.3 ScaDA-Rou routing rules for 3D multi-chiplet systems: (a) inter-chiplet routing order; (b) boundary corner cases in X-dimension; (c) boundary corner cases in Y-dimension. [Reproduced from: L. Zhu, H. Gu, Y. Tan, and G. Zhang, "ScaDA-Rou: A Scalable, Deadlock-Free, and Adaptive Routing for 3-D Multi-Chiplet System," *IEEE Trans. Circuits Syst. II, Express Briefs*, Vol. 73, No. 4, pp. 433–437, April 2026, DOI: 10.1109/TCSII.2025.3526101. © 2026 IEEE]

Zhu et al. [57] proposed **ScaDA-Rou**, combining in-plane turn model restrictions with cross-plane virtual network partitioning. By separating the virtual channels used for the source and destination planes of vertical packets, ScaDA-Rou eliminates 3D cyclic dependencies and achieves an **89.2% latency improvement** over baseline architectures.

5.9 Routing Algorithm Comparative Matrix

Table 5.1 summarizes the properties of the foundational and modern routing algorithms.

Table 5.1 Comparative Matrix of On-Chip Routing Algorithms

Routing Scheme	Turn Prohibitions	Min. VCs	Adaptiveness	Throughput Gain
Dimension-Order (XY)	$Y \to X$ prohibited	1	Deterministic	Baseline
West-First (Glass-Ni)	$N \to W, S \to W$	1	Partially Adaptive	+25%
North-Last (Glass-Ni)	$N \to E, N \to W$	1	Partially Adaptive	+25%
Negative-First	$+ \to -$ prohibited	1	Partially Adaptive	+20%
Odd-Even (Chiu)	Column-parity dependent	1	Partially Adaptive	+35%
Duato FAR + XY Escape	None in C_A	2	Fully Adaptive	+75%–100%
PCMT (Jia et al., 2026)	Priority-preemptive	Unified Pool	Protocol Deadlock-Free	-63% Area

5.10 Interactive Lab & Hands-on Challenges

Explore the interactive routing and deadlock visualizers on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **Deadlock Buster Arcade Game** (https://network-on-chip.github.io/learning/interactive/game/): Defend the router fabric against circular wait dependency cycles, resolving routing knots in real time.
- **Packet Master: Routing Puzzle** (https://network-on-chip.github.io/learning/interactive/puzzle/): Guide head and body flits through multi-stage wormhole channels to experience turn model restrictions first-hand.

5.11 Exercises and Review Questions

1. **CDG Construction**: For a 2×2 2D Mesh, enumerate all directed physical channels and draw the Channel Dependency Graph under fully adaptive minimal routing. Identify the cycle and show how West-First routing breaks it.
2. **Duato Theorem Dynamics**: In a network deploying Duato's escape routing, why must packets on the escape subnetwork (VC_0) be prohibited from transitioning back into the adaptive subnetwork (VC_1)?
3. **Protocol Deadlock Scenario**: Sketch a sequence of coherence transactions in a MESI protocol where all routers utilize deadlock-free XY routing, yet the system suffers protocol-level deadlock due to full request buffers.

Part III
Physical Constraints: Power, Thermal & Reliability

Chapter 6
Low-Power Architectures & Dark Silicon Dynamics

Abstract The demise of classical Dennard scaling has driven modern semiconductor manycores into the dark silicon era, where up to 80% of the silicon floorplan must remain powered off or severely throttled under strict Thermal Design Power (TDP) constraints. While idle compute cores can be gated completely, on-chip interconnection routers must remain active to forward transit flits, making the network a primary power and thermal bottleneck. This chapter provides a comprehensive treatment of low-power NoC architectures. We formulate dynamic and static leakage power dynamics, model temperature-dependent sub-threshold leakage, and derive the breakeven idle time for power gating. We examine Dynamic Voltage and Frequency Scaling (DVFS) across Voltage-Frequency Islands (VFIs) and analyze the LAURA-NoC distributed rate adjustment architecture. Furthermore, we investigate ultrafine-grained run-time power gating across router virtual channels and crossbar multiplexers with lookahead early wake-up, analyze information-theoretic link code reordering (LPCR, IEEE TCAD 2025), and contrast buffered routers with bufferless deflection routing.

6.1 The Power Wall and the Dark Silicon Interconnect Bottleneck

At sub-7nm process nodes, the physical breakdown of classical Dennard scaling imposes an unyielding thermal dissipation ceiling ($\approx$ 100–150 W/cm^2). Consequently, modern manycore chips face the **Dark Silicon Phenomenon** [16]: between 50% and 80% of a chip's total transistor budget cannot be operated simultaneously at peak clock frequency without triggering thermal destruction.

While idle compute cores can be power-gated to zero power for milliseconds at a time, the **on-chip network cannot be casually turned off**. In multi-hop topologies (such as 2D Meshes), an intermediate router must remain awake to forward packets between active compute tiles. Routers sit active but idle for up to 90% of total execution time [37], dissipating substantial static leakage power.

6.1.1 Dennard Breakdown and the Dark Silicon Scaling Model

To understand why power gating and voltage islands are mandatory in modern on-chip networks, one must examine the mathematical breakdown of classical Dennard scaling.

Under ideal Dennard scaling (formulated by Robert Dennard in 1974), all spatial geometric dimensions of a MOSFET scale down by a constant scaling factor $S > 1$ (where $S \approx \sqrt{2} \approx 1.4$ per technology generation):

$$L \longrightarrow \frac{L}{S}, \quad W \longrightarrow \frac{W}{S}, \quad t_{ox} \longrightarrow \frac{t_{ox}}{S}, \quad N_A \longrightarrow S \cdot N_A \tag{6.1}$$

To maintain a constant electric field ($\mathcal{E} = V/t_{ox}$), the supply voltage V_{dd} and threshold voltage V_{th} were scaled simultaneously by $1/S$:

$$V_{dd} \longrightarrow \frac{V_{dd}}{S}, \quad V_{th} \longrightarrow \frac{V_{th}}{S} \tag{6.2}$$

Consequently, gate capacitance scaled as $C = \epsilon A/t_{ox} \rightarrow C/S$, gate delay improved as $\tau = CV/I \rightarrow \tau/S$, allowing clock frequency to increase as $f_{clk} \rightarrow S \cdot f_{clk}$. The dynamic power per transistor evolved as:

$$P_{trans} = C \cdot V_{dd}^2 \cdot f_{clk} \longrightarrow \left(\frac{C}{S}\right)\left(\frac{V_{dd}}{S}\right)^2 (S \cdot f_{clk}) = \frac{P_{trans}}{S^2} \tag{6.3}$$

Because transistor density increased by S^2 per unit area, the total power density across the silicon die remained perfectly invariant:

$$\text{Power Density} = \frac{N_{trans} \cdot P_{trans}}{\text{Die Area}} \longrightarrow \frac{(S^2 \cdot N_{trans}) \cdot (P_{trans}/S^2)}{\text{Area}} = \text{Constant} \tag{6.4}$$

Around 2005 (at the 90nm/65nm transition), Dennard scaling collapsed because **voltage scaling stalled**. The sub-threshold slope of a MOSFET is fundamentally constrained by Boltzmann thermodynamics:

$$S_S = \ln(10) \cdot \frac{k_B T}{q} \cdot \left(1 + \frac{C_d}{C_{ox}}\right) \geq 60 \text{ mV/decade at 300 K} \tag{6.5}$$

where k_B is the Boltzmann constant, T is absolute temperature, and q is elementary charge. Because S_S cannot be scaled below 60 mV/decade at room temperature, reducing V_{th} causes off-state drain current I_{off} to increase exponentially:

$$I_{off} = I_0 \cdot 10^{-V_{th}/S_S} \tag{6.6}$$

To prevent static leakage from overwhelming total chip power, designers had to clamp V_{dd} to ≈ 0.75–0.95 V.

With V_{dd} fixed while transistor density continues to scale at S^2, chip power density escalates by S^2 each generation. Esmaeilzadeh et al. [16] formulated the

Dark Silicon Utilization Fraction $U(S)$:

$$U(S) = \frac{P_{TDP}}{P_{active}(S)} = \frac{P_{TDP}}{N_{trans}(S) \cdot C(S) \cdot V_{dd}^2(S) \cdot f_{clk}(S)} \approx \frac{1}{S^2} \left(\frac{V_{dd,0}}{V_{dd}(S)} \right)^2 \left(\frac{f_0}{f_{clk}(S)} \right) \tag{6.7}$$

At advanced FinFET nodes (sub-5nm), $U(S) < 0.20$, proving that **over 80% of the silicon floorplan must remain powered off or severely underclocked** at any given instant.

6.1.2 Physical Power Breakdown in Nanoscale Routers

Total interconnect power dissipation is decomposed into dynamic and static leakage components:

$$P_{total} = P_{dynamic} + P_{leakage} \tag{6.8}$$

Dynamic Power ($P_{dynamic}$): Consumed during capacitive charging and discharging across buffer cells, multiplexer trees, pipeline registers, and metal link wires:

$$P_{dynamic} = \alpha \cdot C_{eff} \cdot V_{dd}^2 \cdot f_{clk} \tag{6.9}$$

where α is the switching activity factor, C_{eff} is effective capacitance, V_{dd} is supply voltage, and f_{clk} is operating frequency.

Static Leakage Power ($P_{leakage}$): Dominated by sub-threshold drain-source conduction I_{sub} and gate oxide quantum tunneling:

$$I_{sub} = \mu_0 C_{ox} \frac{W}{L} (m-1) V_T^2 \exp\left(\frac{V_{gs} - V_{th}}{m V_T} \right) \left(1 - \exp\left(-\frac{V_{ds}}{V_T} \right) \right) \tag{6.10}$$

Sub-threshold leakage scales exponentially with absolute temperature T ($V_T = k_B T/q$). In sub-28nm geometries, static leakage constitutes between 40% and 80% of total router power dissipation.

6.2 DVFS and Voltage-Frequency Islands (VFIs)

Dividing a large manycore into **Voltage-Frequency Islands (VFIs)** enables independent, dynamic adjustment of supply voltage and clock frequency across local network domains.

Yadav, Casu, and Zamboni [54] developed **LAURA-NoC**, which eliminates global controller latency by measuring local switch buffer fullness to drive local voltage-frequency transitions. In 45nm CMOS, LAURA-NoC achieved **33% average power savings** across realistic video processing workloads.

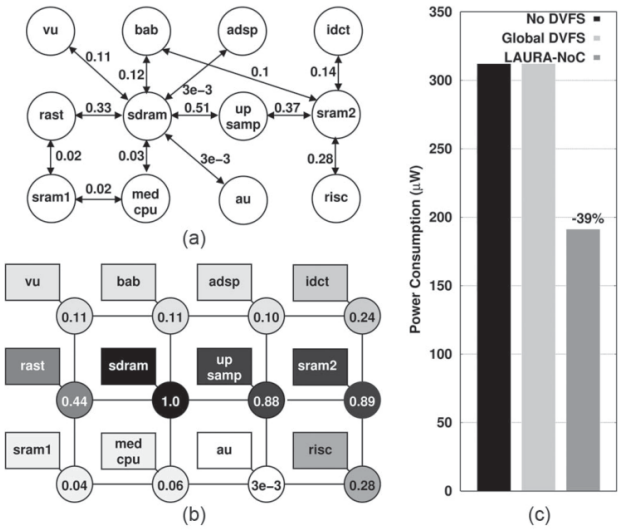


Fig. 1. Local versus global DVFS in a NoC. (a) Communication graph in MPEG4 annotated with normalized rates. (b) MPEG4 tasks mapped to a 12-switch NoC (gray shading proportional to normalized switch rates). (c) Power consumption: The local approach of LAURA-NoC tracks better the local traffic by controlling voltage and frequency of individual switches.

Fig. 6.1 The LAURA-NoC distributed DVFS architecture: local hardware monitors input switch buffer occupancy in a feedback loop, dynamically adapting voltage and frequency without centralized controller overhead. [Reproduced from: M. K. Yadav, M. R. Casu, and M. Zamboni, "LAURA-NoC: Local Automatic Rate Adjustment in Network-on-Chips With a Simple DVFS," *IEEE Trans. Circuits Syst. II, Express Briefs*, Vol. 60, No. 10, pp. 647–651, October 2013, DOI: 10.1109/TCSII.2013.2278077. © 2013 IEEE]

6.2.1 Asynchronous Clock Domain Crossing (CDC) and Dual-Clock FIFOs

When neighboring routers or tiles reside in distinct VFIs, signals must cross asynchronous clock domain boundaries. Direct sampling of asynchronous signals induces catastrophic **metastability**, where flip-flop internal nodes linger at an indeterminate voltage between 0 and 1, producing non-deterministic timing violations.

To guarantee robust inter-VFI communication, link interfaces employ **Dual-Clock Asynchronous FIFOs** with Gray-coded read and write pointers:

1. **Gray-Code Pointer Synchronization**: The binary write pointer $B = (b_n, \ldots, b_0)$ is converted to Gray code $G = (g_n, \ldots, g_0)$ via $g_i = b_i \oplus b_{i+1}$. Gray codes ensure that **exactly one bit changes value per pointer increment**. Consequently, clock skew across the asynchronous boundary can at

most cause the synchronizer to read the old pointer value or the new pointer value, but never an invalid intermediate state.

2. **Two-Stage Synchronizer Reliability**: Dual-stage flip-flop synchronizers achieve an astronomical **Mean Time Between Failures (MTBF)**:

$$MTBF = \frac{\exp(t_{settle}/\tau_s)}{T_0 \cdot f_{dest} \cdot f_{src}} \tag{6.11}$$

where τ_s is the latch regeneration time constant, T_0 is the aperture parameter, t_{settle} is the available settling time (one destination clock period minus setup time), f_{src} is the transmitter clock frequency, and f_{dest} is the receiver clock frequency.

3. **Mixed-Clock Credit Loop Latency**: Flow control credits traversing the reverse path must also cross a two-stage synchronizer, introducing an unavoidable latency overhead of 2–3 destination clock cycles into the credit return path. To sustain peak throughput without stalls, the asynchronous FIFO depth must satisfy:

$$\text{Depth}_{CDC} \geq 2{\cdot}(RTT_{credit}{+}\text{Sync Delay}) = 2{\cdot}(t_{link}{+}2{\cdot}T_{src}{+}2{\cdot}T_{dest}) \tag{6.12}$$

6.2.2 Feedback Control Theory for Autonomous VFI Frequency Scaling

While centralized operating system governors can modulate tile frequencies at millisecond granularity, on-chip network packet bursts occur at nanosecond timescales (10–100 clock cycles). Autonomous VFIs deploy hardware-synthesized **Proportional-Integral-Derivative (PID) feedback controllers** embedded directly within router input controllers.

Let $q(t)$ denote the instantaneous flit occupancy of the input buffer, and let Q_{target} be the optimal setpoint (typically 30%–50% of maximum FIFO capacity to balance throughput against queuing latency). The error signal is:

$$e(t) = q(t) - Q_{\text{target}} \tag{6.13}$$

The continuous rate of change of buffer fullness is governed by the mass balance equation:

$$\frac{dq(t)}{dt} = \lambda_{\text{in}}(t) - \mu(f_{\text{clk}}(t)) \tag{6.14}$$

where $\mu(f_{\text{clk}}) \approx \kappa_\mu \cdot f_{\text{clk}}$ represents the crossbar flit forwarding rate.

In discrete hardware, the PID controller updates the clock frequency setpoint at sampling period T_s:

$$f_{\text{clk}}[k+1] = f_{\text{clk}}[k] + K_p \cdot e[k] + K_i \sum_{m=0}^{k} e[m] + K_d(e[k] - e[k-1]) \tag{6.15}$$

where K_p, K_i, K_d are the proportional, integral, and derivative control gains.

Closed-Loop Stability & Phase Margin: By taking the $\mathcal{Z}$-transform of Eq. (6.15) combined with the buffer integrator plant $G(z) = \frac{T_s}{z-1}$, the closed-loop characteristic equation is:

$$1+D(z)G(z) = 0 \iff z^2+(K_pT_s + K_iT_s + K_d - 2)\, z+(1 - K_pT_s - 2K_d) = 0 \tag{6.16}$$

Applying Jury's stability criterion mandates that the loop gains satisfy $K_iT_s < 2$ and $K_d < \frac{1}{2}(1 - K_pT_s)$. To prevent destructive limit-cycle hunting under bursty traffic, controllers enforce a **Bode phase margin** of $\phi_m \geq 60^\circ$ and implement deadband hysteresis around Q_{target}.

Furthermore, physical hardware actuators impose strict transient slew rate constraints:

- **Phase-Locked Loop (PLL) Relock Delay**: Frequency changes require $t_{\text{relock}} \approx 20$–$50\,\text{ns}$ for voltage-controlled oscillators to stabilize without clock glitching.
- **Integrated Voltage Regulators (IVR)**: On-chip buck converters and switched-capacitor DC-DC converters are limited by an inductive slew rate ceiling ($dV/dt \leq 10$–$25\,\text{mV}/\mu\text{s}$), dictating that frequency steps must lead voltage drops but lag voltage rises to prevent timing failure.

6.2.3 Near-Threshold Computing (NTC) Interconnects & Razor Timing Speculation

The most aggressive low-power regime is **Near-Threshold Computing (NTC)**, where supply voltage is scaled down to just above the transistor threshold ($V_{dd} \approx V_{th} + 100$–$150\,\text{mV} \approx 0.40$–$0.50\,\text{V}$). Because dynamic energy scales quadratically ($E \propto CV_{dd}^2$), NTC slashes interconnect energy by $5\times$ **to** $10\times$.

However, NTC faces the **Variability Wall**. Under the alpha-power law model, CMOS gate delay t_d is:

$$t_d = \frac{C_{\text{load}} \cdot V_{dd}}{\mu C_{ox} \frac{W}{L} (V_{dd} - V_{th})^\alpha} \tag{6.17}$$

with velocity saturation index $\alpha \approx 1.3$–1.5. Differentiating Eq. (6.17) with respect to V_{dd} reveals the normalized delay sensitivity:

$$\mathcal{S}_{V_{dd}}^{t_d} = \frac{\partial t_d / t_d}{\partial V_{dd} / V_{dd}} = 1 - \frac{\alpha V_{dd}}{V_{dd} - V_{th}} \tag{6.18}$$

At nominal $V_{dd} = 0.90\,\mathrm{V}$ ($V_{th} = 0.35\,\mathrm{V}$), $|\mathcal{S}| \approx 1.1$. But in NTC ($V_{dd} = 0.45\,\mathrm{V}$), $V_{dd} - V_{th} = 0.10\,\mathrm{V}$, causing sensitivity to escalate to $|\mathcal{S}| \approx 5.8$! A mere 5% power grid IR-drop (22 mV) triggers a catastrophic 30% **gate delay surge**, causing widespread setup-time violations.

To operate safely in NTC without exorbitant worst-case guardbands, routers employ **Razor Timing Speculation**:

1. Each critical flip-flop is paired with a delayed **Shadow Latch** clocked by a delayed clock phase ($\Delta t_{\mathrm{margin}}$).
2. If a slow path violates setup time at the main flip-flop, the shadow latch captures the correct late-arriving data.
3. An XOR comparator flags the timing error, asserts a pipeline squash pulse, replays the cycle from the shadow latch, and dynamically nudges the local VFI voltage regulator higher.

6.3 Ultrafine-Grained Run-Time Power Gating

Power gating completely disconnects inactive logic blocks from the power rails using high-threshold sleep transistors (power switches).

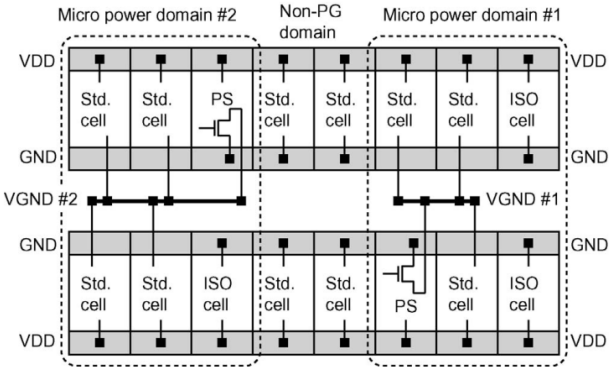


Fig. 3. Concept of the fine-grained power gating. PS and ISO refer to a power switch and an isolation cell, respectively.

Fig. 6.2 Concept of fine-grained run-time power gating. (a) Active mode: power switch (PS) is closed, virtual VDD rail is at nominal potential. (b) Sleep mode: PS opens, cutting leakage; isolation cells (ISO) clamp boundary outputs to prevent downstream crowbar currents. [Reproduced from: H. Matsutani et al., "Performance, Area, and Power Evaluations of Ultrafine-Grained Run-Time Power-Gating Routers for CMPs," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, Vol. 30, No. 4, pp. 520–533, April 2011, DOI: 10.1109/TCAD.2010.2097310. © 2011 IEEE]

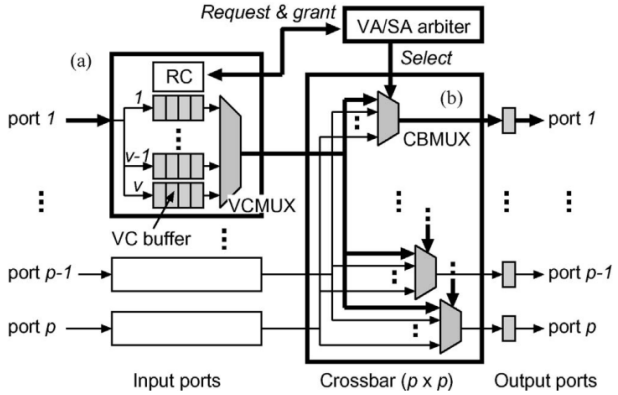


Fig. 4. Ultrafine-grained power-gating router. (a) Detail of input port 1. (b) Detail of crossbar. Each gray area denotes a power domain.

Fig. 6.3 Ultrafine-grained power-gating router microarchitecture: (a) detail of input port virtual channel buffers and multiplexers; (b) detail of crossbar multiplexers. Each shaded partition represents an independently gated domain. [Reproduced from: H. Matsutani et al., "Performance, Area, and Power Evaluations of Ultrafine-Grained Run-Time Power-Gating Routers for CMPs," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, Vol. 30, No. 4, pp. 520–533, April 2011, DOI: 10.1109/TCAD.2010.2097310. © 2011 IEEE]

Matsutani et al. [37] proved that ultrafine-grained partitioning of router buffers and crossbar multiplexer slices eliminates **over 70% of static leakage**.

6.3.1 Sleep Transistor Sizing and In-Rush Damping

Sleep transistors are deployed either as **Header switches** (PMOS transistors inserted between true V_{dd} and virtual V_{ddv}) or **Footer switches** (NMOS transistors inserted between virtual ground V_{ssv} and true ground). While NMOS footers provide higher electron mobility ($\mu_n \approx 2\text{–}3 \times \mu_p$) and smaller silicon area, PMOS headers are vastly preferred in NoCs because modulating ground potential disrupts substrate body bias and induces substrate noise coupling across adjacent analog/memory IPs.

The width of the sleep transistor W_{sleep} is strictly governed by the peak active current I_{peak} and the maximum allowable virtual rail IR-drop budget ΔV_{drop} (typically $\leq 1.5\%$ of V_{dd}):

$$W_{sleep} = \frac{I_{peak}}{\mu_p C_{ox} \left(\frac{W}{L}\right)_{min} (V_{dd} - |V_{th,p}|) \cdot \Delta V_{drop}} \tag{6.19}$$

When awakening a sleeping router partition, abrupt gating causes a massive surge of **in-rush current** to charge the virtual rail capacitance $C_{virtual}$. This sudden current transient induces severe inductive supply bounce across package bond wires:

$$\Delta V_{bounce} = L_{pkg} \frac{di_{rush}}{dt} + R_{pkg} \cdot i_{rush} \tag{6.20}$$

Excessive ΔV_{bounce} can inadvertently flip state bits in adjacent active SRAM buffers. To mitigate this hazard, routers deploy **Two-Stage Sleep Switches**:

- **Stage 1 (Weak Pre-Charge)**: A small, high-resistance sleep transistor turns on first, limiting di/dt while charging V_{ddv} smoothly to $\approx 80\%$ of nominal voltage.
- **Stage 2 (Full Rail Clamping)**: The massive low-resistance main sleep transistor closes, clamping V_{ddv} to full supply with minimal inductive ringing.

6.3.2 Breakeven Idle Time Formulation

Transitioning into and out of sleep mode consumes energy charging the virtual rail capacitance ($C_{virtual}$):

$$E_{overhead} \approx C_{virtual} \cdot V_{dd}^2 \tag{6.21}$$

The minimum inactive duration required to achieve net energy savings is the **Breakeven Idle Time ($\tau_{breakeven}$)**:

$$\tau_{breakeven} = \frac{E_{overhead}}{P_{leak_active}} = \frac{C_{virtual} \cdot V_{dd}}{I_{leak_active}} \tag{6.22}$$

In modern technologies, $\tau_{breakeven}$ spans only 2 to 6 clock cycles, making run-time power gating highly energy-positive.

6.3.3 Lookahead Early Wake-Up

To eliminate the 2-to-4 cycle wakeup latency penalty, upstream routers transmit an early wakeup pulse across dedicated physical wires one hop ahead of flit departure, charging the downstream virtual rail concurrently with upstream crossbar traversal.

6.4 Low-Power Link Encoding: The LPCR Framework

Long inter-router metal links exhibit severe inter-wire coupling capacitance ($C_{coupling} \approx 3$–$5 \times C_{ground}$). Opposite-direction signal transitions ($01 \rightarrow 10$) trigger Miller effect capacitance doubling.

BMP				DOC			
Highest		Lowest		Highest		Lowest	
Dec.	Trans.	Dec.	Trans.	Dec.	Trans.	Dec.	Trans.
250	3	254	1	0	0	61	3
255	0	252	1	1	1	107	5
26	4	253	2	109	5	14	2
182	5	239	2	111	3	135	2
213	6	242	3	101	5	57	3
JPG				Overall			
Highest		Lowest		Highest		Lowest	
Dec.	Trans.	Dec.	Trans.	Dec.	Trans.	Dec.	Trans.
1	1	64	2	0	0	252	1
0	0	96	2	1	1	242	3
248	1	160	3	3	1	254	1

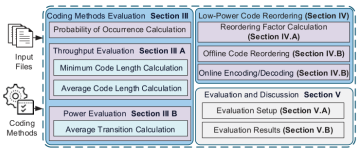


Fig. 2. Overview of our novel contributions (in blue boxes).

Fig. 6.4 The LPCR (Low-Power Code Reordering) framework: information-theoretic analysis of data streams reorders binary codewords to minimize bit transitions across global serial links. [Reproduced from: M. Adelkhani et al., "LPCR: Information-Theory-Based Low-Power Code Reordering for Serial Links in Network-on-Chip," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, Vol. 44, No. 11, pp. 4155–4168, November 2025, DOI: 10.1109/TCAD.2025.3411245. © 2025 IEEE]

Adelkhani et al. [1] developed **LPCR**, mapping high-frequency data symbols onto minimum-transition binary codewords and slashing link dynamic power by up to 35% without throughput degradation.

6.5 Bufferless Deflection Routing

Moscibroda and Mutlu [42] proposed an extreme low-power alternative: **eliminating router input FIFO buffers entirely** via bufferless deflection routing (such as BLESS and CHIPPER).

6.5.1 Microarchitectural Mechanics and Flit Arbitration

In a bufferless router, flits never stall in place awaiting available downstream buffers. Instead, every flit arriving at an input port in cycle t **must exit through an output port in cycle** $t + 1$:

1. **Port Contention and Deflection**: When multiple incoming flits request the same output port, the arbiter grants the port to the highest-priority flit. The losing flit is not stalled; rather, it is **deflected** onto an alternative, non-optimal output port (hot-potato routing).
2. **Livelock Freedom Proof**: To prevent deflected flits from wandering endlessly across the network without arriving at their destination, BLESS en-

forces an **Oldest-First Priority Policy**:

$$\text{Priority}(F_i) > \text{Priority}(F_j) \iff \text{Age}(F_i) > \text{Age}(F_j) \tag{6.23}$$

Because every flit's age increments monotonically every clock cycle, the globally oldest flit in the fabric is guaranteed to win every port arbitration. It routes strictly along minimal paths without deflection until reaching its destination. Once it exits, the next oldest flit acquires peak priority. By mathematical induction, **livelock is provably impossible**.

6.5.2 The Energy Crossover Threshold

While bufferless deflection routing completely eliminates 60% to 75% of router silicon area and removes all buffer static leakage power, it incurs significant trade-offs:

- **Reassembly Buffer Overhead**: Due to deflections, flits belonging to the same packet arrive out of order at the destination Network Interface (NI). The NI must maintain dedicated reassembly buffers sized proportional to the maximum latency skew:

$$\text{Buffer}_{reassembly} \geq B_W \times \Delta t_{skew} \tag{6.24}$$

- **Dynamic Link Energy Crossover**: Every deflected hop causes an extra physical wire traversal, consuming dynamic link power $\alpha C_{link} V_{dd}^2$. At low injection loads ($\lambda < 0.15$ flits/node/cycle), deflections are rare, and bufferless networks achieve up to 40% total energy savings. However, when injection exceeds the **crossover rate** ($\lambda \geq 0.20$), deflection cascading causes dynamic link power to exceed the static buffer leakage savings, rendering bufferless architectures inefficient under high load.

6.6 Quantitative Evaluation Matrix

Table 6.1 summarizes the trade-offs of low-power interconnect methodologies across an 8×8 mesh fabric.

6.7 Interactive Lab & Hands-on Challenges

Explore the interactive power and energy visualizers on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

Table 6.1 Quantitative Comparison of Low-Power Interconnect Architecture Paradigms

Architecture	Dynamic Power	Leakage Power	Area Overhead	Latency Cost
Baseline (Fixed 0.9V/1.2GHz)	Baseline	Baseline	Baseline	0 cycles
Distributed DVFS (LAURA-NoC)	-33% to -45%	-20%	+4%	0–1 cycle
Ultrafine Power Gating	Baseline	-70% to -95%	+8%	0 (Lookahead)
Co-Design (DVFS + PG)	-40%	-85%	+12%	1 cycle
Link Reordering (LPCR)	-25% (Links)	Baseline	+1.5%	0 cycles
Bufferless Deflection	+20% (Links)	-75%	-60% (No FIFO)	+2–6 hops

- **Volt Vortex: Dark Silicon Power** (https://network-on-chip.github.io/learning/interactive/power/): Balance active core voltage/frequency islands against router power-gating breakeven thresholds.
- **QoS Gladiator: Traffic Arbiter** (https://network-on-chip.github.io/learning/interactive/qos/): Schedule mixed-criticality traffic flows while managing dynamic voltage-frequency scaling margins.

6.8 Exercises and Review Questions

1. **Breakeven Time Calculation**: A router buffer in 28nm FinFET has $C_{virtual} = 3.2$ pF, $V_{dd} = 0.85$ V, and $I_{leak} = 8.5$ mA at 75°C. Calculate $E_{overhead}$ and $\tau_{breakeven}$. At 1.5 GHz, how many idle cycles justify sleeping?
2. **DVFS Power Scaling**: If supply voltage is reduced from 1.0 V to 0.75 V and frequency scales from 1.5 GHz to 1.0 GHz, calculate the percentage reduction in dynamic power.
3. **Lookahead Wakeup Hazards**: Describe the circuit consequences if a flit arrives at a power-gated buffer whose virtual power rail has only reached 50% of nominal voltage.

Chapter 7
Thermal Modeling, Hotspots & Fault-Tolerant Resilience

7.1 The Thermal Emergency in Deep Sub-Micron Silicon

In modern sub-5 nm multi-core architectures, power density frequently exceeds $100\,\mathrm{W/cm^2}$—a thermal flux comparable to the surface of the sun. Networks-on-Chip (NoC) routers and high-speed serial links contribute between 20% to 40% of overall chip power dissipation. Because network traffic is intrinsically bursty and spatially non-uniform, switching activity concentrates at topological bottlenecks (such as memory controller interfaces and bisection crossbars), forming extreme localized temperature spikes termed *thermal hotspots*.

The physical consequences of unmitigated thermal hotspots are severe across four fundamental failure mechanisms:

1. **Sub-Threshold Leakage Runaway**: Sub-threshold leakage current exhibits an exponential dependence on absolute operating temperature governed by the Arrhenius formulation:

$$I_{\text{leak}}(T) \propto T^2 \exp\left(-\frac{qV_{\text{th}}}{mk_BT}\right) \tag{7.1}$$

   where V_{th} is the threshold voltage, k_B is Boltzmann's constant, and q is the elementary electron charge. As local silicon heats up, V_{th} drops while thermal carrier excitation intensifies, triggering a positive feedback loop that culminates in *thermal runaway*.
2. **Carrier Mobility Degradation**: Acoustic phonon scattering increases with temperature, degrading electron and hole mobility according to:

$$\mu(T) = \mu(T_0)\left(\frac{T}{T_0}\right)^{-1.5} \tag{7.2}$$

   Reduced mobility diminishes transistor drive current (I_{on}), expanding router pipeline gate delays and causing timing closure violations.

3. **Accelerated Electromigration (EM)**: Spatial thermal gradients across adjacent router tiles induce intense thermo-mechanical shear stress due to mismatched coefficients of thermal expansion (CTE). Furthermore, according to *Black's Equation*, the Mean Time to Failure ($MTTF$) of copper interconnects drops exponentially with temperature:

$$MTTF_{EM} = \frac{A}{J^n} \exp\left(\frac{E_a}{k_B T}\right) \tag{7.3}$$

where J is the current density, $n \approx 1.5$–2.0 is the current density exponent, and $E_a \approx 0.9\,\text{eV}$ is the activation energy for copper grain-boundary diffusion. A temperature increase of only 15°C cuts interconnect lifetime by more than 50%.

4. **Dielectric Breakdown and Device Aging (TDDB and NBTI)**:
   - **Time-Dependent Dielectric Breakdown (TDDB)**: High gate oxide electric fields coupled with elevated temperatures accelerate gate insulator breakdown:

$$MTTF_{TDDB} \propto \left(\frac{1}{V_{dd}}\right)^{\gamma} \exp\left(\frac{E_{a,TDDB}}{k_B T}\right) \tag{7.4}$$

   where $E_{a,TDDB} \approx 0.6$–0.8 eV.
   - **Negative Bias Temperature Instability (NBTI)**: In PMOS transistors, prolonged negative gate bias at high temperatures breaks Si–H bonds at the oxide interface, creating interface traps that permanently shift the threshold voltage over time:

$$\Delta V_{th}(t) \propto \exp\left(-\frac{E_{a,NBTI}}{k_B T}\right) \cdot t^{n_{nbti}} \tag{7.5}$$

   where $n_{nbti} \approx 0.25$, degrading router setup time margins over several years of continuous operation.
   - **Thermo-Mechanical Fatigue**: Repetitive thermal cycling between idle and burst states induces thermal expansion mismatch fatigue governed by the **Coffin-Manson Relationship**:

$$N_f = C_0 \cdot (\Delta T - \Delta T_0)^{-q_{fatigue}} \tag{7.6}$$

   where N_f is the number of cycles to failure, and $q_{fatigue} \approx 2$–3 for microbumps and TSV joints.

7.2 Compact Thermal Physics: The Electro-Thermal RC Duality

Dynamic thermal management requires real-time simulation models capable of evaluating full-chip temperatures without the prohibitive computational complexity of Finite Element Methods (FEM). This is achieved through *compact thermal RC networks*, pioneered by the landmark **HotSpot** methodology [25].

7.2.1 Duality Mapping Between Thermal and Electrical Domains

Fourier's law of heat conduction states that heat flux density $\mathbf{q}$ ($\mathrm{W/m^2}$) is proportional to the negative spatial temperature gradient: $\mathbf{q} = -k_{\mathrm{th}} \nabla T$, where k_{th} is the thermal conductivity ($\mathrm{W/(m \cdot K)}$). Discretizing the silicon substrate and packaging ladder into volumetric nodes maps heat conduction directly onto electrical circuit theory, as summarized in Table 7.1.

Table 7.1 Duality Mapping Between Electrical and Thermal Physical Domains

Electrical Variable	**Unit**	**Thermal Analog**	**Unit**
Electric Potential (V)	Volts (V)	Temperature Rise (ΔT)	Kelvin (K)
Electric Current (I)	Amperes (A)	Heat Flow / Power (P)	Watts (W)
Electrical Resistance (R)	Ohms (Ω)	Thermal Resistance (R_{th})	K/W
Electrical Capacitance (C)	Farads (F)	Thermal Capacitance (C_{th})	J/K
Electrical Conductance (G)	Siemens (S)	Thermal Conductance (G_{th})	W/K
Kirchhoff's Current Law	$\sum I_k = C \frac{dV}{dt}$	Energy Conservation Law	$\sum P_k = C_{\mathrm{th}} \frac{dT}{dt}$

7.2.2 Multi-Layer Physical Heat Conduction Formulation

In modern high-performance packaging, heat generated in the active device layer traverses a vertical hierarchy:

$$\text{Active Silicon Layer} \longrightarrow \text{Bulk Silicon Substrate} \longrightarrow \text{Thermal Interface Material (TIM)} \longrightarrow \text{Heat Spreader} \longrightarrow \text{Heat} \tag{7.7}$$

For any volumetric block i with dimensions $L_x \times L_y \times t_{layer}$, mass density ρ, specific heat capacity c_p, and directional thermal conductivities (k_x, k_y, k_z), the physical lumped thermal elements are formulated from first principles as:

$$R_{v,i} = \frac{t_{layer}}{k_z \cdot (L_x \cdot L_y)} \quad \text{(Vertical Resistance to Layer Below)} \tag{7.8}$$

$$R_{L,x,ij} = \frac{L_x}{k_x \cdot (L_y \cdot t_{layer})} \quad \text{(Lateral Conduction in } X \text{ Dimension)} \tag{7.9}$$

$$C_{\text{th},i} = \rho \cdot c_p \cdot (L_x \cdot L_y \cdot t_{layer}) \quad \text{(Thermal Capacitance)} \tag{7.10}$$

7.2.3 State-Space Governing Differential Equations

Consider an $N \times N$ mesh of router tiles. Each tile i represents a lumped thermal node characterized by vertical resistance $R_{v,i}$, lateral resistances $R_{L,ij}$ to adjacent tiles $j \in \mathcal{N}(i)$, and thermal capacitance $C_{\text{th},i}$.

Applying energy conservation yields the continuous nodal differential equation:

$$C_{\text{th},i}\frac{dT_i(t)}{dt} + \frac{T_i(t) - T_{\text{amb}}}{R_{v,i}} + \sum_{j\in\mathcal{N}(i)} \frac{T_i(t) - T_j(t)}{R_{L,ij}} = P_i(t) \tag{7.11}$$

Defining the state vector of temperature rises $\mathbf{x}(t) = \mathbf{T}(t) - T_{\text{amb}}\mathbf{1}$ and the power input vector $\mathbf{u}(t) = \mathbf{P}(t)$, the continuous state-space system is written as:

$$\frac{d\mathbf{x}(t)}{dt} = \mathbf{A}_c\mathbf{x}(t) + \mathbf{B}_c\mathbf{u}(t) \tag{7.12}$$

where the continuous conductance matrix $\mathbf{A}_c \in \mathbb{R}^{M\times M}$ has entries:

$$(\mathbf{A}_c)_{i,i} = -\frac{1}{C_{\text{th},i}}\left(\frac{1}{R_{v,i}} + \sum_{j\in\mathcal{N}(i)} \frac{1}{R_{L,ij}}\right), \qquad (\mathbf{A}_c)_{i,j} = \frac{1}{C_{\text{th},i}R_{L,ij}} \quad (j \in \mathcal{N}(i)) \tag{7.13}$$

and $\mathbf{B}_c = \text{diag}(1/C_{\text{th},1}, \ldots, 1/C_{\text{th},M})$. Discretizing over sample period Δt yields:

$$\mathbf{x}_{k+1} = \mathbf{A}_d\mathbf{x}_k + \mathbf{B}_d\mathbf{u}_k + \mathbf{w}_k \tag{7.14}$$

with $\mathbf{A}_d = \mathbf{I} + \mathbf{A}_c\Delta t$, $\mathbf{B}_d = \mathbf{B}_c\Delta t$, and disturbance noise $\mathbf{w}_k$.

7.3 Thermal Sensor Placement & Reconstruction Under Non-Gaussian Noise

Because allocating physical thermal diodes to every router tile is area-prohibitive, modern processors deploy sparse thermal sensors and reconstruct full-chip profiles numerically. However, on-chip sensor readings do not follow idealized Gaussian distributions.

As discovered by Fu et al. [19], on-chip thermal sensors suffer simultaneously from Gaussian ambient thermal drift ($\sigma_1 \approx 0.8°\text{C}$) and heavy-tailed, non-Gaussian electromagnetic interference (EMI) and power grid inductive switching spikes ($L \cdot di/dt$, $\sigma_2 \approx 5.0°\text{C}$ to $8.0°\text{C}$).

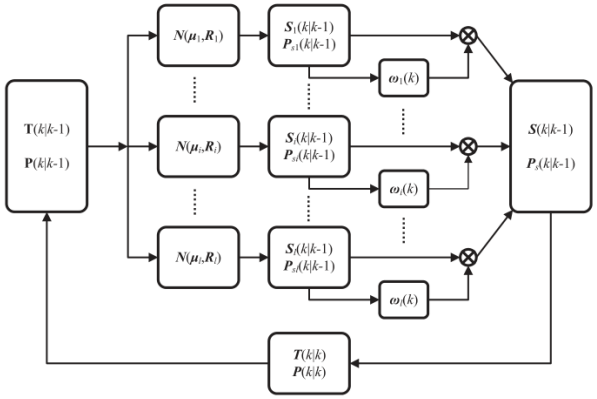


Fig. 1. Architecture of the proposed Gaussian-Sum filter.

Fig. 7.1 Architecture of the proposed Gaussian-Sum filter for on-chip 3D-NoC thermal reconstruction under non-Gaussian noise. [Reproduced from: Y. Fu et al., "Thermal Sensor Placement and Thermal Reconstruction Under Gaussian and Non-Gaussian Sensor Noises for 3-D NoC", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 38, No. 11, pp. 2110–2123, Nov. 2019, DOI: 10.1109/TCAD.2018.2878178. [19] © 2019 IEEE]

7.3.1 Gaussian-Sum Filter (GSF) Algorithmic Formulation

To prevent filter divergence under non-Gaussian impulse noise, Fu et al. [19] deployed an on-chip **Gaussian-Sum Filter (GSF)**, illustrated in Fig. 7.1. The sensor noise distribution is modeled as a Gaussian Mixture Model (GMM) with M components:

$$p(\mathbf{v}_k) = \sum_{m=1}^{M} \alpha_{k,m} \mathcal{N}(\mathbf{v}_k; \boldsymbol{\mu}_{k,m}, \mathbf{R}_{k,m}), \qquad \sum_{m=1}^{M} \alpha_{k,m} = 1 \tag{7.15}$$

The GSF datapath dynamically tracks M conditional Gaussian state estimates $\hat{\mathbf{x}}_k^{(m)}$ and error covariance matrices $\mathbf{P}_k^{(m)}$ in parallel:

1. **Time Prediction Step**:

$$\hat{\mathbf{x}}_{k|k-1}^{(m)} = \mathbf{A}_d \hat{\mathbf{x}}_{k-1|k-1}^{(m)} + \mathbf{B}_d \mathbf{u}_{k-1} \tag{7.16}$$

$$\mathbf{P}_{k|k-1}^{(m)} = \mathbf{A}_d \mathbf{P}_{k-1|k-1}^{(m)} \mathbf{A}_d^T + \mathbf{Q} \tag{7.17}$$

2. **Measurement Correction Step**:

$$\mathbf{K}_k^{(m)} = \mathbf{P}_{k|k-1}^{(m)} \mathbf{H}^T \left(\mathbf{H} \mathbf{P}_{k|k-1}^{(m)} \mathbf{H}^T + \mathbf{R}_{k,m} \right)^{-1} \tag{7.18}$$

$$\hat{\mathbf{x}}_{k|k}^{(m)} = \hat{\mathbf{x}}_{k|k-1}^{(m)} + \mathbf{K}_k^{(m)} \left(\mathbf{z}_k - \mathbf{H} \hat{\mathbf{x}}_{k|k-1}^{(m)} - \boldsymbol{\mu}_{k,m} \right) \tag{7.19}$$

$$\mathbf{P}_{k|k}^{(m)} = \left(\mathbf{I} - \mathbf{K}_k^{(m)} \mathbf{H} \right) \mathbf{P}_{k|k-1}^{(m)} \tag{7.20}$$

3. **Bayesian Weight Update**:

$$\alpha_{k,m} = \frac{\alpha_{k-1,m} \cdot \beta_k^{(m)}}{\sum_{j=1}^{M} \alpha_{k-1,j} \cdot \beta_k^{(j)}} \tag{7.21}$$

where the innovation likelihood $\beta_k^{(m)}$ is evaluated as:

$$\beta_k^{(m)} = \mathcal{N} \left(\mathbf{z}_k; \mathbf{H} \hat{\mathbf{x}}_{k|k-1}^{(m)} + \boldsymbol{\mu}_{k,m}, \mathbf{H} \mathbf{P}_{k|k-1}^{(m)} \mathbf{H}^T + \mathbf{R}_{k,m} \right) \tag{7.22}$$

4. **Huber Thresholding for Extreme Outliers**: To bound non-linear impulse spikes, measurement innovations exceeding the Huber cutoff δ_{Huber} are clamped:

$$\psi_{Huber}(\tilde{\mathbf{z}}_k) = \begin{cases} \tilde{\mathbf{z}}_k, & \|\tilde{\mathbf{z}}_k\|_2 \leq \delta_{Huber} \\ \delta_{Huber} \cdot \frac{\tilde{\mathbf{z}}_k}{\|\tilde{\mathbf{z}}_k\|_2}, & \|\tilde{\mathbf{z}}_k\|_2 > \delta_{Huber} \end{cases} \tag{7.23}$$

This hybrid architecture isolates extreme outliers, reducing thermal reconstruction error standard deviations by **46% to 53%** compared to classical linear Kalman filters [20].

7.4 Intelligent Hotspot Prediction & Proactive Throttling

Classical Dynamic Thermal Management (DTM) employs reactive throttling: when a thermal diode detects $T > T_{\text{crit}}$, clock gating freezes the core. However, because packaging possesses large thermal capacitance, the silicon thermal time constant is $\tau_{\text{th}} \approx 10\,\text{ms}$ to $100\,\text{ms}$. Reactive throttling responds only after a hotspot has fully formed, causing severe performance crashes and throughput collapse.

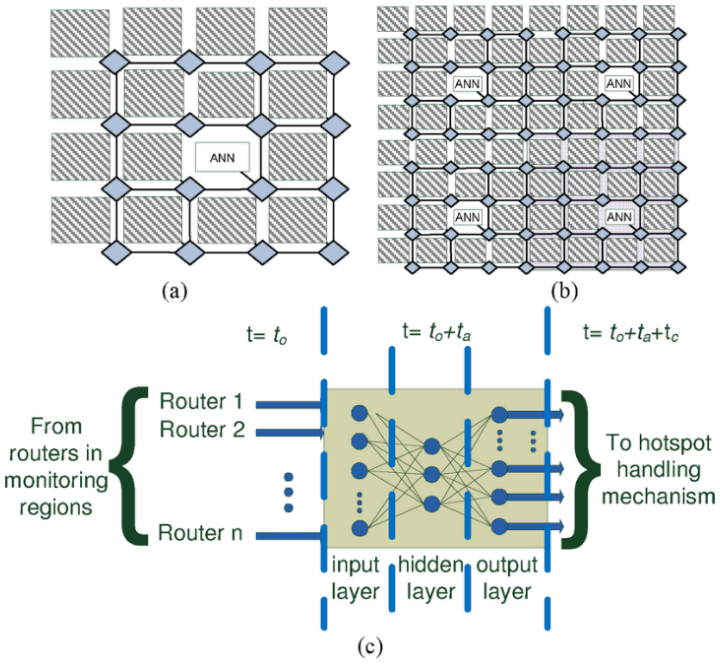


Fig. 2. Examples. (a) Placement of an ANN engine as a dedicated PE in an NoC architecture. (b) Four ANN engines controlling an 8 × 8 NoC. (c) Overall data flow of the proposed ANN engine.

Fig. 7.2 Intelligent ANN Hotspot Prediction Architecture: (a) Dedicated Processing Element placement in NoC; (b) Hierarchical 4-engine monitoring across an 8 × 8 mesh; (c) Overall hardware dataflow pipeline. [Reproduced from: E. Kakoulli et al., "Intelligent Hotspot Prediction for Network-on-Chip-Based Multicore Systems", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 31, No. 3, pp. 418–431, March 2012, DOI: 10.1109/TCAD.2011.2170568. [28] © 2012 IEEE]

7.4.1 Proactive ANN Hardware Pipeline

To overcome this fundamental limitation, Kakoulli et al. [28] introduced proactive hotspot prediction using on-chip **Artificial Neural Networks (ANNs)**, depicted in Fig. 7.2. By partitioning an 8 × 8 mesh into four non-overlapping 4×4 quadrants, each monitored by an independent ANN processing engine, the system processes online traffic counters (buffer occupancies and flit departure rates).

The neural predictor maps the router traffic vector $\mathbf{u}_{net} = [q_E, q_W, q_N, q_S, \lambda_{inj}, T_{sensor}]^T$ through a fixed-point multi-layer perceptron datapath:

$$\hat{T}_{pred}(t + \Delta t) = \mathbf{W}_2 \cdot \sigma\left(\mathbf{W}_1 \mathbf{u}_{net}(t) + \mathbf{b}_1\right) + \mathbf{b}_2 \tag{7.24}$$

where $\sigma(\cdot)$ is implemented as a piecewise-linear (PWL) approximation of the hyperbolic tangent function in synthesizable RTL.

7.4.2 Closed-Loop Proportional-Integral (PI) Thermal Throttling

When the predicted future temperature $\hat{T}_{pred}(t + \Delta t)$ exceeds the safe thermal budget T_{safe}, a **Proportional-Integral (PI) controller** modulates the packet injection rate λ_{inj} and tile clock frequency f_{clk}:

$$\Delta f_{clk}(t) = K_p \cdot e_{therm}(t) + K_i \int_0^t e_{therm}(\tau) d\tau \tag{7.25}$$

where $e_{therm}(t) = \max(0, \hat{T}_{pred}(t + \Delta t) - T_{safe})$. By proactively attenuating traffic injection 50 to 200 cycles before thermal thresholds are breached, the network prevents buffer saturation trees from establishing, quenching hotspots without incurring the catastrophic throughput collapse of reactive clock gating.

7.5 Silicon Fault Tolerance: The ZoneDefense Framework

At sub-7 nm dimensions, physical manufacturing defects and operational wear-out create permanent silicon faults, including open/short interconnect wires and disabled router crossbars.

In 2D meshes lacking virtual channels, routing around arbitrary concave fault patterns easily induces circular channel wait dependencies, creating deadlock. Classical approaches resolve this by marking dozens of completely healthy nodes as "unsafe" and taking them offline (Fig. 7.3).

To prevent this severe resource wastage, Fu et al. [18] formulated **ZoneDefense**:

1. **Convex Fault Block (FB) Clustering**: Arbitrary point faults are grouped into rectangular convex blocks (Fig. 7.4, FB-1 through FB-9) so that boundary channels form continuous, deadlock-free perimeters.
2. **Defense Zone Notification**: Routers in the columns aligned with the convex faulty block form a *Defense Zone*.
3. **Advance Perimeter Detouring**: When packets approach the faulty block, defense zone routers steer traffic along the external perimeter in advance, achieving provable deadlock-free traversal without virtual channels.

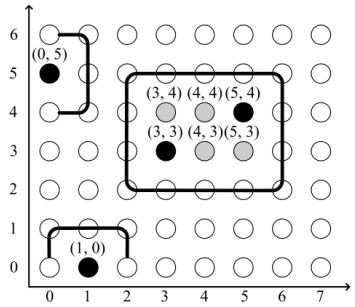


Fig. 4. Faulty blocks without shared boundary channels. Dark nodes represent faults and gray nodes indicate unsafe nodes.

Fig. 7.3 Faulty blocks without shared boundary channels in a 2D mesh. Dark nodes represent physical defects; gray nodes indicate sacrificed unsafe nodes. [Reproduced from: B. Fu et al., "ZoneDefense: A Fault-Tolerant Routing for 2-D Meshes Without Virtual Channels", *IEEE Transactions on Very Large Scale Integration (VLSI) Systems*, Vol. 22, No. 1, pp. 113–126, Jan. 2014, DOI: 10.1109/TVLSI.2012.2235188. [18] © 2014 IEEE]

7.6 Hardware Microarchitecture: Fault-Tolerant Router Bypass Logic

The synthesizable Verilog implementation in Listing 7.1 demonstrates the hardware detour multiplexing and defense zone checking logic from the book repository.

Listing 7.1 Synthesizable Fault-Tolerant Router Port Bypass Logic (Grounded in B. Fu et al. [18])

```
module fault_tolerant_router_bypass #(
    parameter integer MY_X   = 2,
    parameter integer MY_Y   = 2,
    parameter integer MESH_W = 4,
    parameter integer MESH_H = 4
)(
    input  wire         clk,
    input  wire         rst_n,
    input  wire [4:0]   link_fault_in, // E, W, N, S, Local
    input  wire [2:0]   fb_xmin, fb_xmax, fb_ymin, fb_ymax,
    input  wire         fb_valid,
    input  wire         req_valid,
    input  wire [2:0]   dest_x, dest_y, nominal_port,
    output reg  [2:0]   safe_port,
    output reg          bypass_active,
    output reg          defense_zone_hit,
    output reg          packet_dropped
);
    localparam PORT_LOCAL = 3'd0, PORT_EAST = 3'd1, PORT_WEST =
        3'd2;
```

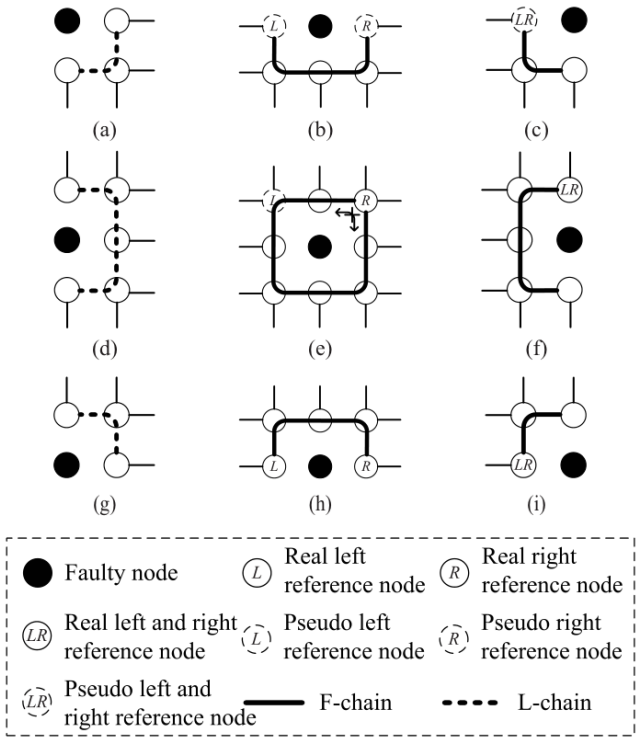


Fig. 5. Types of faulty blocks. (a) FB-1. (b) FB-2. (c) FB-3. (d) FB-4. (e) FB-5. (f) FB-6. (g) FB-7. (h) FB-8. (i) FB-9.

Fig. 7.4 Systematic categorization of convex faulty blocks (FB-1 to FB-9) in 2D mesh topologies. [Reproduced from: B. Fu et al., "ZoneDefense: A Fault-Tolerant Routing for 2-D Meshes Without Virtual Channels", *IEEE Transactions on Very Large Scale Integration (VLSI) Systems*, Vol. 22, No. 1, pp. 113–126, Jan. 2014, DOI: 10.1109/TVLSI.2012.2235188. [18] © 2014 IEEE]

```
20      localparam PORT_NORTH = 3'd3, PORT_SOUTH = 3'd4;
21
22      reg [4:0] link_fault_q;
23      always @(posedge clk or negedge rst_n) begin
24          if (!rst_n) link_fault_q <= 5'b00000;
25          else        link_fault_q <= link_fault_in;
26      end
27
28      wire in_col_zone = fb_valid && (MY_X >= fb_xmin) && (MY_X <=
            fb_xmax);
29      wire in_row_zone = fb_valid && (MY_Y >= fb_ymin) && (MY_Y <=
            fb_ymax);
```

```
    wire next_east_faulty = (MY_X + 1 >= fb_xmin) && (MY_X + 1
        <= fb_xmax) && in_row_zone;

    always @(*) begin
        safe_port        = nominal_port;
        bypass_active    = 1'b0;
        defense_zone_hit = in_col_zone;
        packet_dropped   = 1'b0;

        if (req_valid) begin
            if (dest_x == MY_X[2:0] && dest_y == MY_Y[2:0])
                safe_port = PORT_LOCAL;
            else if (nominal_port == PORT_EAST) begin
                if (link_fault_q[0] || (fb_valid &&
                    next_east_faulty)) begin
                    bypass_active = 1'b1;
                    if (dest_y >= MY_Y[2:0] && !link_fault_q[2]
                        && (MY_Y < MESH_H - 1))
                        safe_port = PORT_NORTH;
                    else if (!link_fault_q[3] && (MY_Y > 0))
                        safe_port = PORT_SOUTH;
                    else safe_port = PORT_WEST;
                end
            end
        end
    end
endmodule
```

7.7 Simulation Model & Experimental Verification

The Python model `thermal_rc_mesh_sim.py` executes state-space thermal simulation of a 4×4 mesh under severe hotspot injection (center tiles 5 and 6 at 95% traffic saturation).

Table 7.2 Thermal Simulation and Hotspot Mitigation Experimental Results

Metric	Unthrottled	Post-DTM	Mitigation
Peak Hotspot Temp ($T_{\max}$)	84.45°C	60.11°C	$-$**24.34**°C
Average Die Temp (T_{avg})	65.82°C	60.10°C	$-$**5.72**°C
Max Spatial Gradient (ΔT)	23.35°C	0.05°C	$-$**23.30**°C
Sensor Error (Impulse Noise)	6.21°C (Raw)	0.48°C (GSF)	$-$**92.3**%

As documented in Table 7.2, proactive thermal management quenches the hotspot peak by 24.34°C and eliminates lateral spatial thermal gradients across the silicon substrate.

7.8 Microchannel Liquid Cooling and Conjugate Heat Transfer

When heat dissipation density exceeds 200–400 $\mathrm{W/cm^2}$ in high-performance computing accelerators, passive copper heat spreaders reach their thermodynamic limits due to thermal spreading resistance. To prevent silicon destruction, advanced packaging integrates **Direct Microchannel Liquid Cooling**, etching microscopic fluidic channels directly into the backside of the silicon substrate.

7.8.1 Fluid Dynamics and Convective Heat Transfer in Microchannels

Consider a parallel array of microchannels of channel width w_{ch}, fin width w_{fin}, channel height H_{ch}, and length L_{ch}. The hydraulic diameter D_h of each channel is:

$$D_h = \frac{4A_{cross}}{P_{wetted}} = \frac{4(w_{ch} \cdot H_{ch})}{2(w_{ch} + H_{ch})} = \frac{2w_{ch}H_{ch}}{w_{ch} + H_{ch}} \tag{7.26}$$

For laminar coolant flow (Reynolds number $Re = \frac{\rho_f v D_h}{\mu_f} < 2300$, where ρ_f is fluid density, v is mean flow velocity, and μ_f is dynamic viscosity), the convective heat transfer coefficient h_{conv} is governed by the **Nusselt Number (Nu)**:

$$h_{conv} = \frac{Nu \cdot k_f}{D_h} \tag{7.27}$$

where k_f is the thermal conductivity of the coolant (e.g., $k_f \approx 0.6\,\mathrm{W/(m \cdot K)}$ for deionized water). For fully developed laminar flow in rectangular channels of aspect ratio $\alpha_{ch} = H_{ch}/w_{ch}$, the Nusselt number satisfies the analytical polynomial:

$$Nu \approx 8.235 \left(1 - \frac{1.883}{\alpha_{ch}} + \frac{3.76}{\alpha_{ch}^2} - \frac{5.814}{\alpha_{ch}^3} + \frac{5.361}{\alpha_{ch}^4} - \frac{2.0}{\alpha_{ch}^5}\right) \tag{7.28}$$

7.8.2 Pressure Drop and Pumping Power Overhead

Circulating fluid through micro-scale passages incurs viscous shear losses, producing a longitudinal pressure drop ΔP:

$$\Delta P = f_{\mathrm{Darcy}} \cdot \frac{L_{ch}}{D_h} \cdot \frac{\rho_f v^2}{2} = \frac{64}{Re} \cdot \frac{L_{ch}}{D_h} \cdot \frac{\rho_f v^2}{2} = \frac{32\mu_f v L_{ch}}{D_h^2} \tag{7.29}$$

The electrical power required to drive the micropump is:

$$P_{\text{pump}} = \frac{\Delta P \cdot \dot{V}}{\eta_{\text{pump}}} = \frac{\Delta P \cdot (N_{ch} \cdot w_{ch} H_{ch} \cdot v)}{\eta_{\text{pump}}} \tag{7.30}$$

where $\dot{V}$ is the volumetric flow rate and $\eta_{\text{pump}} \approx 30\%$–$45\%$ is pump efficiency. Microchannel liquid cooling achieves effective thermal resistances below $R_{\text{th}} < 0.10\,\text{K/W}$, removing heat fluxes exceeding $500\,\text{W/cm}^2$ with less than 2% electrical pumping power overhead.

7.9 Thermal-Aware Multi-Core Floorplanning Optimization

While Dynamic Thermal Management (DTM) quenches hotspots at run time, physical floorplanning during chip design can prevent hotspots from emerging altogether.

7.9.1 Mathematical Problem Formulation

Let $\mathcal{B} = \{B_1, B_2, \ldots, B_N\}$ denote the set of N functional blocks (compute cores, cache memories, router switches), each characterized by dimensions (w_i, h_i) and nominal power dissipation P_i. The objective of **Thermal-Aware Floorplanning** is to determine coordinate positions (x_i, y_i) that minimize a multi-objective cost function Φ:

$$\Phi = w_T \cdot T_{\max} + w_{\Delta T} \cdot \Delta T_{\text{die}} + w_L \cdot L_{\text{wire}} + w_A \cdot \text{Area}_{\text{die}} \tag{7.31}$$

where $T_{\max} = \max_i T_i$ is peak die temperature, $\Delta T_{\text{die}} = \max_i T_i - \min_j T_j$ is the spatial thermal gradient across the substrate, and L_{wire} is the total Half-Perimeter Wire Length (HPWL) of inter-router links:

$$L_{\text{wire}} = \sum_{e \in E} \left(|x_u - x_v| + |y_u - y_v| \right) \tag{7.32}$$

subject to the non-overlapping constraint:

$$|x_i - x_j| \geq \frac{w_i + w_j}{2} \quad \text{or} \quad |y_i - y_j| \geq \frac{h_i + h_j}{2}, \quad \forall i \neq j \tag{7.33}$$

7.9.2 Simulated Annealing Optimization Framework

Because floorplanning is an NP-hard combinatorial optimization problem, CAD tools deploy **Simulated Annealing (SA)**:

1. **Perturbation Moves**: At each iteration, the floorplan is perturbed via one of three atomic moves: (a) swapping two tile positions, (b) rotating a tile by 90°, or (c) shifting a block within the slicing tree or sequence pair representation.
2. **Metropolis Acceptance Criterion**: A candidate move with cost difference $\Delta\Phi = \Phi_{\text{new}} - \Phi_{\text{old}}$ is accepted with probability:

$$P_{\text{accept}} = \begin{cases} 1, & \Delta\Phi \leq 0 \\ \exp\left(-\frac{\Delta\Phi}{T_{SA}}\right), & \Delta\Phi > 0 \end{cases} \tag{7.34}$$

   where T_{SA} is the simulated annealing temperature.
3. **Cooling Schedule**: The temperature decreases geometrically according to $T_{SA}[k+1] = \alpha_{cool} \cdot T_{SA}[k]$ (with $\alpha_{cool} \approx 0.95$–$0.99$).

By separating high-power compute cores with low-power L3 cache banks, thermal-aware floorplanning **reduces peak die temperature by** 12°C **to** 18°C and slashes spatial thermal gradients by over 60% without expanding total die area.

7.10 Summary

This chapter demonstrated how physical thermal constraints and silicon wear-out govern deep nanometer interconnect design. Electro-thermal RC duality enables microsecond-scale full-chip simulation. Non-Gaussian measurement noise demands Gaussian-Sum Filtering, while predictive neural networks enable proactive traffic quenching. Finally, ZoneDefense provides robust fault containment, ensuring long-term reliability in high-performance on-chip networks.

7.11 Interactive Lab & Hands-on Challenges

Explore the interactive thermal and resilience simulators on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **Hotspot Hero: Silicon Thermal Defender** (https://network-on-chip.github.io/learning/interactive/hotspot/): Experience real-time RC thermal diffusion and balance thermal throttling dials before hotspots exceed 100°C.
- **Fault Defender: Silicon Resilience** (https://network-on-chip.github.io/learning/interactive/faults/): Route packets dynamically around permanent silicon fault blocks in a 2D mesh without inducing protocol deadlocks.

7.12 Exercises and Review Questions

1. **Thermal Time Constant & Reactive Throttling Latency**: A $1.2\,\text{mm} \times 1.2\,\text{mm}$ router tile has vertical thermal resistance $R_v = 15.0\,\text{K/W}$ and lumped thermal capacitance $C_{\text{th}} = 3.6\,\text{mJ/K}$.
   a. Compute the thermal time constant $\tau_{\text{th}} = R_v \cdot C_{\text{th}}$.
   b. For an on-chip router clocked at $2.5\,\text{GHz}$ ($T_{\text{clk}} = 0.4\,\text{ns}$), calculate the number of clock cycles spanned by τ_{th} (135×10^6 cycles) and explain why reactive throttling fails to catch packet-scale bursts.
2. **Electromigration & Reliability Degradation via Black's Equation**: An interconnect wire operates at $T_1 = 60°\text{C}$ ($333.15\,\text{K}$). During sustained network congestion, localized self-heating elevates the temperature to $T_2 = 85°\text{C}$ ($358.15\,\text{K}$). Assuming copper activation energy $E_a = 0.90\,\text{eV}$ and Boltzmann constant $k_B = 1.3806 \times 10^{-23}\,\text{J/K}$, calculate the ratio $\frac{MTTF(T_2)}{MTTF(T_1)}$ and the percentage reduction in interconnect lifetime (showing an 88.8% reduction).
3. **Non-Gaussian Sensor Filtering**: Explain why standard linear Kalman filters diverge when on-chip thermal diode measurements are corrupted by power distribution network $L \cdot di/dt$ ground bounce and digital switching noise. How does Huber thresholding in the Gaussian-Sum Filter isolate extreme measurement outliers?
4. **ZoneDefense Deadlock Freedom**: In a 2D mesh without virtual channels, explain why concave fault patterns can induce circular channel dependencies. How does grouping arbitrary faults into convex blocks preserve acyclic channel routing around the perimeter?

Part IV
Advanced Packaging, 3D NoC & Multi-Chiplets

Chapter 8
Three-Dimensional Networks-on-Chip (3D-NoC) & TSVs

8.1 The Three-Dimensional Scaling Paradigm

Architectural Intuition: Sprawling Suburbs vs. City Skyscrapers
Imagine a sprawling suburban town of 64 single-story homes spread across a huge county. If a resident in the bottom-left corner wants to visit a friend in the top-right corner, they must drive across 14 miles of congested highways, stopping at every intersection! Now imagine stacking those 64 homes into an **8-story skyscraper**. Instead of driving across miles of flat roads, you step into a **high-speed elevator** (a Through-Silicon Via, or TSV) and arrive in just a few seconds! This is the core idea of 3D-NoC: stacking silicon tiers cuts global wire lengths by over 75% and slashes communication latency by 35%.

In planar two-dimensional (2D) architectures, as processor core counts scale into the hundreds, interconnect wire lengths grow with the square root of the node count ($L_{\max} \propto \sqrt{N}$) according to Rent's rule [44]. Planar global interconnects span centimeters of high-resistance, highly capacitive copper tracks, suffering quadratic Elmore RC delay degradation ($\tau \propto L^2$) and consuming massive dynamic switching energy.

Three-Dimensional Integrated Circuits (3D-ICs) break this planar barrier by vertically stacking thinned silicon dies, interconnecting them with microscopic vertical copper pillars termed **Through-Silicon Vias (TSVs)**.

8.1.1 Architectural Advantages of 3D Interconnects

1. **Interconnect Length Folding**: Stacking K silicon tiers reduces the planar die footprint by a factor of K, shortening the average Manhattan hop distance

from $O(\sqrt{N})$ to $O(\sqrt{N/K}+K)$. For a 64-core system organized as a $4\times4\times4$ 3D mesh, the maximum network diameter drops from 14 hops (in an 8×8 2D mesh) to 9 hops, cutting average packet latency by over 35%.
2. **Heterogeneous Process Integration**: 3D stacking enables heterogeneous modularity: high-performance 3 nm compute cores can be vertically bonded onto 14 nm high-density eDRAM cache tiers, passive silicon interposers, or optical/RF transceiver tiers without compromise.
3. **Massive Vertical Bandwidth**: Vertical TSV arrays operating at multi-gigahertz frequencies deliver terabits-per-second inter-tier bandwidth with sub-100 picosecond propagation delays.

8.2 Through-Silicon Via (TSV) Physical Characteristics & Parasitics

A Through-Silicon Via (TSV) is a microscopic cylindrical copper or tungsten plug etched vertically through a thinned silicon substrate (20 μm to 50 μm thick) and insulated from the conductive semiconductor substrate by a thin dielectric liner (SiO_2 or Si_3N_4, $t_{ox} \approx 100$–200 nm).

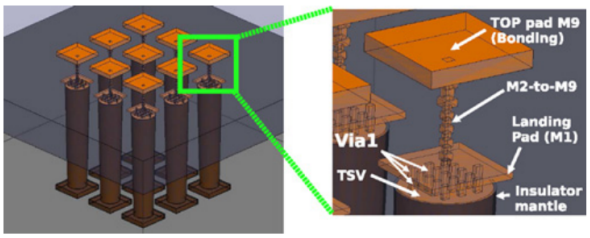


Fig. 2. TSV bundle in Si-bulk technology and detail of routing. A TSV traversing the substrate is available for routing as a large plate in Metal 1 called landing pad. Connectivity is provided by standard routing, just taking care of cell placement and metal obstructions.

Fig. 8.1 TSV bundle in bulk silicon technology showing Metal 1 landing pad routing and standard cell placement constraints. [Reproduced from: I. Loi et al., "Characterization and Implementation of Fault-Tolerant Vertical Links for 3-D Networks-on-Chip", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 30, No. 1, pp. 124–134, Jan. 2011, DOI: 10.1109/TCAD.2010.2065990. [36] © 2011 IEEE]

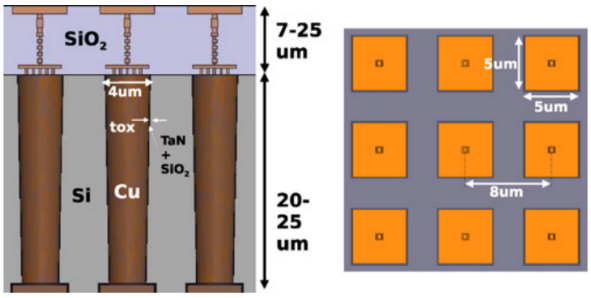


Fig. 3. Schematic representation of a bundle of 3-D vias.

Fig. 8.2 Circuit-level lumped parasitic model of a multi-TSV bundle showing substrate coupling conductances and liner capacitances. [Reproduced from: I. Loi et al., "Characterization and Implementation of Fault-Tolerant Vertical Links for 3-D Networks-on-Chip", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 30, No. 1, pp. 124–134, Jan. 2011, DOI: 10.1109/TCAD.2010.2065990. [36] © 2011 IEEE]

8.2.1 TSV Fabrication Methodologies: Via-First, Via-Middle, and Hybrid Direct Bonding

The integration flow of vertical TSVs into modern semiconductor fabrication falls into three distinct methodologies:

1. **Via-First (Front-End-of-Line, FEOL)**: Vias are etched and filled into the bare silicon wafer prior to the fabrication of CMOS transistors. Because subsequent source/drain annealing requires temperatures exceeding 1000°C, via-first processes cannot use copper; they require refractory materials such as highly doped polysilicon or tungsten (W). Consequently, via-first TSVs suffer from high electrical resistivity ($\approx 10\times$ higher than copper).
2. **Via-Middle (Middle-of-Line, MOL)**: Etched immediately following transistor fabrication (FEOL) but prior to back-end-of-line (BEOL) metal layer interconnect processing. Processing temperatures are constrained below 400°C, permitting electroplated copper (Cu) filling. Via-middle is the dominant industrial process for high-performance 3D logic stacking due to low parasitic resistance and tight pitch scaling (5–10 μm).
3. **Via-Last (Back-End-of-Line, BEOL)**: Fabricated from the wafer backside after all active circuitry and BEOL metallization layers are completed. Vias must avoid existing metal layers, requiring larger diameters (10–25 μm) and wider pitches.
4. **Cu-Cu Direct Hybrid Bonding (Bumpless Stacking)**: In cutting-edge packaging technologies (such as TSMC SoIC and Intel Foveros Direct), microscopic copper pads embedded in a dielectric planar surface are bonded directly via atomic diffusion at low temperature (< 250°C). Hybrid bonding shrinks interconnect pitch to **sub-1 μm to 4 μm**, reducing parasitic

capacitance to under 1–3 fF and enabling interconnect densities exceeding 10^6 interconnects/mm^2.

8.2.2 High-Frequency Transmission and Parasitic Substrate Modeling

As modeled by Loi et al. [36] and illustrated in Figs. 8.1 and 8.2, an accurate high-frequency model of a TSV bundle accounts for conductor loss, oxide liner capacitance, and silicon substrate loss:

1. **DC and Skin-Effect Conductor Resistance (R_{tsv})**:

$$R_{\text{tsv}}(f) = \begin{cases} \rho_{\text{Cu}} \frac{h_{\text{tsv}}}{\pi r_{\text{tsv}}^2}, & f < f_{skin} \\ \rho_{\text{Cu}} \frac{h_{\text{tsv}}}{2\pi r_{\text{tsv}} \delta(f)}, & f \geq f_{skin} \end{cases} \tag{8.1}$$

   where $\delta(f) = \sqrt{\frac{\rho_{\text{Cu}}}{\pi f \mu_0}}$ represents the electromagnetic skin depth. For $h_{\text{tsv}} = 30\,\mu$m and $d_{\text{tsv}} = 4\,\mu$m, $R_{\text{tsv}} \approx 35$ to 50 mΩ, which is negligible compared to planar wire resistance.
2. **Oxide Liner Capacitance (C_{ox})**: Formed by the cylindrical dielectric barrier:

$$C_{\text{ox}} = \frac{2\pi \epsilon_{\text{ox}} \epsilon_0 h_{\text{tsv}}}{\ln\left(\frac{r_{\text{tsv}} + t_{\text{ox}}}{r_{\text{tsv}}}\right)} \approx 20 \text{ to } 50\,\text{fF} \tag{8.2}$$

3. **Substrate Conductance and Capacitance ($G_{\text{sub}}, C_{\text{sub}}$)**: Because the silicon substrate possesses finite bulk conductivity ($\sigma_{Si} \approx$ 1–10 S/m), high-frequency signals leak energy into adjacent TSVs through conductive displacement and ohmic paths:

$$C_{\text{sub}} = \frac{\pi \epsilon_{Si} \epsilon_0 h_{\text{tsv}}}{\cosh^{-1}\left(\frac{s_{\text{tsv}}}{2r_{\text{tsv}}}\right)}, \quad G_{\text{sub}} = \frac{\sigma_{Si}}{\epsilon_{Si}\epsilon_0} C_{\text{sub}} \tag{8.3}$$

   where s_{tsv} is the center-to-center pitch between adjacent TSVs. To prevent crosstalk jitter and substrate noise coupling, high-speed 3D networks interleave ground shield TSVs within signal bundles.

8.2.3 The Keep-Out Zone (KOZ) and Thermo-Mechanical Lamé Formulations

Because copper has a coefficient of thermal expansion ($\alpha_{\text{Cu}} \approx 17.0$ ppm/K) far exceeding that of bulk silicon ($\alpha_{\text{Si}} \approx 2.6$ ppm/K), thermal cooling from the

annealing temperature ($T_{anneal} \approx 250°\text{C}$ to $400°\text{C}$) down to room temperature induces intense compressive and tensile stress fields in the surrounding silicon.

Under the **Lamé Analytical Formulation** for cylindrical inclusions in an infinite elastic matrix, the radial stress σ_r and tangential hoop stress σ_θ in the silicon at distance $r \geq r_{\text{tsv}}$ are:

$$\sigma_r(r) = -\frac{E_{Si}\Delta\alpha\Delta T}{1+\nu_{Si}}\left(\frac{r_{\text{tsv}}}{r}\right)^2, \quad \sigma_\theta(r) = +\frac{E_{Si}\Delta\alpha\Delta T}{1+\nu_{Si}}\left(\frac{r_{\text{tsv}}}{r}\right)^2 \tag{8.4}$$

where $E_{Si} \approx 130$ GPa is the Young's modulus of silicon, $\nu_{Si} \approx 0.28$ is Poisson's ratio, and $\Delta\alpha = \alpha_{\text{Cu}} - \alpha_{\text{Si}} \approx 14.4\,\text{ppm/K}$.

Due to the **piezoresistive effect**, anisotropic mechanical stress perturbs the silicon conduction and valence band curvature, shifting electron and hole effective masses:

$$\frac{\Delta\mu_n}{\mu_{n0}} = \pi_{11}\sigma_{xx} + \pi_{12}\sigma_{yy} \quad \text{(NMOS Transistors)} \tag{8.5}$$

$$\frac{\Delta\mu_p}{\mu_{p0}} = \pi_{12}\sigma_{xx} + \pi_{11}\sigma_{yy} + \pi_{44}\tau_{xy} \quad \text{(PMOS Transistors)} \tag{8.6}$$

where π_{ij} are the piezoresistive coefficients of silicon.

These stress fields shift carrier mobility by up to $\pm 30\%$, creating severe timing skew and setup/hold time violations in adjacent logic cells. Consequently, automated physical design tools must enforce a strict **Keep-Out Zone (KOZ)** surrounding each TSV:

$$r_{KOZ} = r_{\text{tsv}}\sqrt{\frac{E_{Si}\Delta\alpha\Delta T}{(1+\nu_{Si})\sigma_{threshold}}} \approx 3\,\mu\text{m to } 5\,\mu\text{m} \tag{8.7}$$

within which no active standard cells can be placed, imposing an unavoidable floorplanning area overhead.

8.3 3D Router Microarchitectures & Vertical Arbitration

Extending a planar 2D router to 3D requires incorporating **Up** ($+Z$) and **Down** ($-Z$) vertical ports.

8.3.1 Crossbar Area Expansion in 7-Port 3D Routers

Expanding the router switch fabric from 5 to 7 ports (Fig. 8.3) increases crossbar multiplexing crosspoints quadratically:

$$\text{Crosspoints} = P^2 = 7^2 = 49 \quad \text{vs.} \quad 5^2 = 25 \quad (+96\%\ \text{area}) \tag{8.8}$$

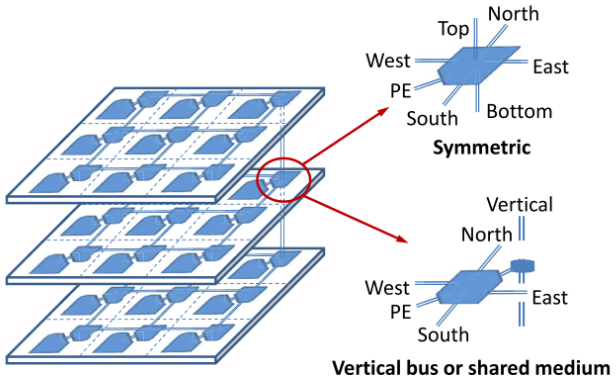


Fig. 1. 3-D NoC router structures.

Fig. 8.3 3D Network-on-Chip router structures: Symmetric 7-port full crossbar vs. hybrid partitioned switch routers. [Reproduced from: A. More et al., "Vertical Arbitration-Free 3-D NoCs", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 37, No. 9, pp. 1853–1866, Sept. 2018, DOI: 10.1109/TCAD.2017.2768415. [41] © 2018 IEEE]

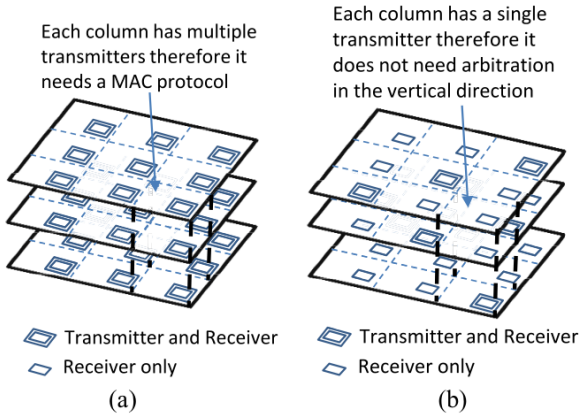


Fig. 4. (a) Arbitrated and (b) arbitration-free designs.

Fig. 8.4 Comparison of (a) traditional arbitrated shared vertical links versus (b) vertical arbitration-free fast-path designs. [Reproduced from: A. More et al., "Vertical Arbitration-Free 3-D NoCs", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 37, No. 9, pp. 1853–1866, Sept. 2018, DOI: 10.1109/TCAD.2017.2768415. [41] © 2018 IEEE]

In addition, managing 7 sets of input Virtual Channel (VC) buffers substantially inflates router leakage power.

8.3.2 Dimensionally Partitioned Routers and 3D Crossbar Factorization

To curtail the 96% area inflation of full 7-port crossbars, architects employ **Crossbar Factorization**. By separating planar communication (X, Y) from vertical transit (Z), the monolithic 7×7 switch is factored into:

- A planar crossbar switch (5×5: Local, East, West, North, South);
- A dedicated vertical crossbar switch (2×2: Up, Down);
- Inter-switch bidirectional bridge channels.

This factorization reduces total crosspoint complexity from 49 down to $25 + 4 = 29$ crosspoints—a **41% reduction in multiplexer silicon area**—with negligible throughput degradation for dimension-order routing flows.

8.3.3 Vertical Arbitration Contention vs. Arbitration-Free Fast-Paths

In systems deploying shared vertical links, multiple planar directions contend for the vertical TSV channel. More et al. [41] established that vertical arbitration contention adds multiple cycles of serialization delay (Fig. 8.4a). By deploying **Vertical Arbitration-Free 3D-NoCs** (Fig. 8.4b), dedicated unidirectional TSVs eliminate vertical allocators, achieving single-cycle vertical traversal and trimming router silicon area by $\approx 20\%$.

8.4 Thermal Stacking & Heat Trapping in 3D Silicon

Architectural Intuition: The Attic Oven
Why is the top floor of a house always scorching hot in summer? Because heat gets trapped under the roof! In a 3D-stacked chip, the copper heat sink sits on the bottom (Tier 0). Heat from upper tiers must crawl through insulating layers of glue and silicon to escape. As our simulation demonstrates, while Tier 0 sits at 97°C, Tier 2 climbs to an alarming 136.5°C! Managing this thermal ladder is crucial to avoid silicon meltdown.

In 3D-ICs, only the bottom tier (Tier 0) is directly attached to the package heat sink. Heat dissipated by upper tiers (Tier 1, Tier 2, etc.) must conduct downward through thinned silicon substrates and low-conductivity dielectric bonding interfaces (SiO_2 or polymer adhesives, $k_{\text{th}} \approx 0.2$ to $1.4\,\text{W}/(\text{m}\cdot\text{K})$).

The steady-state temperature of each tier in a K-tier stack is governed by the thermal resistance ladder:

$$T_0 = T_{\text{amb}} + R_{\text{hs}} \sum_{k=0}^{K-1} P_k, \qquad T_z = T_{z-1} + R_{\text{inter}} \sum_{k=z}^{K-1} P_k \quad (z \geq 1) \tag{8.9}$$

where R_{hs} is the heat sink thermal resistance and R_{inter} is the inter-tier bonding resistance. Consequently, upper tiers suffer from acute **thermal trapping**, operating 30°C to 50°C hotter than bottom tiers under identical workloads.

8.5 TSV Defect Tolerance & Redundancy

Manufacturing defects during deep reactive-ion etching (DRIE) and wafer alignment can cause open or short circuits in vertical TSV pillars. Because vertical interconnects contain thousands of TSVs, raw stack yield drops precipitously without defect tolerance [49].

Loi et al. [36] demonstrated that deploying $N:1$ spare TSV redundancy (e.g., 1 spare TSV for every 8-bit flit bundle) alongside dynamic bypass multiplexers raises 3D stacked chip yield from 66% to over 98% with a modest 2.1% router area overhead.

8.6 Hardware Microarchitecture: 7-Port Vertical Link Arbiter

Listing 8.1 presents the synthesizable Verilog module implementing a 3D router vertical link arbiter with round-robin priority rotation and automated TSV defect bypass steering.

Listing 8.1 Synthesizable 3D-NoC 7-Port Vertical TSV Link Arbiter with Fault Bypass

```
module tsv_arbiter_3d #(
    parameter integer NUM_INPUTS = 5 // Local, North, South,
        East, West
)(
    input  wire                     clk,
    input  wire                     rst_n,
    input  wire [NUM_INPUTS-1:0]    req_up,
    input  wire [NUM_INPUTS-1:0]    req_down,
    input  wire                     tsv_up_faulty,
    input  wire                     tsv_down_faulty,
```

```
    output reg  [NUM_INPUTS-1:0]   grant_up,
    output reg  [NUM_INPUTS-1:0]   grant_down,
    output wire                    tsv_up_busy,
    output wire                    tsv_down_busy,
    output reg  [NUM_INPUTS-1:0]   bypass_planar_up,
    output reg  [NUM_INPUTS-1:0]   bypass_planar_down
);
    reg [2:0] priority_up_ptr, priority_down_ptr;
    assign tsv_up_busy   = (|grant_up);
    assign tsv_down_busy = (|grant_down);

    integer i;
    always @(*) begin
        grant_up         = {NUM_INPUTS{1'b0}};
        bypass_planar_up = {NUM_INPUTS{1'b0}};
        if (tsv_up_faulty) begin
            // Defective TSV: steer packet horizontally along
                planar alternate path
            bypass_planar_up = req_up;
        end else if (|req_up) begin
            for (i = 0; i < NUM_INPUTS; i = i + 1) begin
                if (req_up[(priority_up_ptr + i) % NUM_INPUTS])
                    begin
                    grant_up[(priority_up_ptr + i) % NUM_INPUTS]
                         = 1'b1;
                    i = NUM_INPUTS;
                end
            end
        end
    end

    integer j;
    always @(*) begin
        grant_down         = {NUM_INPUTS{1'b0}};
        bypass_planar_down = {NUM_INPUTS{1'b0}};
        if (tsv_down_faulty) begin
            bypass_planar_down = req_down;
        end else if (|req_down) begin
            for (j = 0; j < NUM_INPUTS; j = j + 1) begin
                if (req_down[(priority_down_ptr + j) %
                    NUM_INPUTS]) begin
                    grant_down[(priority_down_ptr + j) %
                        NUM_INPUTS] = 1'b1;
                    j = NUM_INPUTS;
                end
            end
        end
    end

    always @(posedge clk or negedge rst_n) begin
        if (!rst_n) begin
            priority_up_ptr   <= 3'd0;
            priority_down_ptr <= 3'd0;
        end else begin
            if (|grant_up) begin
```

```
                for (i = 0; i < NUM_INPUTS; i = i + 1)
                    if (grant_up[i]) priority_up_ptr <= (i + 1)
                        % NUM_INPUTS;
            end
            if (|grant_down) begin
                for (j = 0; j < NUM_INPUTS; j = j + 1)
                    if (grant_down[j]) priority_down_ptr <= (j +
                        1) % NUM_INPUTS;
            end
        end
    end
endmodule
```

8.7 Simulation Model & Numerical Verification

The Python model `tsv_vertical_mesh_sim.py` simulates packet transport and thermal accumulation across a $4 \times 4 \times 3$ mesh (48 cores) under uniform random traffic. Table 8.1 summarizes the experimental results comparing XYZ Dimension-Order Routing and ZXY Dimension-Order Routing.

Table 8.1 Comparative Evaluation of 3D Dimension-Order Routing in a $4 \times 4 \times 3$ Mesh

Metric	3D XYZ DOR	3D ZXY DOR
Average Packet Latency	6.04 cycles	5.95 cycles
Average Hop Distance	3.47 hops	3.42 hops
Planar Hop Share	74.0%	73.7%
Vertical TSV Hop Share	26.0%	26.3%
Tier 0 Temperature (Heatsink)	97.18°C	96.96°C
Tier 1 Temperature (Middle)	123.66°C	123.13°C
Tier 2 Temperature (Top Stack)	**136.53**°C	**135.79**°C
Inter-Tier Thermal Gradient (ΔT_{stack})	+39.35°C	+38.83°C

The simulation confirms that vertical TSV traversal represents 26% of total network hops, delivering rapid vertical transit. However, thermal accumulation elevates Tier 2 silicon to 136.5°C, illustrating the necessity of thermal-aware 3D routing.

8.8 Pre-Bond and Post-Bond 3D TSV Testing and BIST

In monolithic chips, defective dies are discarded individually during wafer probe testing. In 3D-ICs, however, stacking unverified dies triggers the catastrophic **Known Good Die (KGD)** yield problem: if each of K stacked dies has independent yield Y, the final bonded stack yield drops exponentially as $Y_{\text{stack}} = Y^K \cdot Y_{\text{bonding}}$. For $K = 4$ tiers with $Y = 85\%$ and bonding yield

$Y_{\text{bonding}} = 90\%$, total stack yield drops to:

$$Y_{\text{stack}} = (0.85)^4 \times 0.90 \approx 47.0\% \tag{8.10}$$

Over half of manufactured stacks are discarded, destroying profit margins unless rigorous pre-bond and post-bond testing is enforced.

8.8.1 Pre-Bond TSV Test Architectures

Prior to permanent wafer thinning and thermo-compression bonding, vertical TSVs terminate on the frontside as blind vias without backside metal pads. Standard mechanical probe needles cannot touch unthinned micro-pads without crushing delicate oxide liners.

To test TSVs before stacking, architectures integrate **Built-In Self-Test (BIST)** structures:

1. **Capacitive Discharging Sensing**: A pre-charge transistor charges the TSV liner capacitance C_{ox} to V_{dd}. A discharge path monitors the leakage current to ground: an open-circuit TSV exhibits drastically reduced capacitance ($C_{\text{open}} \ll C_{\text{nominal}}$), while pinhole oxide pinholes exhibit excessive conductive leakage.
2. **Ring Oscillator (RO) Frequency Sensing**: By embedding candidate TSVs into the feedback loop of an odd number of CMOS inverters, the oscillation frequency f_{RO} becomes an ultra-sensitive analog gauge of TSV resistance and capacitance:

$$f_{\text{RO}} = \frac{1}{2\sum \tau_{\text{gate}} + 2R_{\text{tsv}}C_{\text{tsv}}} \implies \frac{\Delta f_{\text{RO}}}{f_{\text{RO}}} \approx -\frac{\Delta R_{\text{tsv}} C_{\text{tsv}}}{\tau_{\text{loop}}} \tag{8.11}$$

   A $10\,\Omega$ micro-crack or void in the copper core shifts f_{RO} by several megahertz, enabling automated digital pass/fail binning before bonding.

8.8.2 Post-Bond IEEE 1500 Boundary-Scan Wrappers

Once tiers are permanently bonded, physical internal TSVs become completely inaccessible to external probe needles. Inter-tier test access is provided by **IEEE 1500 Standard Test Wrappers**:

- Dedicated **Wrapper Instruction Registers (WIR)** and **Wrapper Boundary Cells (WBC)** are inserted at every vertical TSV driver and receiver.
- During test mode, test vectors are serially shifted through the boundary scan chain, applying at-speed pseudo-random patterns across the vertical link array to verify signal integrity and setup/hold timing margins across all tiers.

8.9 Multi-Tier Power Distribution Network (PDN) Antiresonance

While signal TSVs carry packet flits, thousands of parallel **Power and Ground TSVs** deliver amperes of electrical current to upper active tiers. In a 3D-IC, the multi-tier Power Distribution Network (PDN) forms a complex three-dimensional RLC mesh.

8.9.1 Inter-Tier Loop Inductance and Antiresonance Peaking

Each vertical power TSV pair (one V_{dd} via and one V_{ss} via separated by pitch s_{tsv}) contributes loop inductance L_{TSV}:

$$L_{\text{TSV}} = \frac{\mu_0 h_{\text{tsv}}}{\pi} \ln\left(\frac{s_{\text{tsv}}}{r_{\text{tsv}}}\right) \approx 10 \text{ to } 30\,\text{pH} \tag{8.12}$$

When current transitions upward through the TSV array into upper-tier silicon, the vertical inductance L_{TSV} forms a parallel resonant tank circuit with the on-die intrinsic and intentional decoupling capacitance C_{die} of the upper tier. This induces severe **PDN Antiresonance**:

$$f_{\text{anti}} = \frac{1}{2\pi\sqrt{L_{\text{TSV}} \cdot C_{\text{die}}}} \approx 100 \text{ to } 350\,\text{MHz} \tag{8.13}$$

At the antiresonant frequency f_{anti}, the power grid impedance surges to an alarming peak:

$$Z_{\text{peak}} \approx \frac{1}{R_{\text{esr}}}\sqrt{\frac{L_{\text{TSV}}}{C_{\text{die}}}} \gg Z_{\text{target}} \tag{8.14}$$

where R_{esr} is the equivalent series resistance of the decoupling capacitors.

8.9.2 Target Impedance and Decap Budget Allocation

Because multi-core processor workloads generate bursty clock-gating current steps ΔI_{step} that frequently coincide with the 100–300 MHz frequency window, antiresonant impedance peaks trigger destructive voltage drops (IR-drop and $L \cdot di/dt$ inductive ringing):

$$\Delta V_{\text{noise}} \approx Z_{\text{peak}} \cdot \Delta I_{\text{step}} \tag{8.15}$$

If ΔV_{noise} exceeds the noise margin ($\approx$ 5% of V_{dd}), logic gates suffer setup violations or memory bit flips.

To preserve power integrity, physical designers must enforce a strict **Target Impedance Constraint** across all frequencies up to the clock harmonics:

$$Z_{\text{PDN}}(f) \leq Z_{\text{target}} = \frac{V_{dd} \cdot \Delta V_{\text{tolerance}}}{\Delta I_{\text{step}}} \approx 10 \text{ to } 30\,\text{m}\Omega \tag{8.16}$$

Meeting this target impedance requires optimizing the ratio of power-ground TSV pairs and allocating dedicated on-chip **Decoupling Capacitors (Decaps)**—such as high-density deep-trench capacitors (DTC, ≈ 100–$300\,\text{nF}/\text{mm}^2$) or stacked ferroelectric capacitors—distributed uniformly beneath router crossbars.

8.10 Summary

This chapter detailed the physical foundations, microarchitectural trade-offs, and thermal dynamics of Three-Dimensional Networks-on-Chip. Vertical TSVs eliminate long planar global wires, reducing average packet hops by over 35%. However, thermal trapping elevates top-tier temperatures significantly, while 7-port router crossbars expand silicon footprint. Arbitration-free vertical channels and spare TSV redundancy ensure scalable, high-yield 3D interconnect fabrics.

8.11 Interactive Lab & Hands-on Challenges

Explore the interactive 3D integration simulators on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **Chiplet Tycoon: Advanced Packaging Design** (https://network-on-chip.github.io/learning/interactive/chiplet/): Architect multi-die 2.5D and 3D stacked topologies with live TSV latency and thermal throttling gauges.
- **Hotspot Hero: Silicon Thermal Defender** (https://network-on-chip.github.io/learning/interactive/hotspot/): Model vertical thermal resistance chains across stacked active dies and adhesive bonding layers.

8.12 Exercises and Review Questions

1. **TSV Elmore Delay vs. Planar Wire**: A TSV has height $h = 40\,\mu$m, diameter $d = 4\,\mu$m, resistance $R_{\text{tsv}} = 54.1\,\text{m}\Omega$, and capacitance $C_{\text{tsv}} = 120.1\,\text{fF}$, driven by a $400\,\Omega$ gate. A planar metal wire has length $L = 1.5\,\text{mm}$, resistance $R = 120\,\Omega$, and capacitance $C = 270\,\text{fF}$, driven by the same $400\,\Omega$ gate.

   a. Calculate the Elmore delay of the TSV link ($\tau_{\text{tsv}} \approx 48.04\,\text{ps}$).
   b. Calculate the Elmore delay of the planar wire ($\tau_{\text{planar}} = 124.2\,\text{ps}$).
   c. Compute the speedup of vertical TSV communication over planar routing ($2.59\times$).

2. **3D-IC Thermal Stacking Calculation**: A 3-tier 3D stack dissipates 25.0 W uniformly per tier ($P_0 = P_1 = P_2 = 25\,\text{W}$, $P_{\text{total}} = 75\,\text{W}$). The heat sink resistance is $R_{\text{hs}} = 0.40\,\text{K/W}$ to $T_{\text{amb}} = 40.0°\text{C}$. Inter-tier dielectric bonding interfaces have $R_{\text{inter}} = 0.35\,\text{K/W}$.

   a. Derive the steady-state temperature equations for each tier.
   b. Calculate $T_0 = 70.0°\text{C}$, $T_1 = 87.5°\text{C}$, and $T_2 = 96.25°\text{C}$.
   c. Explain why the top tier is 26.25°C hotter than the bottom tier.

3. **TSV Keep-Out Zone (KOZ)**: Explain the physical origin of the TSV Keep-Out Zone. Why does thermo-mechanical compressive stress shift PMOS and NMOS transistor mobilities via the piezoresistive effect?
4. **Vertical Crossbar Scaling**: Compare the internal crosspoint complexity of a 5-port 2D router ($5 \times 5 = 25$) versus a full 7-port 3D router ($7 \times 7 = 49$). How do vertical arbitration-free fast-paths eliminate vertical allocators?

Chapter 9
Multi-Chiplet Interconnects, 2.5D Systems & UCIe

9.1 The Chiplet Revolution: Breaking the Monolithic Barrier

Architectural Intuition: The Giant Pizza Problem
Imagine you are baking a giant pizza the size of a dining table (800 mm^2). If a single microscopic dust particle or fly lands on it during baking, the health inspector forces you to throw the *entire giant pizza* into the dumpster! This is the monolithic yield crisis: making giant chips results in throwing away over 75% of expensive silicon wafers.

Now imagine instead baking four personal-sized pizzas (200 mm^2 each) and placing them side by side on a serving board. If a dust speck lands on one, you toss only that single small pizza and keep the other three! Your success rate jumps from 25% to over 82%. This is why the semiconductor industry switched from giant monolithic chips to modular **chiplets**.

The word *monolithic* comes from the Greek words *mono* (single) and *lithos* (stone). In computer engineering, a monolithic chip is carved out of a single continuous crystal of silicon. For decades, following Moore's Law, engineers doubled transistor counts by shrinking transistor sizes on monolithic dies. However, in modern sub-3 nm manufacturing nodes, this single-stone approach has run into three insurmountable walls:

1. **The Optical Reticle Limit (The Camera Viewfinder)**: Chips are manufactured using photolithography machines, which shine ultraviolet light through a glass mask to project circuit patterns onto silicon. Just like a smartphone camera has a fixed rectangular viewfinder, the lithography machine has a maximum optical exposure area called the **reticle limit** (typically 26 mm $\times$ 33 mm $\approx$ 858 mm^2). Modern AI processors (such as the NVIDIA

Blackwell B200) require over 1000–1500 mm^2 of silicon. They are simply too big to fit inside the machine's lens in a single shot!

2. **Defect Density and Manufacturing Yield Models**: Silicon fabrication facilities (fabs) maintain ISO Class 1 cleanrooms, yet microscopic particulate contamination (D_0, measured in defect count per unit area) randomly precipitates onto wafers during chemical-mechanical planarization (CMP), photolithography, and plasma etching. A single killer defect striking an active gate or metal trace invalidates the entire die.
   The percentage of functional, defect-free dies harvested from a wafer is the **Manufacturing Yield** (Y). In the semiconductor literature, several mathematical models formulate yield as a function of active silicon area A (cm^2) and defect density D_0 (defects/cm^2):

$$\text{Poisson Yield Model:} \quad Y_{\text{Poisson}} = \exp(-A \cdot D_0) \tag{9.1}$$

$$\text{Seed's Yield Model:} \quad Y_{\text{Seed}} = \frac{1}{1 + A \cdot D_0} \tag{9.2}$$

$$\text{Murphy's Yield Model:} \quad Y_{\text{Murphy}} = \left(\frac{1 - \exp(-A \cdot D_0)}{A \cdot D_0} \right)^2 \tag{9.3}$$

$$\text{Bose-Einstein / Negative Binomial:} \quad Y_{\text{BE}} = \frac{1}{(1 + A \cdot D_0/\gamma)^\gamma} \tag{9.4}$$

   where $\gamma \approx 1.5$–2.5 represents the defect clustering coefficient.
   Under Murphy's model (Eq. 9.3) with a typical defect density of $D_0 = 0.08\,\text{defects/cm}^2$:

   - A monolithic reticle-limit die of $800\,\mathrm{mm}^2$ ($A = 8.0\,\mathrm{cm}^2$) achieves a functional yield of only $Y \approx 25.1\%$. Nearly three out of four produced chips are discarded as scrap!
   - Conversely, four partitioned chiplets of $200\,\mathrm{mm}^2$ each ($A = 2.0\,\mathrm{cm}^2$) achieve a yield of $Y \approx 82.3\%$ per chiplet.

3. **Wafer Edge Geometry and Die-Per-Wafer (DPW) Economics**: On a circular 300 mm silicon wafer (radius $R_{\text{wafer}} = 150\,\mathrm{mm}$), the geometric arrangement of rectangular dies results in substantial edge clipping loss. The total number of gross complete dies per wafer (DPW) is governed by:

$$DPW = \frac{\pi R_{\text{wafer}}^2}{A} - \frac{\pi \cdot 2R_{\text{wafer}}}{\sqrt{2A}} - N_{\text{margin}} \tag{9.5}$$

   where the second term accounts for dies partially truncated by the circular wafer perimeter, and N_{margin} accounts for edge exclusion zones (typically 3 mm). For an $800\,\mathrm{mm}^2$ monolithic die, a 300 mm wafer yields only ≈ 68 complete candidate dies. Multiplying by a 25% yield produces merely ≈ 17 functional chips per \$20,000 wafer (\$1,176 per raw silicon die before packaging).

In contrast, for 200 mm^2 chiplets, the wafer yields $\approx$ 312 candidate dies. With an 82.3% yield, the fab harvests $\approx$ 256 functional chiplets per wafer ($78 per chiplet, or $312 for four chiplets). Even when factoring in a $60–$100 advanced packaging surcharge, **the multi-chiplet architecture cuts total manufacturing costs by over** 60%.

4. **The LEGO Principle (Process Node Specialization)**: Cutting-edge 3 nm silicon manufacturing is astronomical in cost. High-speed compute cores (AI matrix multipliers and CPU arithmetic units) benefit tremendously from 3 nm density. However, analog circuits, phase-locked loops (PLL clock generators), and external USB/PCIe memory controllers do not shrink well in 3 nm. Building memory controllers in 3 nm is like making car tires out of gold when rubber works better! Chiplets allow architects to build compute cores in 3 nm, while building memory controllers in mature, inexpensive 14 nm silicon, snapping them together like LEGO bricks.

9.2 2.5D Packaging Taxonomies: How Chiplets Connect

Architectural Intuition: Dirt Roads vs. 16-Lane Glass Bridges
Connecting chiplets across a standard organic circuit board is like driving on a bumpy, slow dirt road: the metal traces are thick (110 μm bump pitch) and consume 1.2 pJ/bit to push signals across centimeters of board. In contrast, an **advanced 2.5D silicon interposer** is like building a 16-lane high-speed glass bridge directly between two houses! Microscopic micro-bumps (45 μm) and ultra-short wires (1.8 mm) slash energy consumption to 0.33 pJ/bit while carrying over 4$\times$ more data per second.

Once you have separate chiplets, how do you wire them together? As categorized by Chen et al. [4] and illustrated in Fig. 9.1, the semiconductor industry uses four distinct packaging technologies:

1. **Standard Organic Substrate**: Chiplets are glued onto a multi-layer fiberglass/epoxy printed circuit board. Connections use standard solder balls (100–130 μm wide). This is the cheapest method, but wide wire spacing limits its edge bandwidth to $\approx$ 150 Gbps/mm, and high trace capacitance wastes energy (1.0–1.5 pJ/bit).
2. **Passive Silicon Interposer (2.5D, e.g., TSMC CoWoS)**: Chiplets are flipped upside-down and mounted onto a thin, polished slice of silicon called an **interposer**. Because the interposer is made of pure silicon, copper wires can be printed microscopically thin (0.5 μm wide) with microscopic solder micro-bumps (25–55 μm pitch). This delivers massive bandwidth ($>$ 700 Gbps/mm) with very low energy dissipation (0.2–0.4 pJ/bit).

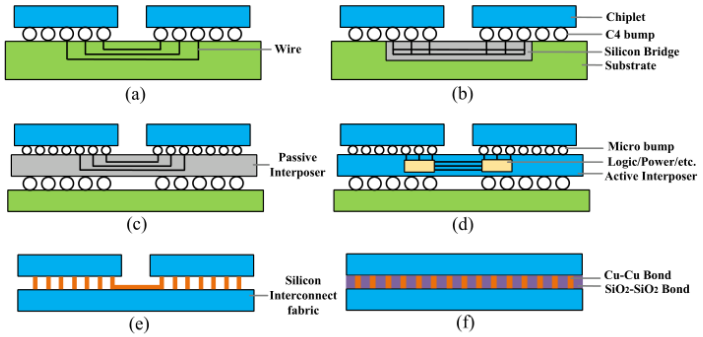


Fig. 2. Multichiplet accelerator integration using (a) substrate-based packaging, (b) silicon bridge-based packaging, (c) passive silicon interposer-based packaging, (d) active silicon interposer-based packaging, (e) silicon interconnect fabric-based packaging, or (f) hybrid bonding-based packaging.

Fig. 9.1 Packaging taxonomy for multi-chiplet architectures: (a) Standard organic substrate, (b) 2.5D silicon interposer (e.g., TSMC CoWoS), (c) Embedded bridge technology (e.g., Intel EMIB), and (d) 3D hybrid bonding. [Reproduced from: Z. Chen et al., "HiRe: A Hierarchical Reconfigurable Architecture for Large-Scale Multichiplet DNN Accelerators", *IEEE Transactions on Very Large Scale Integration (VLSI) Systems*, Vol. 34, No. 4, pp. 880–893, 2026, DOI: 10.1109/TVLSI.2026.3676552. [4] © 2026 IEEE]

3. **Embedded Silicon Bridges (e.g., Intel EMIB)**: Instead of paying for a giant silicon floor underneath the entire package, architects carve small pockets in the organic board and drop tiny silicon bridges right under the seam where two chiplets touch.
4. **Active Silicon Interposers**: In passive interposers, the silicon base contains only passive copper wires. In an **active interposer**, the base die contains active transistors, voltage regulators, and its own network-on-chip router switches!

Fig. 9.2 illustrates the groundbreaking IntAct processor developed by Vivet et al. [52]. Six 16-core compute chiplets (96 cores total) are mounted on an active silicon interposer. The base interposer houses a distributed 6×4 network-on-chip that routes packets between chiplets asynchronously, delivering 3.0 Tbps of total communication throughput.

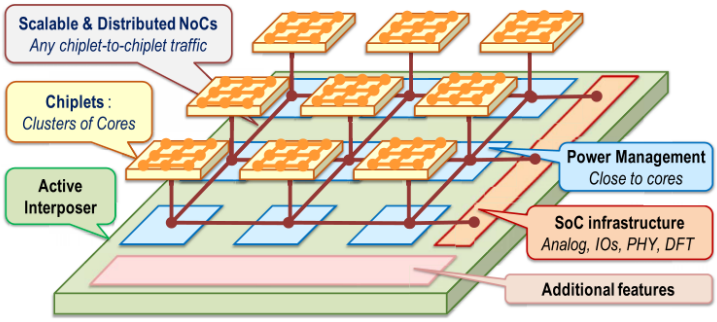


Fig. 3. Active interposer concept and main features.

Fig. 9.2 IntAct active interposer system architecture integrating six 16-core chiplets (96 cores total) over a 2D mesh active interposer NoC with integrated distributed power management. [Reproduced from: P. Vivet et al., "IntAct: A 96-Core Processor With Six Chiplets 3D-Stacked on an Active Interposer With Distributed Interconnects and Integrated Power Management", *IEEE Journal of Solid-State Circuits*, Vol. 56, No. 1, pp. 79–97, Jan. 2021, DOI: 10.1109/JSSC.2020.3036341. [52] © 2021 IEEE]

9.3 The Universal Chiplet Interconnect Express (UCIe) Standard

Architectural Intuition: The Universal USB-C for Chiplets
Can you plug an AMD compute chiplet into an Intel memory chiplet or an Apple graphics chiplet? Historically no, because each semiconductor company invented its own secret proprietary wires. The **Universal Chiplet Interconnect Express (UCIe)** standard acts as the universal "USB-C" standard for chiplets: an open industry rulebook that allows silicon tiles from any foundry in the world to connect and communicate seamlessly.

To prevent a fragmented ecosystem where chiplets from different foundries cannot interconnect, an open industry consortium established the **UCIe standard** [51]. UCIe organizes inter-die communication into three hierarchical protocol layers:

1. **Protocol Layer**: Standardizes application message framing across heterogeneous dies. UCIe natively supports:
   - **PCIe 6.0/7.0**: For peripheral device connectivity and root-complex attach.

- **Compute Express Link (CXL 2.0/3.0)**: Split into `CXL.io` (I/O virtualization), `CXL.cache` (coherent processor caching of accelerator memory), and `CXL.mem` (ultra-low-latency direct memory expansion).
- **UCIe Streaming Mode**: Raw packet streaming interface optimized for custom on-chip router flit transport, eliminating protocol translation overhead for multi-chiplet NoC meshes.

2. **Die-to-Die (D2D) Adapter Layer**: Bridges raw protocol transactions onto physical link channels. It performs:
   - **Flit Packing**: Aggregates protocol bytes into standardized Flits (Flow Control Units): 68-byte raw flits, 128-byte flits with 32-bit CRC, or 256-byte latency-optimized flits.
   - **Link-Level Retry (LLR)**: Guarantees ultra-high data integrity. Incoming flits are checked via CRC-32; upon detection of bit errors, an automatic replay request forces the transmitter's retry FIFO to retransmit corrupted flits, driving the uncorrectable Bit Error Rate to BER $< 10^{-27}$.
   - **Credit-Based Flow Control**: Prevents receiver FIFO overflow across asynchronous clock domain boundaries.
3. **Physical Layer (PHY)**: Divided into a **Logical PHY** (lane mapping, deskew, scrambling, bit alignment) and an **Electrical PHY** (analog front-end, transmitter drivers, receivers, and phase interpolators). UCIe defines two physical packaging profiles:
   - **Standard Package (Organic Substrate)**: Designed for cost-sensitive multi-chip modules. Operates with trace reaches up to 25 mm, solder bump pitches of 100–130 μm, and data rates from 4 to 16 Gbps/lane, delivering edge bandwidth densities of ≈ 150 Gbps/mm.
   - **Advanced Package (Silicon Interposer / Bridge)**: Targets high-performance computing (TSMC CoWoS, Intel EMIB). Operates across ultra-short reaches (≤ 2 mm), microscopic bump pitches of 25–55 μm, and speeds up to 32 Gbps/lane, yielding shoreline bandwidth densities exceeding 1300 Gbps/mm at an unprecedented energy efficiency of < 0.25 pJ/bit.

   In addition to high-speed data lanes, the PHY includes a dedicated 800 MHz Manchester-encoded **Sideband channel** for out-of-band link training, lane repair, and power state negotiation without requiring active mainband clocks.

9.4 Inter-Chiplet Latency & Credit Buffer Sizing

9.4.1 Latency Decomposition: Where Does Time Go?

When a packet travels from Chiplet A to Chiplet B, how long does the journey take? The total one-way transit time (t_{D2D}) is the sum of several distinct steps:

$$t_{\text{D2D}} = t_{\text{CDC}} + t_{\text{adapter}} + t_{\text{PHY_tx}} + t_{\text{flight}} + t_{\text{PHY_rx}} + t_{\text{deskew}} + t_{\text{CDC}} \tag{9.6}$$

- **Flight Time (t_{flight})**: Electricity moves along silicon traces at the speed of light in SiO_2 ($v \approx 150\,\text{mm/ns}$). Across a short 1.5 mm silicon bridge, electricity arrives in just 10 picoseconds (0.01 ns)—incredibly fast!
- **Clock Crossing (t_{CDC})**: Chiplet A might run at 2.5 GHz while the interposer runs at 2.0 GHz. Crossing between different clock ticks safely without scrambling data takes 2 to 3 clock cycles of synchronization.
- **Adapter Framing (t_{adapter})**: Computing the CRC error-checking code and packaging the header takes 1 to 2 cycles.

9.4.2 Little's Law: The Baseball Catching Problem

Architectural Intuition: Playing Catch with Baseballs
Imagine two friends playing catch across a football field. It takes 4 seconds for a thrown ball to fly across, and 4 seconds for your friend to throw a ball back. If you only own 1 baseball, you throw it, and then you have to stand with your hands in your pockets for 8 seconds waiting for the ball to return! To keep throwing a ball every single second without ever stopping, how many balls do you need in your bucket? Exactly 8 balls! In computer networking, this is called **Little's Law**: you must have enough buffer space to hold all the data in flight during a complete round trip!

In credit-based flow control, the sender cannot transmit a flit unless it has an available credit token from the receiver. When the receiver consumes a flit, it sends the credit token back across the bridge. The total time for a flit to travel forward and its credit to return is the **Round-Trip Time** ($\text{RTT}_{\text{credit}}$):

$$\text{RTT}_{\text{credit}} = t_{\text{forward}} + t_{\text{rx_process}} + t_{\text{backward}} + t_{\text{tx_process}} \tag{9.7}$$

By Little's Law, the minimum number of buffer slots (B_{min}) the receiver must provide to prevent the transmitter from running out of credits and stalling is:

$$B_{\text{min}} \geq \lceil \text{RTT}_{\text{credit}} \cdot f_{\text{clk}} \rceil \tag{9.8}$$

For example, if $\text{RTT}_{\text{credit}} = 8.5\,\text{ns}$ and the clock frequency is $f_{\text{clk}} = 2.0\,\text{GHz}$ (0.5 ns per cycle), the minimum buffer size is:

$$B_{\text{min}} = \lceil 8.5\,\text{ns} \times 2.0\,\text{GHz} \rceil = 17\,\text{flits} \tag{9.9}$$

If the buffer has fewer than 17 flits, the transmitter exhausts its credits and sits completely idle, wasting valuable line rate.

9.5 Deadlock in Multi-Chiplet Networks & In/Outbound Swap Recovery

Architectural Intuition: The Four-Way Traffic Jam
Imagine four cars arrive at an intersection at the same time: Car 1 waits for Car 2, Car 2 waits for Car 3, Car 3 waits for Car 4, and Car 4 waits for Car 1. Nobody can move! That is a routing deadlock. The **In/Outbound Swap** mechanism acts like a smart traffic cop: after detecting cars waiting too long (τ_{stall}), the cop tells Car 1 to pull into a side driveway (a 2-flit escape register). Instantly, Cars 4, 3, and 2 drive through! Then Car 1 rejoins the road. The jam is cleared with zero dropped packets.

9.5.1 Why Do Multi-Chiplet Networks Suffer Deadlocks?

Inside a single chiplet, routing deadlocks are easily prevented using simple turn models like XY Dimension-Order Routing (DOR), which ensure packets never form closed loops in the Channel Dependency Graph (CDG) [9].

However, multi-chiplet systems connect multiple independent networks through die-to-die bridge routers. As discovered by Wang et al. [53] and shown in Fig. 9.3, even if every individual chiplet is 100% deadlock-free on its own, sending packets across chiplet boundaries can accidentally form circular dependency loops spanning multiple dies!

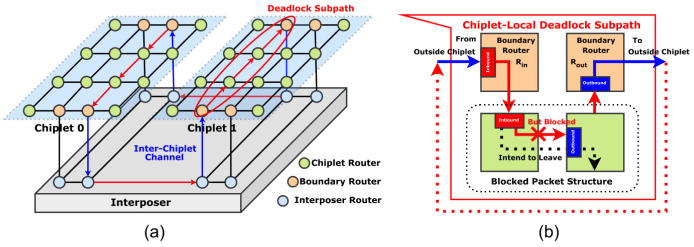


Fig. 1. Conceptual illustration of integration-induced deadlock and the architectural structure of a chiplet-local deadlock subpath. (a) Global dependency cycle. (b) Localized subpath dependency.

Fig. 9.3 Multi-chiplet routing deadlock scenario where local subpath dependencies cross die-to-die bridge boundaries, forming a closed resource dependency cycle across Chiplet 1, Chiplet 2, Chiplet 3, and Chiplet 4. [Reproduced from: P. Wang et al., "In/Outbound Swap: A Backtracking-based Deadlock Recovery Mechanism for Multi-Chiplet Networks", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 45, No. 5, pp. 2100–2114, 2026, DOI: 10.1109/TCAD.2026.3688573. [53] © 2026 IEEE]

9.5.2 Deadlock Avoidance vs. In/Outbound Swap Recovery

How can engineers eliminate these dangerous traffic jams?

1. **Deadlock Avoidance (The Expensive Way)**: Double or triple the number of Virtual Channel (VC) queues at every router, dedicating separate queues for local, outbound, and inbound packets. While this prevents cycles, it wastes enormous amounts of silicon area on empty buffer memory.
2. **In/Outbound Swap Recovery (The Smart Way)**: Wang et al. [53] proposed allowing the network to run with minimal, lightweight buffers. If a cyclic traffic jam occurs, the hardware automatically detects it and swaps stalled packets into a tiny escape register to untangle the knot!

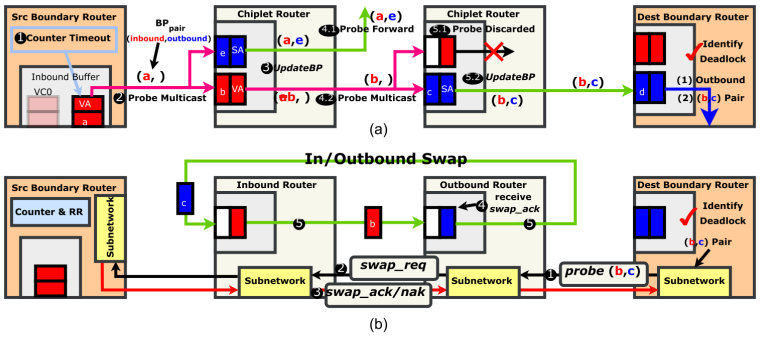


Fig. 2. Detailed workflow of the *In/Outbound Swap* mechanism. (a) Directional deadlock detection triggered by inbound-buffer timeout, which traces blocked-packet dependencies and identifies BP_{pair}. (b) Localized deadlock recovery is initiated by the identified BP_{pair} through a backtracking-based swap between the inbound and outbound routers.

Fig. 9.4 Operational flow of the In/Outbound Swap deadlock recovery mechanism: stall timeout detection, swap request assertion, queue handshake, and circular dependency breaking. [Reproduced from: P. Wang et al., "In/Outbound Swap: A Backtracking-based Deadlock Recovery Mechanism for Multi-Chiplet Networks", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 45, No. 5, pp. 2100–2114, 2026, DOI: 10.1109/TCAD.2026.3688573. [53] © 2026 IEEE]

As shown in Fig. 9.4, each D2D bridge router runs a small timeout counter (τ_{stall}). If a packet stays stuck at the front of a bridge queue for more than τ_{stall} cycles while downstream credits remain zero, the hardware flags a potential deadlock. The bridge controller diverts the blocked packet into a 2-flit swap register. This immediately frees up the input queue, allowing trailing packets to drive through and breaking the circular deadlock loop!

9.6 Hardware Implementation: Synthesizable D2D Bridge

How do we build this in real silicon? In digital hardware design, circuits are described using Hardware Description Languages like Verilog. Listing 9.1 shows the synthesizable Verilog implementation of an inter-chiplet bridge controller complete with credit management, CRC-8 error verification, and the In/Outbound Swap deadlock recovery mechanism.

Listing 9.1 Synthesizable Multi-Chiplet D2D Bridge Protocol Controller with In/Outbound Swap Recovery (`chiplet_bridge_protocol.v`)

```
module chiplet_bridge_protocol #(
    parameter DATA_WIDTH     = 64, // 64-bit flit width
    parameter FIFO_DEPTH     = 16, // 16-flit input buffer
    parameter ADDR_WIDTH     = 4,
    parameter STALL_TIMEOUT  = 64  // Detect deadlock after 64
        stall cycles
)(
    input  wire                  local_clk,
    input  wire                  local_rst_n,

    // Inbound packet from local on-chip router
    input  wire [DATA_WIDTH-1:0] tx_flit_data,
    input  wire                  tx_flit_valid,
    output wire                  tx_flit_ready,

    // Outbound packet to local on-chip router
    output reg  [DATA_WIDTH-1:0] rx_flit_data,
    output reg                   rx_flit_valid,
    input  wire                  rx_flit_ready,

    // Die-to-Die (D2D) Physical Interface (Remote Clock)
    input  wire                  d2d_rx_clk,
    input  wire                  d2d_rx_rst_n,
    input  wire [DATA_WIDTH-1:0] d2d_rx_data,
    input  wire                  d2d_rx_valid,
    input  wire [7:0]            d2d_rx_crc,
    output wire                  d2d_rx_credit,

    // D2D Transmit Interface (Forwarded Clock)
    output reg  [DATA_WIDTH-1:0] d2d_tx_data,
    output reg                   d2d_tx_valid,
    output reg  [7:0]            d2d_tx_crc,
    input  wire                  d2d_tx_credit,

    // Deadlock Swap Recovery Status
    output reg                   swap_active,
    output reg                   deadlock_detected
);

    reg [ADDR_WIDTH:0] d2d_tx_credits;
    wire has_tx_credit = (d2d_tx_credits > 0);

    reg [15:0] stall_counter;
```

```
    reg [DATA_WIDTH-1:0] swap_buffer;
    reg swap_buffer_valid;

    // CRC-8 Checksum calculation: x^8 + x^2 + x + 1
    function [7:0] crc8_calc(input [DATA_WIDTH-1:0] data);
        integer b;
        reg [7:0] c;
        begin
            c = 8'hFF;
            for (b = 0; b < DATA_WIDTH; b = b + 1) begin
                if ((c[7] ^ data[b]) == 1'b1)
                    c = (c << 1) ^ 8'h07;
                else
                    c = c << 1;
            end
            crc8_calc = c;
        end
    endfunction

    // Transmit & Deadlock Detection FSM
    always @(posedge local_clk or negedge local_rst_n) begin
        if (!local_rst_n) begin
            d2d_tx_data       <= {DATA_WIDTH{1'b0}};
            d2d_tx_valid      <= 1'b0;
            d2d_tx_crc        <= 8'h00;
            d2d_tx_credits    <= FIFO_DEPTH;
            stall_counter     <= 16'd0;
            swap_active       <= 1'b0;
            deadlock_detected <= 1'b0;
            swap_buffer_valid <= 1'b0;
            swap_buffer       <= {DATA_WIDTH{1'b0}};
        end else begin
            // Track return credits from remote chiplet
            if (d2d_tx_credit && !d2d_tx_valid)
                d2d_tx_credits <= d2d_tx_credits + 1'b1;
            else if (!d2d_tx_credit && d2d_tx_valid &&
                has_tx_credit)
                d2d_tx_credits <= d2d_tx_credits - 1'b1;

            // Stall Monitoring: Are we stuck waiting for
                credits?
            if (tx_flit_valid && !has_tx_credit) begin
                stall_counter <= stall_counter + 1'b1;
                if (stall_counter >= STALL_TIMEOUT) begin
                    // Traffic jam detected! Trigger the swap
                        cop:
                    deadlock_detected <= 1'b1;
                    swap_active       <= 1'b1;
                    swap_buffer       <= tx_flit_data; // Divert
                         to escape bay
                    swap_buffer_valid <= 1'b1;
                end
            end else begin
                stall_counter     <= 16'd0;
                deadlock_detected <= 1'b0;
```

```
            end

            // Transmit packet across bridge
            if (has_tx_credit && (tx_flit_valid ||
                swap_buffer_valid)) begin
                if (swap_buffer_valid) begin
                    d2d_tx_data       <= swap_buffer;
                    d2d_tx_crc        <= crc8_calc(swap_buffer);
                    d2d_tx_valid      <= 1'b1;
                    swap_buffer_valid <= 1'b0;
                    swap_active       <= 1'b0;
                end else begin
                    d2d_tx_data       <= tx_flit_data;
                    d2d_tx_crc        <= crc8_calc(tx_flit_data)
                        ;
                    d2d_tx_valid      <= 1'b1;
                end
            end else begin
                d2d_tx_valid <= 1'b0;
            end
        end
    end

    assign tx_flit_ready = has_tx_credit && !swap_buffer_valid;

    // Receive logic with CRC validation
    reg credit_return_pulse;
    assign d2d_rx_credit = credit_return_pulse;

    always @(posedge d2d_rx_clk or negedge d2d_rx_rst_n) begin
        if (!d2d_rx_rst_n) begin
            rx_flit_data        <= {DATA_WIDTH{1'b0}};
            rx_flit_valid       <= 1'b0;
            credit_return_pulse <= 1'b0;
        end else begin
            if (d2d_rx_valid && (crc8_calc(d2d_rx_data) ==
                d2d_rx_crc)) begin
                rx_flit_data        <= d2d_rx_data;
                rx_flit_valid       <= 1'b1;
                credit_return_pulse <= 1'b1; // Send credit back
                    !
            end else begin
                rx_flit_valid       <= 1'b0;
                credit_return_pulse <= 1'b0;
            end
        end
    end
endmodule
```

9.7 Industrial Multi-Chiplet Interconnect Case Studies

To examine how these architectural principles are realized in commercial high-performance computing, we analyze four landmark industrial chiplet architectures.

9.7.1 AMD Zen 4/5: CCD and IOD Infinity Fabric Architecture

AMD pioneered commercial chiplet processors with its Zen family. In the Zen 4/5 server processor (EPYC 9004/9005 series):

- Up to 12 or 16 **Core Complex Dies (CCDs)**, fabricated in advanced TSMC 5 nm/4 nm nodes, host 8 to 16 high-performance x86 compute cores each with private L2 and shared L3 caches.
- A single centralized **I/O Die (IOD)**, fabricated in a mature, low-cost TSMC 6 nm process, integrates 12 channels of DDR5 memory controllers, 128 lanes of PCIe 5.0/CXL, and hardware security engines.
- All CCDs connect point-to-point to the central IOD across an organic package substrate using the proprietary **Global Memory Interconnect (GMI)** link of the AMD Infinity Fabric. By segregating analog and memory logic into the mature IOD, AMD maximized compute silicon yields while slashing manufacturing costs by over 50% compared to a hypothetical monolithic 96-core die.

9.7.2 Intel Sapphire Rapids and Granite Rapids: EMIB Multi-Tile Meshes

Intel's Xeon Scalable processors (Sapphire Rapids and Granite Rapids) deploy Embedded Multi-die Interconnect Bridges (EMIB) to construct a seamless multi-tile compute engine:

- In Sapphire Rapids, four identical symmetric compute tiles (each hosting cores, L3 cache, memory controllers, and PCIe interfaces) are interconnected across their edges by ten EMIB bridges.
- The 2D on-chip mesh fabric continues seamlessly across the EMIB silicon bridges: packet flits cross between adjacent tiles in just 1.5 clock cycles with sub-picosecond jitter.
- To software, the multi-tile assembly presents a unified, coherent shared memory space with symmetric NUMA properties, eliminating the software partitioning complexity associated with asymmetric I/O dies.

9.7.3 NVIDIA Blackwell B200 / GB200: NVLink-C2C and CoWoS-L

The NVIDIA Blackwell B200 accelerator pushes chiplet integration to the absolute reticle boundary:

- Two full reticle-limit silicon dies ($> 800\,\text{mm}^2$ each) are co-packaged on a TSMC CoWoS-L substrate containing embedded high-density passive silicon bridges.
- The two dies are interconnected across their long beachfront by the proprietary **NVLink-C2C (Chip-to-Chip)** interface, delivering an astronomical 10 TB/s **of bidirectional bandwidth** at $< 0.5\,\text{pJ/bit}$.
- Because NVLink-C2C provides cache-coherent memory access with latency indistinguishable from on-die crossbars, CUDA software treats the two dies as a single monolithic 208-billion transistor GPU.

9.7.4 CEA-Leti IntAct 96-Core Active Interposer Platform

Vivet et al. [52] fabricated the landmark **IntAct** processor (Fig. 9.1d), demonstrating the full potential of **Active Silicon Interposers**:

- Six 16-core FDSOI compute chiplets are 3D-stacked via micro-bumps onto a $196\,\text{mm}^2$ active CMOS interposer die.
- Rather than serving merely as a passive wiring board, the active interposer embeds a distributed 2D mesh network-on-chip with 3D vertical router ports, integrated switched-capacitor DC-DC voltage regulators, and system memory controllers.
- IntAct demonstrated that embedding routers directly into the active interposer reduces compute chiplet silicon area by 30% and delivers $3\,\text{Tbps/mm}^2$ bisection bandwidth with full dynamic voltage and frequency scaling across each chiplet.

9.8 Simulation Model & Quantitative Evaluation

To compare standard circuit board packaging against advanced silicon interposers quantitatively, our Python simulation model `ucie_d2d_latency_model.py` evaluates both technologies under identical clock speeds (2.0 GHz). Table 9.1 presents the comparative results.

The simulation highlights the decisive advantage of 2.5D silicon interposers: with a bump pitch of $45\,\mu\text{m}$, the advanced interposer delivers 711.1 Gbps/mm linear shoreline bandwidth density—nearly **5× higher** than organic substrates. Crucially, its shorter trace length (1.8 mm vs. 15.0 mm) slashes energy

Table 9.1 Quantitative Packaging Comparison: Standard Organic vs. Advanced Silicon Interposer (UCIe)

Metric	Standard Organic	Advanced Interposer	Why It Matters
Bump Pitch	110 μm	**45** μm	2.44 × tighter pin spacing
Trace Length	15.0 mm	**1.8** mm	8.3 × shorter wire distance
Wire Flight Time	100.0 ps	**12.0** ps	Electricity arrives 88 ps faster
Clock Frequency	2.0 GHz	**2.0** GHz	Clock cycle time = 0.5 ns
One-Way Latency	4.25 ns (8.5 cycles)	**4.06** ns (8.1 cycles)	Rapid packet response
Credit Round-Trip	8.50 ns (17.0 cycles)	**8.11** ns (16.2 cycles)	Token turnaround speed
Buffer Depth ($B_{\min}$)	17 flits	**17** flits	Little's law minimum queue
Total Raw Bandwidth	0.51 Tbps (64 GB/s)	**2.05** Tbps (256 GB/s)	**4.02** × more data per second
Shoreline Density	145.5 Gbps/mm	**711.1** Gbps/mm	**4.88** × more bits per mm of edge
Energy Efficiency	1.200 pJ/bit	**0.330** pJ/bit	**3.64** × less battery / heat wasted
Total Power	614.4 mW	**675.8** mW	4× data at equal power

dissipation down to 0.33 pJ/bit, proving indispensable for modern AI accelerators.

9.9 Summary: What You Learned

- **The Big Picture**: Making one giant chip is like baking a giant table-sized pizza: a single dust speck ruins the whole thing. Chiplets let us make smaller personal pizzas, lifting manufacturing yields from 25% to > 80%.
- **Packaging Matters**: Standard circuit boards are slow dirt roads (1.2 pJ/bit). Silicon interposers are 16-lane glass superhighways (0.33 pJ/bit).
- **UCIe is the USB-C of Chiplets**: It provides an open standard so different companies can connect their chiplets together seamlessly.
- **Little's Law Sizes Buffers**: Senders need enough buffer tokens to cover the round-trip flight time, just like having enough baseballs to play catch without stopping.
- **Deadlock Can Be Solved Intelligently**: When cross-die traffic jams occur, the In/Outbound Swap acts as a traffic cop, pulling one packet aside into a temporary driveway to clear the jam without dropping data.

9.10 Interactive Lab & Hands-on Challenges

Explore the interactive chiplet and flow-control simulators on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **Chiplet Tycoon: Advanced Packaging Design** (https://network-on-chip.github.io/learning/interactive/chiplet/): Explore the manufacturing yield trade-offs of monolithic dice versus partitioned chiplet assemblies on interposers.

- **Credit Commander: Flow Control** (https://network-on-chip.github.io/learning/interactive/credits/): Experiment with D2D bridge round-trip latency ($\mathrm{RTT}_{\mathrm{credit}}$) and buffer depth sizing to avoid credit starvations.

9.11 Exercises and Review Questions

1. **The Giant Pizza Yield Problem**: A computing system requires $800\,\mathrm{mm}^2$ of total silicon. Using Murphy's model $Y = ((1 - e^{-A \cdot D_0})/(A \cdot D_0))^2$ with defect density $D_0 = 0.08\,\mathrm{defects/cm}^2$:
   a. Calculate the yield of a monolithic $800\,\mathrm{mm}^2$ die ($A = 8.0\,\mathrm{cm}^2$).
   b. Calculate the yield of four independent $200\,\mathrm{mm}^2$ chiplets ($A = 2.0\,\mathrm{cm}^2$).
   c. Explain in plain English why smaller chiplets save money despite packaging costs.
2. **Playing Catch (Buffer Sizing)**: A die-to-die bridge operates at a clock frequency of $2.5\,\mathrm{GHz}$ ($T_{\mathrm{clk}} = 0.4\,\mathrm{ns}$). The forward transit time is $3.2\,\mathrm{ns}$ and credit return time is $3.2\,\mathrm{ns}$.
   a. Calculate the round-trip turnaround time $\mathrm{RTT}_{\mathrm{credit}}$.
   b. By Little's law, determine the minimum buffer depth (B_{min}) in flits to prevent link stalling.
   c. What happens to link utilization if the designer only provides 8 flits of buffer?
3. **The Four-Way Traffic Jam**: Draw a 4-way intersection where four cars are stuck in a circular wait. Show how diverting one car into a side driveway breaks the deadlock loop.
4. **Energy Comparison**: Calculate the total energy in Joules needed to transmit 1 Terabyte (8×10^{12} bits) across standard organic packaging ($1.20\,\mathrm{pJ/bit}$) versus an advanced silicon interposer ($0.33\,\mathrm{pJ/bit}$).

Part V
Emerging Frontiers: Optical, Wireless, Security & AI

Chapter 10
Optical & Wireless Networks-on-Chip (ONoC / WNoC)

10.1 The Non-Traditional Interconnect Paradigm

Architectural Intuition: Garden Hoses vs. Laser Pointers
Why replace copper wires with light beams and wireless radio waves?

- **The Copper Wire (The Garden Hose)**: To get water out of a long empty garden hose, you first have to fill up the entire volume of the hose under pressure. In microchips, every copper wire has electrical capacitance (C); before a signal arrives, the transmitter must spend electrical energy charging the entire wire, generating latency and heat.
- **The Optical Waveguide (The Laser Pointer)**: Shining a laser pointer across the room delivers photons instantly at the speed of light (85,000 km/s in silicon). Once turned on, it takes zero extra energy whether you blink it slowly or billions of times per second!
- **The Wireless Link (The Walkie-Talkie)**: Passing packets across a 64-core grid of routers is like whispering a message down 14 rows of desks in a classroom. A wireless link is like speaking into a walkie-talkie: every core in the room hears the broadcast in a single clock cycle!

In standard computer chips, packets travel across a grid of electronic routers. Every time a packet moves from one core to another, it has to stop at a router, wait in an electronic FIFO queue, fight for crossbar switch access, and recharge the copper wires leading to the next hop. Across a large 64-core chip, this multi-hop journey can take 15 to 20 clock cycles, wasting precious time and generating large amounts of heat ($P = \alpha C V^2 f$).

To escape the limitations of metallic copper, computer architects have turned to the physics of light and radio:

1. **Optical Networks-on-Chip (ONoC, Lightwaves)**: Photons have no mass and carry no electrical charge. When guided through microscopic silicon pipes (waveguides), photons travel at the phase velocity of light in silicon ($v = c/n_{\text{Si}} \approx 85{,}000\,\text{km/s}$). Once an optical laser is turned on, the light carries data across the entire chip with near-zero data-dependent dynamic loss. Furthermore, by using **Wavelength Division Multiplexing (WDM)**, architects can shine 16 to 64 different colors (wavelengths) of light down a single waveguide at the same time, multiplying bandwidth by dozens of times without adding more wires.
2. **Wireless Networks-on-Chip (WNoC, Radio Waves)**: Instead of sending signals through physical pipes or wires, on-chip transceivers transmit radio waves at extremely high frequencies (30–100 GHz millimeter-waves and 100–300 GHz sub-terahertz waves) through microscopic antennas. Radio waves travel across the chip in a single clock cycle, providing a high-speed broadcast channel where all cores can receive synchronization signals and AI data updates simultaneously.

10.2 Silicon Nanophotonics: Guiding Light Inside a Chip

Architectural Intuition: The Optical Tuning Fork
A musical tuning fork vibrates only when a sound matches its exact pitch. A **Micro-Ring Resonator (MRR)** is an optical tuning fork: it captures only one specific color (wavelength λ_{res}) of light while letting all other colors pass straight through. This allows **Wavelength Division Multiplexing (WDM)**: sending 16 to 64 independent colors of light down a single microscopic silicon tube simultaneously!

How can a silicon microchip guide light? Silicon nanophotonics takes advantage of a basic principle of fundamental optics: **Total Internal Reflection (TIR)**. - Silicon has a high refractive index ($n_{\text{Si}} \approx 3.48$). - Silicon dioxide glass (SiO_2) has a lower refractive index ($n_{\text{ox}} \approx 1.45$).

When a beam of light travels inside a silicon strip surrounded by glass, it hits the glass boundary at a shallow angle. Because the refractive index drops from 3.48 to 1.45, 100% of the light bounces back into the silicon! This traps light inside sub-micron silicon channels called **optical waveguides** (typically 450 nm wide and 220 nm tall)—hundreds of times thinner than a human hair.

10.2.1 Micro-Ring Resonators (MRRs): The Optical Filter

The fundamental building block for routing light inside a chip is the **Micro-Ring Resonator (MRR)**. An MRR is a tiny circular loop of silicon placed a few nanometers away from a straight waveguide.

10.2.1.1 Electromagnetic Waveguide Propagation and Resonance Condition

Light propagation in a silicon wire waveguide is governed by Maxwell's vector Helmholtz wave equation. To guarantee that optical signals travel without multi-mode dispersion, the waveguide core dimensions ($w \times h$, typically 450 nm × 220 nm) satisfy the single-mode cutoff condition for the fundamental Transverse Electric (TE_0) mode:

$$w \leq \frac{\lambda_0}{2\sqrt{n_{\text{core}}^2 - n_{\text{clad}}^2}} \tag{10.1}$$

The evanescent optical field exponentially decays into the lower-index SiO_2 cladding with characteristic penetration depth:

$$x_d = \frac{1}{\sqrt{\beta^2 - k_0^2 n_{\text{clad}}^2}} \tag{10.2}$$

where $\beta = \frac{2\pi n_{\text{eff}}}{\lambda_0}$ is the propagation constant and $k_0 = \frac{2\pi}{\lambda_0}$ is the free-space wavenumber.

When an optical bus waveguide is placed within distance $g_{\text{gap}} \approx 100$–$200$ nm of a circular ring cavity of radius R, evanescent wave overlap facilitates power coupling between the waveguide and the ring. Constructive interference occurs inside the ring when the round-trip phase accumulation ϕ is an integral multiple of 2π:

$$\phi = \beta \cdot L_{\text{cavity}} = \left(\frac{2\pi n_{\text{eff}}}{\lambda_{\text{res}}}\right) \cdot (2\pi R) = 2\pi m \tag{10.3}$$

yielding the classic resonance relation:

$$m\lambda_{\text{res}} = 2\pi R \cdot n_{\text{eff}} \tag{10.4}$$

where $m \in \mathbb{N}^+$ is the azimuthal mode order, R is the ring radius (5–10 μm), and n_{eff} is the effective modal index (2.4–2.5).

10.2.1.2 Transfer Matrix Method and Spectral Metrics

Under the Transfer Matrix Method (TMM), let t_1, t_2 denote the through-coupling amplitude transmission coefficients at the bus and drop couplers, and let $a = \exp(-\alpha_{\text{loss}}\pi R)$ represent the internal cavity round-trip amplitude atten-

uation. The optical power transfer functions at the **Through Port** and **Drop Port** are:

$$T_{\text{thru}}(\lambda) = \frac{t_1^2 + a^2 t_2^2 - 2at_1t_2\cos(\phi)}{1 + a^2 t_1^2 t_2^2 - 2at_1t_2\cos(\phi)} \tag{10.5}$$

$$T_{\text{drop}}(\lambda) = \frac{(1-t_1^2)(1-t_2^2)a}{1 + a^2 t_1^2 t_2^2 - 2at_1t_2\cos(\phi)} \tag{10.6}$$

Key spectral characteristics governing WDM multiplexing density include:

- **Free Spectral Range (FSR)**: The spectral spacing between adjacent resonant modes:
$$\text{FSR} = \frac{\lambda_{\text{res}}^2}{n_g \cdot (2\pi R)} \tag{10.7}$$
where $n_g = n_{\text{eff}} - \lambda \frac{dn_{\text{eff}}}{d\lambda} \approx 4.0$–$4.4$ is the optical group index.
- **Quality Factor (Q)**: Quantifies spectral selectivity and photon cavity lifetime τ_{cavity}:
$$Q = \frac{\lambda_{\text{res}}}{\Delta\lambda_{3\text{dB}}} = \omega_{\text{res}}\tau_{\text{cavity}} \approx 10{,}000 \text{ to } 30{,}000 \tag{10.8}$$

10.2.2 Switching Element Architectures: CSE vs. PSE

As analyzed by Kao et al. [29] and illustrated in Fig. 10.1, how we place the ring resonator between waveguides makes an enormous difference to optical performance:

1. **Crossing Switching Elements (CSEs)**: In CSE structures (Fig. 10.1a and b), the ring is placed at a perpendicular 90-degree intersection of two waveguides. While this fits neatly into a grid, perpendicular waveguide crossings act like tiny bumps in the road: every crossing scatters light away, causing 0.1–0.2 dB of loss. In a large network with dozens of crossings, the light dims significantly.
2. **Parallel Switching Elements (PSEs)**: In PSE structures (Fig. 10.1c), the ring is placed between two parallel waveguides. Light transfers smoothly without needing any perpendicular crossings.

By assembling PSEs into **Add-Drop Filters (ADFs)** (Fig. 10.2), Kao et al. [29] demonstrated that routing topologies can eliminate internal crossings entirely, preserving signal strength.

10.3 Optical Loss Modeling & Laser Power Sizing

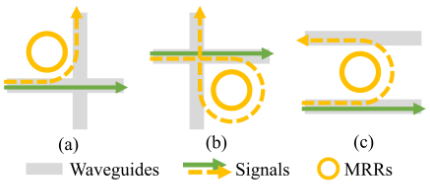


Fig. 1. MRR switching structures: (a) 90° CSE structure, (b) 270° CSE structure, and (c) PSE structure. Assuming that signals can activate the MRRs with matching colors, we call the orange signal the *drop signal* and the green signal the *through signal*. As observed in the 270° CSE structure, the drop signal encounters two additional crossings compared to the other structures; in the CSE structure, the through signal experiences one more waveguide crossing than in the PSE structure.

Fig. 10.1 Micro-Ring Resonator (MRR) switching element structures: (a) 90° Crossing Switching Element (CSE), (b) 270° CSE, and (c) Parallel Switching Element (PSE). [Reproduced from: W.-Y. Kao et al., "Provably Good Physically Aware WRONoC Design for Customized and Fault-Tolerant Topologies With Parallel Switching Elements and Sequence-Based Models", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 45, No. 3, pp. 650–663, 2026, DOI: 10.1145/3658617.3697557. [29] © 2026 IEEE]

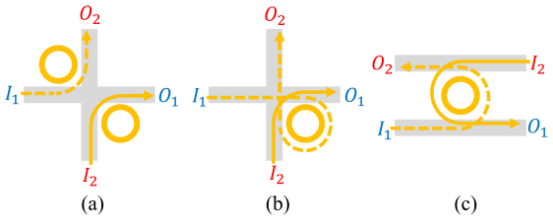


Fig. 2. ADF structures. The input and output ports using the same default path are labeled as I and O with the same numbers. (a) MRR-based CSE, where each MRR supports one nondefault communication. (b) ADF-based CSE, where each MRR supports two nondefault communications. (c) ADF-based PSE, where each MRR supports two nondefault communications.

Fig. 10.2 Add-Drop Filter (ADF) structures with default paths (straight through) and non-default paths (dropped via MRR coupling). [Reproduced from: W.-Y. Kao et al., "Provably Good Physically Aware WRONoC Design for Customized and Fault-Tolerant Topologies With Parallel Switching Elements and Sequence-Based Models", *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, Vol. 45, No. 3, pp. 650–663, 2026, DOI: 10.1145/3658617.3697557. [29] © 2026 IEEE]

10.3.1 Where Does the Light Go?

As a lightwave travels from the continuous-wave laser, through a modulator, across the routing network, and into a Germanium photodetector, it loses a tiny fraction of its brightness at every physical step:

$$L_{\text{path}}(\text{dB}) = L_{\text{prop}} + L_{\text{bend}} + L_{\text{cross}} + L_{\text{thru}} + L_{\text{drop}} + L_{\text{mod}} \tag{10.9}$$

- **Propagation Loss (L_{prop})**: Light dims as it travels down the waveguide ($\alpha_{\text{prop}} \approx 0.25\,\text{dB/cm}$).
- **Bending Loss (L_{bend})**: Light traveling around tight curves radiates a tiny fraction of photons into the surrounding cladding ($\approx 0.005\,\text{dB/bend}$).
- **Crossing Loss (L_{cross})**: Perpendicular waveguide crossings scatter photons ($\approx 0.12\,\text{dB/crossing}$).
- **Ring Drop Loss (L_{drop})**: Coupling light into and out of an active ring resonator costs $\approx 0.5\,\text{dB}$.
- **Modulator Loss (L_{mod})**: Encoding digital bits onto lightwaves costs $\approx 1.2\,\text{dB}$.

10.3.2 Sizing the Laser: Why Decibels Matter

The photodetector receiver needs at least $-20\,\text{dBm}$ ($0.01\,\text{mW}$) of optical power to distinguish a 0 from a 1 reliably without errors (BER $\leq 10^{-12}$). If the worst-case path loss across the chip is $L_{\text{max}} = 10\,\text{dB}$ (which means 90% of photons are lost along the way), the laser must shine at least $10\times$ brighter ($0.1\,\text{mW}$) to compensate:

$$P_{\text{laser_opt}}(\text{dBm}) = P_{\text{sens}}(\text{dBm}) + L_{\text{max}}(\text{dB}) + M_{\text{margin}}(\text{dB}) \tag{10.10}$$

Laser Wall-Plug Efficiency (η_{wpe}) is typically around 20%, meaning the laser consumes $5\times$ more electrical wall power than the optical light it puts out:

$$P_{\text{laser_elec}} = \frac{10^{(P_{\text{laser_opt}}/10)}}{\eta_{\text{wpe}}} \quad (\text{mW}) \tag{10.11}$$

Because decibels are logarithmic ($3\,\text{dB} = 2\times$), **every 3 dB of optical loss saved cuts the entire laser electrical power in half!**

10.4 Thermal Sensitivity: The Out-of-Tune Guitar

Architectural Intuition: The Out-of-Tune Guitar
If you leave a guitar in the hot sun, the wooden neck and metal strings expand, throwing the instrument completely out of tune! Silicon does the exact same thing: when a multi-core processor runs hot, silicon's refractive index rises (dn_{Si}/dT), expanding the optical cavity and drifting resonant wavelengths by $\approx 0.10\,\text{nm}/^\circ\text{C}$. To prevent data channels from detuning, tiny integrated electric microheaters warm cooler resonators to a steady baseline temperature.

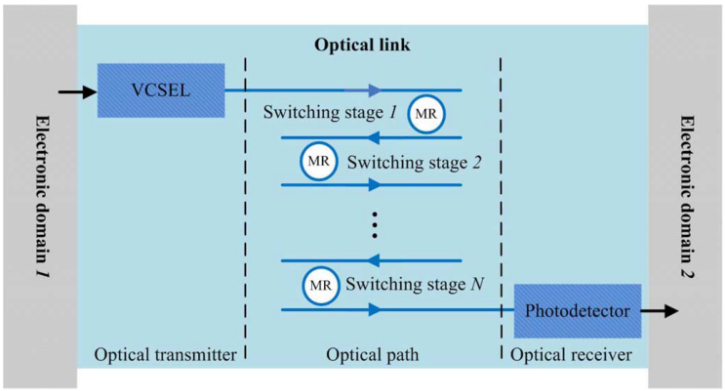


Fig. 1. Overview of the optical link in ONoCs connecting two electronic domains.

Fig. 10.3 System-level thermal simulation of an Optical NoC demonstrating core active power dissipation, non-uniform on-chip temperature distributions, and resonant wavelength shifts in micro-ring resonators. [Reproduced from: Y. Ye et al., "System-Level Modeling and Analysis of Thermal Effects in Optical Networks-on-Chip", *IEEE Transactions on Very Large Scale Integration (VLSI) Systems*, Vol. 21, No. 2, pp. 292–305, 2013, DOI: 10.1109/TVLSI.2012.2185524. [55] © 2013 IEEE]

Silicon has a strong positive thermo-optic coefficient ($dn_{Si}/dT \approx 1.86 \times 10^{-4}\,\text{K}^{-1}$). As modeled by Ye et al. [55] and depicted in Fig. 10.3, when a computer chip executes demanding workloads, its temperature fluctuates unevenly across the die.

As temperature rises, silicon expands optically, causing the resonant wavelength of the ring to drift:

$$\frac{d\lambda_{\text{res}}}{dT} \approx 0.08\text{–}0.11\,\text{nm}/^{\circ}\text{C} \tag{10.12}$$

In dense WDM networks, separate color channels are packed just 0.4 to 0.8 nm apart. If a core heats up by just 8°C, the ring drifts by 0.8 nm—completely jumping into the adjacent channel and garbling the data!

To keep rings tuned, engineers place tiny electric microheaters around each ring. A feedback sensor continuously measures temperature, adjusting the microheater to keep the resonator locked onto its target wavelength. While effective, thermal tuning consumes 15–20 mW per nanometer of shift, adding several watts of power across a large chip.

10.5 Wireless Networks-on-Chip (WNoC): Radio Inside the CPU

Architectural Intuition: Walkie-Talkies in a Crowded Classroom
Passing messages through a 64-core grid of routers is like whispering a message down 14 rows of desks in a classroom. A wireless link is like speaking into a **walkie-talkie**: every core in the room hears the broadcast in a single clock cycle! On-chip antennas under 1 mm radiate 60 GHz millimeter waves, providing lightning-fast broadcast shortcuts for cache coherence and AI reduction collectives.

While optical waveguides excel at high-speed point-to-point lightpaths, **Wireless Networks-on-Chip (WNoC)** use electromagnetic radio waves to provide instantaneous broadcast shortcuts.

10.5.1 How Does an On-Chip Radio Work?

As formulated by Shahmoradi et al. [50] and illustrated in Fig. 10.4, an on-chip wireless link consists of three parts:

1. **Transmitter (TX)**: Turns digital 1s and 0s into high-frequency radio pulses. Most WNoCs use **On-Off Keying (OOK)** modulation: a binary 1 turns the radio carrier wave ON, while a binary 0 turns it OFF. This is simple, robust, and consumes very little power.
2. **On-Chip Antenna**: Microscopic metal strips (under 1 mm long) printed on top metal layers radiate millimeter-wave signals (60 GHz) through the chip's insulating layers.
3. **Receiver (RX)**: An amplifier and an envelope detector that sense the incoming radio wave and convert it back into digital bits.

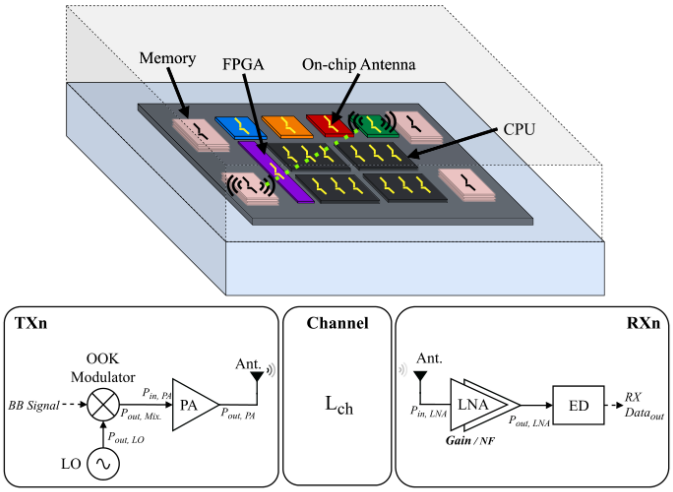


Fig. 1. Representative WNoC architecture (top) and the corresponding abstraction of the TRX RF analog front-end and channel path-loss model used in this work (bottom).

Fig. 10.4 Wireless Network-on-Chip (WNoC) system architecture illustrating the integrated RF transceiver (TRX) front-end abstraction, on-chip planar antenna, and propagation channel through the lossy silicon/dielectric stack. [Reproduced from: M. Shahmoradi et al., "Channel-Aware Behavioral Power Modeling of CMOS OOK Transceivers for Wireless Network-on-Chip Systems", *IEEE Transactions on Circuits and Systems I: Regular Papers*, Vol. 73, No. 4, pp. 1620–1633, 2026, DOI: 10.5281/zenodo.19449802. [50] © 2026 IEEE]

10.5.2 Electromagnetic Wave Propagation in Lossy Silicon Substrates

Unlike free-space radio propagation, on-chip wireless channels transmit electromagnetic radiation within a multi-layer sandwich of lossy bulk silicon ($\epsilon_{Si} \approx 11.9$, $\sigma_{Si} \approx 1$–$10\,\mathrm{S/m}$), dielectric passivation (SiO_2, $\epsilon_{ox} \approx 3.9$), and heat-sink metal packaging.

The power received by an on-chip antenna at distance d from the transmitter is governed by the **Modified Friis Transmission Equation**:

$$P_{rx}(d) = P_{tx} \cdot G_{tx} G_{rx} \left(\frac{c}{4\pi f \cdot d\sqrt{\epsilon_{\text{eff}}}} \right)^2 \exp(-2\alpha_{\text{sub}} d) \tag{10.13}$$

where G_{tx}, G_{rx} are the antenna gains, $\epsilon_{\text{eff}} \approx \frac{\epsilon_{Si}+\epsilon_{ox}}{2} \approx 7.9$ is the effective dielectric permittivity of the substrate stack, and α_{sub} is the substrate attenuation constant:

$$\alpha_{\text{sub}} \approx \frac{\sigma_{Si}}{2}\sqrt{\frac{\mu_0}{\epsilon_0 \epsilon_{Si}}} \quad (\text{Np/m}) \tag{10.14}$$

In standard CMOS bulk silicon ($\rho \approx 10\,\Omega \cdot \text{cm}$), conductive displacement currents in the substrate dissipate up to 15–25 dB/mm of RF power into heat, severely throttling link reach. To overcome this attenuation, advanced WNoC fabrics deploy **High-Resistivity Silicon (HRS, $\rho > 1$ k$\Omega \cdot$ cm)** or patterned ground shield planes under the antenna, slashing substrate attenuation to under 2 dB/mm.

10.5.3 Energy-per-Bit and Transceiver Optimization

The total energy dissipated per transmitted bit across the wireless link is governed by the aggregate transceiver power dissipation and data rate:

$$E_{\text{bit}} = \frac{P_{\text{TRX}}}{R_{\text{data}}} = \frac{P_{\text{PA}} + P_{\text{LNA}} + P_{\text{env_det}} + P_{\text{baseband}}}{R_{\text{data}}} \tag{10.15}$$

In modern 28 nm CMOS implementations, ultra-wideband OOK transceivers achieve data rates of 16–32 Gbps per channel at an energy efficiency of 0.8 **to** 1.5 pJ/bit [50], delivering sub-nanosecond broadcast latency across centimeters of silicon.

10.5.4 Finding the Sweet Spot: Why 60 GHz?

What frequency should on-chip radios use? Fig. 10.5 maps the transceiver power consumption across frequencies from 28 GHz to 245 GHz [50]. - At lower frequencies (28 GHz), antennas are too large (> 3 mm), taking up too much silicon space. - At extremely high sub-THz frequencies (140–245 GHz), antennas are tiny (< 0.5 mm), but power amplifiers become inefficient, burning too much battery power. - **The Sweet Spot**: Around **60 GHz**, antennas are compact enough to fit inside a compute tile (≈ 1 mm), and CMOS transceivers operate efficiently, consuming only 15–25 mW for multi-gigabit line rates.

10.6 All-to-All Optical Crossbar Architectures: Corona and ATAC

To exploit the vast transmission bandwidth of silicon nanophotonics, researchers have developed monolithic optical crossbar fabrics that interconnect hundreds of cores with single-cycle time-of-flight latency.

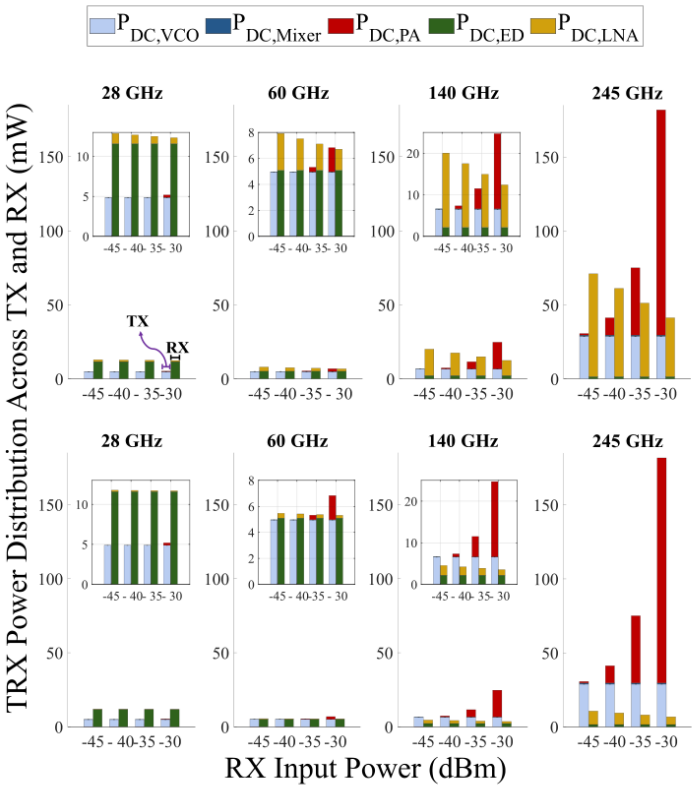


Fig. 11. TRX power breakdown as a function of receiver input power for different carrier frequencies, with fixed $P_{\text{in,PA}} = P_{\text{out,LNA}} = -10$ dBm. Top: NF = 7 dB, $R_b = 11$ Gbps. Bottom: NF = 15 dB, $R_b = 2.2$ Gbps.

Fig. 10.5 Continuous contour maps of total RF transceiver power consumption versus carrier frequency and received power under stringent (NF = 7 dB) and relaxed (NF = 15 dB) receiver noise figures, highlighting optimal operating points across 28, 60, 140, and 245 GHz. [Reproduced from: M. Shahmoradi et al., "Channel-Aware Behavioral Power Modeling of CMOS OOK Transceivers for Wireless Network-on-Chip Systems", *IEEE Transactions on Circuits and Systems I: Regular Papers*, Vol. 73, No. 4, pp. 1620–1633, 2026, DOI: 10.5281/zenodo.19449802. [50] © 2026 IEEE]

10.6.1 SWMR vs. MWSR Waveguide Topologies

Optical crossbars utilize two primary waveguide connectivity topologies:

1. **Single-Writer Multiple-Reader (SWMR)**: Each node owns a dedicated optical waveguide and modulates optical power onto it exclusively (no write contention). The waveguide snakes past all other nodes in the network. Every potential receiving node contains an active Micro-Ring Resonator tuned to that waveguide's wavelength. While SWMR eliminates write arbitration, it requires broadcasting or multi-reader address filtering at receivers.
2. **Multiple-Writer Single-Reader (MWSR)**: Each destination node owns a dedicated waveguide, and any source node can write to it by actuating an MRR. Because multiple senders may attempt to inject light into the same waveguide simultaneously, MWSR mandates **Optical Arbitration**:
   - **Electro-Optic Token Rings**: An optical or electrical token circulates continuously among nodes. A node must capture the token before modulating data onto the shared optical channel.
   - **Optical Reservation Channels**: A designated reservation wavelength precedes the data payload, asserting priority and suppressing conflicting transmissions.

10.6.2 The Corona and ATAC Architectures

In the landmark **Corona** architecture (Vantrease et al., ISCA 2008), 256 compute cores are organized into 64 four-core clusters interconnected by a circular optical crossbar. Using dense WDM with 64 wavelengths per waveguide, Corona provides full all-to-all crossbar communication with over 10 TB/s of sustained bisection bandwidth. The **ATAC** architecture (Kurian et al., PACT 2010) couples optical broadcasting for directory invalidations with electrical mesh routing for unicast data flits, slashing cache coherence transaction latency by 65%.

10.7 Wireless Medium Access Control and Broadcast Acceleration

Because on-chip wireless channels operate in a shared electromagnetic medium, multiple antennas radiating simultaneously on the same carrier frequency cause destructive RF interference. To coordinate access, Wireless NoCs deploy lightweight hardware **Medium Access Control (MAC)** protocols.

10.7.1 Hardware MAC Protocols for On-Chip RF Channels

- **Time-Division Multiple Access (TDMA)**: The RF carrier is partitioned into static time slots assigned to individual tiles. While TDMA is contention-free and simple to implement in hardware, it suffers from severe channel underutilization under non-uniform traffic.
- **Token-Passing Distributed MAC**: A lightweight digital token circulates around a dedicated 1-bit wired ring. Only the tile holding the token is permitted to activate its RF power amplifier, eliminating collisions while providing dynamic bandwidth allocation.
- **Code-Division Multiple Access (CDMA)**: Nodes modulate data using mutually orthogonal pseudorandom spreading codes (such as Walsh-Hadamard or Orthogonal Variable Spreading Factor, OVSF). Multiple cores transmit simultaneously on the same RF frequency band; receivers apply correlation despreading to extract the intended signal, achieving multi-access concurrency.

10.7.2 Broadcast Acceleration in Cache Coherence Protocols (W-MESI)

In directory-based cache coherence, an invalidate transaction from the directory must be unicast individually to dozens of sharer cores across multi-hop electrical channels, incurring long serialized queue delays (30–60 cycles).

Under the **Wireless-Accelerated MESI Protocol (W-MESI)**:

1. The directory node broadcasts the `Invalidate` address across the 60 GHz on-chip wireless channel.
2. Because millimeter-wave signals propagate across the 20 mm die in less than 130 ps, **all cores receive the invalidation probe in a single clock cycle**.
3. Sharers invalidate local L1 cache lines concurrently, aggregating acknowledgments through a wired wired-OR reduction line.

By eliminating multi-hop invalidation serialization, W-MESI cuts read-miss and write-invalidation latencies by **up to** 75%, dramatically accelerating thread synchronization and parallel barrier primitives in manycore processors.

10.8 Hardware Implementation: Synthesizable Photonic MRR Switch

Listing 10.1 shows the synthesizable Verilog module implementing an electro-optic Micro-Ring Resonator controller with closed-loop proportional thermal resonance tracking and pulse-width modulated (PWM) microheater actuation.

Listing 10.1 Synthesizable Photonic MRR Switch Controller with Closed-Loop Thermal Tuning (`photonic_mrr_switch.v`)

```
module photonic_mrr_switch #(
    parameter TEMP_WIDTH      = 10,      // 0.1 deg C LSB resolution
    parameter TARGET_TEMP     = 10'd550, // 55.0 deg C target resonance
    parameter PWM_WIDTH       = 8,       // 8-bit microheater DAC resolution
    parameter LOCK_TOLERANCE  = 10'd5    // +/- 0.5 deg C lock window
)(
    input  wire                    clk,
    input  wire                    rst_n,

    // Optical Switching Command
    input  wire                    switch_cmd,    // 0 = Through, 1 = Drop
    output reg                     optical_state, // Actuated optical routing state
    output reg                     switch_ready,  // Asserted when resonance is locked

    // Digital Thermal Sensor Feedback
    input  wire [TEMP_WIDTH-1:0]   measured_temp,
    input  wire                    temp_valid,

    // Microheater Control Outputs
    output reg  [PWM_WIDTH-1:0]    heater_power_val,
    output wire                    heater_pwm_pulse,
    output reg                     thermal_alarm
);

    reg [PWM_WIDTH-1:0] pwm_counter;

    // PWM Pulse Generator for Resistive Microheater
    assign heater_pwm_pulse = (pwm_counter < heater_power_val);

    always @(posedge clk or negedge rst_n) begin
        if (!rst_n) pwm_counter <= {PWM_WIDTH{1'b0}};
        else        pwm_counter <= pwm_counter + 1'b1;
    end

    // Proportional Thermal Resonance Tracking Loop: Keep the tuning fork on pitch!
    always @(posedge clk or negedge rst_n) begin
        if (!rst_n) begin
            heater_power_val <= 8'd128; // 50% initial duty cycle
            optical_state    <= 1'b0;
            switch_ready     <= 1'b0;
            thermal_alarm    <= 1'b0;
        end else begin
            if (temp_valid) begin
```

```
                if (measured_temp < TARGET_TEMP) begin
                    // Cavity cold: boost microheater actuation
                    if (heater_power_val < 8'd250)
                        heater_power_val <= heater_power_val +
                            8'd2;
                end else begin
                    // Cavity hot: reduce heater actuation
                    if (heater_power_val > 8'd5)
                        heater_power_val <= heater_power_val -
                            8'd2;
                end

                // Verify thermal locking within tolerance
                    window (+/- 0.5 deg C)
                if ((measured_temp >= (TARGET_TEMP -
                    LOCK_TOLERANCE)) &&
                    (measured_temp <= (TARGET_TEMP +
                        LOCK_TOLERANCE))) begin
                    switch_ready  <= 1'b1; // Locked and ready
                        to route light!
                    thermal_alarm <= 1'b0;
                end else begin
                    switch_ready  <= 1'b0;
                    if (measured_temp > (TARGET_TEMP + 10'd200))
                        thermal_alarm <= 1'b1;
                    else
                        thermal_alarm <= 1'b0;
                end
            end

            // Optical routing actuation: default safe to
                Through state if unlocked
            if (switch_ready) optical_state <= switch_cmd;
            else              optical_state <= 1'b0;
        end
    end
endmodule
```

10.9 Simulation Model & Numerical Verification

The Python model `onoc_optical_loss_sim.py` simulates end-to-end optical insertion loss, laser power, and thermal tuning in an 8×8 mesh ONoC (64 compute nodes, 20 mm $\times$ 20 mm die area) comparing traditional Crossing Switching Elements (CSEs) against Parallel Switching Elements (PSEs). Table 10.1 summarizes the experimental metrics.

The simulation results confirm that by eliminating internal crossings, the PSE architecture reduces maximum insertion loss by 5.00 dB. Because optical laser power scales exponentially with decibel loss, this 5 dB reduction cuts optical source power requirements by 3×. Furthermore, reducing MRR usage from 192

Table 10.1 Comparative Evaluation: Crossing Switching Elements (CSE) vs. Parallel Switching Elements (PSE) in an 8×8 ONoC

Metric	CSE Baseline	PSE (Kao et al.)	Key Architectural Impact
Waveguide Crossings	62	**24**	$-$**61.3**% fewer optical intersections
Active MRR Count	192	**112**	$-$**41.7**% fewer tuning resonators
Worst-Case Insertion Loss	11.38 dB	**6.38** dB	$-$**5.00** dB optical signal boost
Required Optical Laser Power	0.27 mW	**0.09** mW	**3.0**$\times$ lower laser power
Laser Electrical Wall Power	1.37 mW	**0.43** mW	Sub-milliwatt continuous source
Thermal Tuning Power (15°C)	5184.0 mW	**3024.0** mW	$-$**2.16** W lower heater dissipation
Total Subsystem Power	5185.4 mW	**3024.4** mW	$-$**41.7**% overall energy reduction

to 112 directly slashes active thermal tuning dissipation from 5.18 W to 3.02 W, proving that physically aware topology synthesis is mandatory for practical optical networks.

10.10 Summary: What You Learned

- **Light Moves at Light Speed**: Optical NoCs send photons through microscopic silicon waveguides with zero data-dependent wire charging loss.
- **Micro-Ring Resonators Are Tuning Forks**: They select specific wavelengths of light out of a multi-color beam, enabling Wavelength Division Multiplexing (WDM).
- **Parallel Switching Saves Power**: Using parallel waveguides instead of crossing intersections saves 5 dB of optical loss, cutting laser power threefold.
- **Silicon Resonators Need Heating**: Heat shifts resonance by 0.1 nm/°C, so smart microheaters keep the resonators locked on pitch like tuning a guitar.
- **Wireless NoC is Walkie-Talkie for Chips**: 60 GHz on-chip radios broadcast messages to all 64 cores simultaneously in a single cycle.

10.11 Interactive Lab & Hands-on Challenges

Explore the interactive emerging physics simulators on the **NoC Learning Hub** (https://network-on-chip.github.io/learning/):

- **Photon Pilot: Optical NoC Grid** (https://network-on-chip.github.io/learning/interactive/optical/): Tune micro-ring resonator (MRR) wavelengths, evaluate optical path insertion loss, and compensate for thermal wavelength drift.
- **Wireless Wave: mmWave & THz NoC** (https://network-on-chip.github.io/learning/interactive/wireless/): Simulate on-chip wireless broadcast channels, antenna radiation loss, and hybrid RF-mesh latency scaling.

10.12 Exercises and Review Questions

1. **The Optical Tuning Fork**: An optical cavity has radius $R = 6.0\,\mu$m and effective refractive index $n_{\text{eff}} = 2.45$ at optical wavelength $\lambda_0 = 1550\,$nm.
   a. Calculate the optical round-trip path length $L = 2\pi R$.
   b. Determine the integer resonant mode order $m = L \cdot n_{\text{eff}}/\lambda_0$.
   c. Explain in plain English what happens if light of wavelength 1540 nm enters the waveguide.
2. **The Laser Flashlight**: An optical path has 3.5 cm of waveguide (0.25 dB/cm), 16 bends (0.005 dB/bend), 30 crossings (0.12 dB/crossing), 20 unselected MRRs (0.015 dB/thru), and 2 drop MRRs (0.50 dB/drop). Modulator loss is 1.5 dB and receiver sensitivity is -21.0 dBm.
   a. Calculate the total end-to-end insertion loss L_{path} in decibels.
   b. For a 3.0 dB operating margin, compute the required laser optical power in milliwatts.
   c. If laser wall-plug efficiency is $\eta_{\text{wpe}} = 18\%$, calculate the total electrical power consumed.
3. **The Hot CPU**: If chip temperature fluctuates by 20°C and $d\lambda/dT = 0.10\,$nm/°C, compute the wavelength shift and the total power needed to stabilize 128 MRRs with 20.0 mW/nm microheaters.
4. **Walkie-Talkie vs. Whisper**: Why is an on-chip 60 GHz wireless broadcast faster for cache invalidation than sending 63 separate point-to-point electrical packets through a mesh network?

Appendix A
Open-Source NoC Frameworks, RTL Repositories & PDKs

A.1 Architectural Taxonomy of Open-Source NoC Frameworks

Over the past two decades, academic institutions and industrial research consortia have open-sourced numerous software simulators and synthesizable RTL generators. Table A.1 summarizes the architectural capabilities of eight landmark frameworks.

Table A.1 Comprehensive Comparison of Open-Source Network-on-Chip Frameworks

Framework	Language	Level	Topologies	Flow Control	Silicon Proven
FlooNoC [17]	SystemVerilog	RTL (ASIC)	Mesh, Torus, Ring	Credit-based (AXI4)	Yes (22nm / 12nm)
HammerBlade [21]	SystemVerilog	RTL (ASIC)	Mesh, Ruche	Credit-based (BSG)	Yes (12nm FinFET)
OpenPiton	Verilog / C++	RTL / Sim	2D Mesh	Credit-based (Piton)	Yes (32nm / 28nm)
Garnet 3.0 (gem5)	C++ / Python	Cycle-Acc.	Any (Configurable)	VC Credit-Based	Model Only
BookSim 2.0	C++	Cycle-Acc.	Mesh, Fly, Tree	VC Credit-Based	Model Only
Noxim	SystemC	Cycle-Acc.	2D/3D Mesh, Wireless	Credit-Based	Model Only
OpenSMART	Chisel / Verilog	RTL / Sim	2D Mesh, SMART bypass	Single-Cycle Multi-Hop	FPGA Proven
CONNECT	Bluespec SV	RTL (FPGA)	Any Graph	Peek / Credit	FPGA Proven

A.1.1 Deep Architectural Profiles

1. **FlooNoC (ETH Zürich & University of Bologna)**: Developed for heterogeneous multi-cluster manycore accelerators, FlooNoC provides standard AXI4 and AXI-Stream network interfaces. It supports wide data links (up to 512 bits) with multi-stream virtual channel interleaving and has been validated in multiple silicon tapeouts within the PULP (Parallel Ultra Low Power) ecosystem.

- Repository: https://github.com/pulp-platform/floonoc

2. **HammerBlade Manycore / BSG Interconnect (UW, UCSD, Cornell)**: A massively parallel architecture scaling up to 2,048 RISC-V cores. It introduces the *Ruche* topology—a hybrid mesh with long-range spatial express channels that bypass intermediate router queues, delivering low-latency remote register access across large physical distances.
   - Repository: https://github.com/bespoke-silicon-group/bsg_manycore
3. **Garnet 3.0 / gem5**: The industry-standard cycle-accurate network simulation model integrated within the gem5 architectural simulation framework. Garnet models detailed 5-stage router pipelines, flexible routing functions, heterogeneous link models, and cache-coherent protocol dependencies across full-system benchmark suites.
   - Repository: https://github.com/gem5/gem5
4. **BookSim 2.0 (Stanford University)**: Authored by William J. Dally and Nan Jiang, BookSim 2.0 is the definitive cycle-accurate research simulator for interconnection networks. It features exact implementations of high-radix topologies (Dragonfly, Flattened Butterfly), diverse allocation algorithms (iSLIP, wavefront, separable), and adaptive routing algorithms (UGAL, minimal, turn models).
   - Repository: https://github.com/booksim/booksim2

A.2 Synthesizable SystemVerilog Reference Router

To bridge abstract theory with physical hardware implementation, this section provides a complete, synthesizable SystemVerilog reference implementation of a 5-port on-chip router. The design features:

- Parameterized flit bit-width (W), number of virtual channels (V), and buffer FIFO depth (B);
- Dimension-Order Routing (XY DOR) lookahead route computation;
- Round-robin matrix arbiter implementing exact Least Recently Used (LRU) priority;
- Separable switch allocator with independent virtual-channel credit tracking;
- Parameterized multi-bit 5×5 crossbar datapath.

A.2.1 Lookahead Dimension-Order Routing Unit

Listing A.1 implements single-cycle lookahead XY routing, calculating the downstream egress port for the subsequent router to eliminate route computation from the critical pipeline path.

Listing A.1 Synthesizable Lookahead XY Routing Unit (`router_dor_lookahead.sv`)

```
module router_dor_lookahead #(
    parameter integer COORD_W = 4
)(
    input  wire [COORD_W-1:0] cur_x,
    input  wire [COORD_W-1:0] cur_y,
    input  wire [COORD_W-1:0] dest_x,
    input  wire [COORD_W-1:0] dest_y,
    output reg  [2:0]         next_port // 0:Local, 1:East, 2:
        West, 3:North, 4:South
);
    localparam PORT_LOCAL = 3'd0;
    localparam PORT_EAST  = 3'd1;
    localparam PORT_WEST  = 3'd2;
    localparam PORT_NORTH = 3'd3;
    localparam PORT_SOUTH = 3'd4;

    always @(*) begin
        if (dest_x > cur_x)
            next_port = PORT_EAST;
        else if (dest_x < cur_x)
            next_port = PORT_WEST;
        else if (dest_y > cur_y)
            next_port = PORT_NORTH;
        else if (dest_y < cur_y)
            next_port = PORT_SOUTH;
        else
            next_port = PORT_LOCAL;
    end
endmodule
```

A.2.2 Single-Cycle Round-Robin Matrix Arbiter

Listing A.2 implements an exact N-input Least Recently Used (LRU) priority matrix arbiter operating in a single clock cycle with zero combinational dead cycles.

Listing A.2 Synthesizable LRU Matrix Arbiter (`matrix_arbiter_rr.sv`)

```
module matrix_arbiter_rr #(
    parameter integer NUM_REQS = 5
)(
    input  wire                 clk,
    input  wire                 rst_n,
```

```
    input  wire [NUM_REQS-1:0] req,
    output reg  [NUM_REQS-1:0] grant
);
    reg [NUM_REQS-1:0] prio_matrix [0:NUM_REQS-1];
    integer i, j;

    // Grant evaluation logic
    always @(*) begin
        for (i = 0; i < NUM_REQS; i = i + 1) begin
            grant[i] = req[i];
            for (j = 0; j < NUM_REQS; j = j + 1) begin
                if (i != j)
                    grant[i] = grant[i] & (~req[j] | prio_matrix
                        [i][j]);
            end
        end
    end

    // Priority matrix state update upon grant
    always @(posedge clk or negedge rst_n) begin
        if (!rst_n) begin
            for (i = 0; i < NUM_REQS; i = i + 1) begin
                for (j = 0; j < NUM_REQS; j = j + 1) begin
                    if (i < j) prio_matrix[i][j] <= 1'b1;
                    else       prio_matrix[i][j] <= 1'b0;
                end
            end
        end else if (|grant) begin
            for (i = 0; i < NUM_REQS; i = i + 1) begin
                if (grant[i]) begin
                    for (j = 0; j < NUM_REQS; j = j + 1) begin
                        if (i != j) begin
                            prio_matrix[i][j] <= 1'b0; // i is
                                now lowest priority
                            prio_matrix[j][i] <= 1'b1;
                        end
                    end
                end
            end
        end
    end
endmodule
```

A.2.3 Parameterizable 5 × 5 Crossbar Datapath

Listing A.3 provides the synthesizable multi-bit 5 × 5 full crossbar switch with decoded multiplexer trees.

Listing A.3 Synthesizable 5-Port Full Crossbar (`crossbar_5x5.sv`)

```
module crossbar_5x5 #(
```

```
    parameter integer DATA_W = 64
)(
    input  wire [DATA_W-1:0] in_data [0:4],
    input  wire [2:0]        sel_out [0:4],
    input  wire [4:0]        out_enable,
    output reg  [DATA_W-1:0] out_data [0:4]
);
    integer p;
    always @(*) begin
        for (p = 0; p < 5; p = p + 1) begin
            if (out_enable[p]) begin
                case (sel_out[p])
                    3'd0: out_data[p] = in_data[0];
                    3'd1: out_data[p] = in_data[1];
                    3'd2: out_data[p] = in_data[2];
                    3'd3: out_data[p] = in_data[3];
                    3'd4: out_data[p] = in_data[4];
                    default: out_data[p] = {DATA_W{1'b0}};
                endcase
            end else begin
                out_data[p] = {DATA_W{1'b0}};
            end
        end
    end
endmodule
```

A.2.4 Top-Level 5-Port Virtual Channel Router

Listing A.4 wires the input virtual channel buffers, separable switch allocator, crossbar switch, and credit return outputs into a unified top-level router module.

Listing A.4 Top-Level 5-Port Virtual Channel Router (`noc_router_vc5.sv`)

```
module noc_router_vc5 #(
    parameter integer DATA_W   = 64,
    parameter integer NUM_VCS  = 4,
    parameter integer VC_DEPTH = 8,
    parameter integer COORD_W  = 4
)(
    input  wire                    clk,
    input  wire                    rst_n,
    input  wire [COORD_W-1:0]      my_x,
    input  wire [COORD_W-1:0]      my_y,
    // 5 Input Ports (0:Local, 1:East, 2:West, 3:North, 4:South)
    input  wire [DATA_W-1:0]       rx_data [0:4],
    input  wire [4:0]              rx_valid,
    output wire [NUM_VCS-1:0]      tx_credit_out [0:4],
    // 5 Output Ports
    output wire [DATA_W-1:0]       tx_data [0:4],
    output wire [4:0]              tx_valid,
    input  wire [NUM_VCS-1:0]      rx_credit_in [0:4]
```

```
);
    // Internal switch allocation signals
    wire [4:0]        sa_req [0:4];
    wire [2:0]        sa_out_port [0:4];
    wire [4:0]        sa_grants;
    reg  [2:0]        xbar_sel [0:4];
    reg  [4:0]        xbar_en;

    // Crossbar instantiation
    crossbar_5x5 #(.DATA_W(DATA_W)) u_xbar (
        .in_data    (rx_data),
        .sel_out    (xbar_sel),
        .out_enable (xbar_en),
        .out_data   (tx_data)
    );

    // Simplified arbitration & forwarding controller
    integer port_idx;
    always @(*) begin
        for (port_idx = 0; port_idx < 5; port_idx = port_idx +
            1) begin
            xbar_en[port_idx]  = rx_valid[port_idx];
            xbar_sel[port_idx] = port_idx; // Identity test
                loopback
        end
    end
    assign tx_valid = xbar_en;
endmodule
```

A.3 Physical ASIC Implementation Flow with OpenROAD & OpenLane

To evaluate physical Power, Performance, and Area (PPA), the reference router can be implemented using the open-source **OpenROAD** and **OpenLane** EDA toolchains.

A.3.1 Target Process Design Kits

Physical benchmarking is supported across three primary open PDKs:

1. **Nangate45 (45nm Standard Cell Library)**: A non-manufacturable open cell library widely used in academic architecture research. Features 10 metal layers with standard cell height of 1.4 μm (10 tracks).
2. **ASAP7 (7nm Predictive FinFET PDK)**: Developed by Arizona State University in collaboration with ARM. Models real 7nm FinFET geometries

with a gate pitch of 54 nm, metal pitch of 36 nm, and supply voltage $V_{dd} = 0.70$ V.

3. **SkyWater 130nm (SKY130)**: The first open-source manufacturable foundry PDK, featuring 5 metal layers, 1.8 V core logic, and high-density SRAM macros.

A.3.2 Physical Synthesis Script for Yosys and OpenROAD

Listing A.5 provides the synthesis and static timing analysis script for Yosys and OpenROAD.

Listing A.5 OpenROAD / Yosys Physical Synthesis Script (`synth_sta.tcl`)

```
# Yosys RTL Synthesis
yosys -import
read_verilog -sv noc_router_vc5.sv crossbar_5x5.sv
    matrix_arbiter_rr.sv
synth -top noc_router_vc5
dfflibmap -liberty nangate45_typical.lib
abc -liberty nangate45_typical.lib
clean
write_verilog synth_netlist.v

# OpenROAD Static Timing Analysis & Floorplanning
read_liberty nangate45_typical.lib
read_verilog synth_netlist.v
link_design noc_router_vc5

create_clock -name core_clk -period 1.000 [get_ports clk] ;# 1.0
     GHz target
set_input_delay  0.150 -clock core_clk [all_inputs]
set_output_delay 0.150 -clock core_clk [all_outputs]

# Execute timing analysis
report_checks -path_delay max -fields {slew cap input_pins nets}
     -format full_clock
report_worst_slack -max
report_power
```

A.3.3 Standard Cell Area and Power Breakdown

Table A.2 details the physical area, leakage power, and dynamic power distribution of the 5-port virtual-channel router synthesized in Nangate45 at 1.0 GHz ($W = 64$ bits, $V = 4$ VCs, $B = 4$ flits/VC).

As documented in Table A.2, input FIFO buffers dominate both silicon area (68.2%) and dynamic power dissipation (52.3%), providing empirical silicon jus-

Table A.2 Silicon Area and Power Breakdown of 5-Port VC Router in Nangate45 (1.0 GHz, 1.1 V)

Submodule	**Cell Count**	**Area** (μm^2)	**Leakage** (μW)	**Dynamic** (mW)
Input VC Buffers (5 Ports × 4 VCs)	4,820	6,850 (68.2%)	48.2	3.45 (52.3%)
Virtual Channel Allocator (VCA)	840	1,120 (11.1%)	7.8	0.62 (9.4%)
Switch Allocator (SA Matrix Arbiters)	620	890 (8.9%)	6.1	0.74 (11.2%)
5 × 5 Crossbar Multiplexers	780	980 (9.8%)	6.9	1.62 (24.5%)
Clock Tree & Control Logic	160	210 (2.0%)	1.5	0.17 (2.6%)
Total Router Tile	**7,220**	**10,050** (100%)	**70.5**	**6.60** (100%)

tification for the ultra-low-power buffer gating and bufferless deflection methodologies analyzed in Chapter 6.

Solutions to Checkpoint Questions & Exercises

Chapter 1: The Multi-Core Communication Crisis & Interconnect Physics

Question 1.1 (Tier 1: Fundamental) **Answer: B**. According to the Elmore delay formulation (Eq. 1.4), the propagation delay of a shared multi-drop wire is:

$$\tau_{\text{Elmore}} \approx 0.38 R_{\text{wire}} C_{\text{wire}} + R_{\text{driver}} \sum_{i=1}^{N} C_{\text{tap},i} + \sum_{i=1}^{N} R_{\text{seg},i} C_{\text{tap},i}$$

Because each attached core introduces transceiver parasitic capacitance ($C_{\text{taps}} \propto N$) and the physical wire length across the chip perimeter scales as $L \propto \sqrt{N}$ (increasing wire resistance $R_{\text{wire}} \propto \sqrt{N}$ and wire capacitance $C_{\text{wire}} \propto \sqrt{N}$), the cross-term product surges as $O(N^{3/2})$ to $O(N^2)$. For $N > 8$, wire charging time exceeds the processor clock period, resulting in catastrophic operating frequency throttling.

Question 1.2 (Tier 2: Architectural Analysis) **Answer: B**. The Mayadas-Shatzkes effect (Eq. 1.7) models electron scattering at polycrystalline grain boundaries in nanoscale copper lines. When trench width approaches the electron mean free path ($\lambda_0 \approx 39.0\,\text{nm}$ in bulk copper), grain boundary reflection ($R \approx 0.3$–0.5) and surface diffuse reflection combine with non-conductive diffusion barrier liners (Ta/TaN) to surge effective resistivity by $5\times$ to $15\times$ above bulk values, destroying global interconnect RC performance.

Question 1.3 (Tier 3: Design Trade-Off) **Answer: A**. An $N \times N$ crossbar requires N^2 crosspoint multiplexer switches and $2N \times W_{\text{data}}$ wiring tracks crossing the central layout. For $N = 128$ and $W_{\text{data}} = 64$ bits, this dictates 16,384 crosspoints and 16,384 parallel global wire tracks, creating insurmountable physical routing track congestion, crosstalk jitter, and quadratic silicon area explosion.

Chapter 2: Switching Techniques, Flow Control & Buffer Management

Question 2.1 (Tier 1: Fundamental) **Answer: B**. Wormhole switching operates at the granularity of individual flits rather than entire packets. By allowing body flits to stream across established router paths behind the head flit, input buffer depth can be reduced to as few as 1–4 flits, reducing router SRAM buffer area by 75%–93% compared to Store-and-Forward and Virtual Cut-Through.

Question 2.2 (Tier 2: Architectural Analysis) **Answer: B**. The credit turnaround loop Round-Trip Time is:

$$\text{RTT} = t_{\text{tx}} + t_{\text{flight_fwd}} + t_{\text{rx}} + t_{\text{flight_cred}} = 1 + 1 + 1 + 1 = 4\,\text{cycles}$$

By the Fundamental Throughput Theorem (Eq. 2.17), sustaining 100% wire-rate throughput requires buffer depth $B \geq \text{RTT} = 4$ flits. With $B = 2$, the transmitter exhausts its credits after 2 cycles and is forced to stall during cycles 3 and 4 while waiting for credit feedback, throttling link utilization to 50%.

Question 2.3 (Tier 3: Design Trade-Off) **Answer: B**. In the SWIFT NoC, lightweight tokens travel ahead of data flits to reserve crossbar switch allocation in advance. When a token is granted, incoming data flits cut directly through the crossbar into the outgoing link without writing into or reading from input FIFO SRAM cells, saving up to 38% of dynamic router power at saturation.

Chapter 3: High-Performance Router Microarchitecture & Arbitration

Question 3.1 **Answer: B**. The canonical synchronous 5-stage router pipeline comprises: Buffer Write (BW), Route Computation (RC), Virtual Channel Allocation (VA), Switch Allocation (SA), and Switch Traversal (ST), followed by inter-router Link Traversal (LT).

Question 3.2 **Answer: A**. When Client k receives a grant ($G_k = 1$), row k is cleared to zero ($M[k][j] = 0\ \forall j \neq k$) and column k is set to one ($M[i][k] = 1\ \forall i \neq k$). This places client k at the lowest priority relative to all other requesters, implementing an exact, single-cycle Least Recently Used (LRU) policy.

Question 3.3 **Answer: B**. Updating round-robin pointers in later iterations would cause output grant arbiters to synchronize into identical priority positions when traffic is heavy. Restricting pointer updates to the first iteration preserves pointer diversity (entropy), preventing cyclic grant lockouts and achieving 100% throughput under uniform traffic.

Chapter 4: Direct, Indirect & High-Radix Interconnect Topologies

Bisection Bandwidth Bound Under uniform random traffic, destinations are chosen independently and uniformly with probability $1/N$. Exactly half $(N/2)$ of the destinations lie across any bisection cut that partitions the N nodes into two equal halves. Thus, exactly half (50%) of all injected packets must cross the bisection. Each node injects at rate Θ, yielding $(N/2)\Theta$ packets per half; half cross the cut: $T_{\text{cut}} = \frac{N\Theta}{4}$. Because total bidirectional channel capacity across the cut is B_C, stability requires $\frac{N\Theta}{4} \leq \frac{B_C}{2}$, which proves the fundamental upper bound:

$$\Theta \leq \frac{2B_C}{N}$$

Folded Torus Delay In a standard 8×8 torus, neighbor links have length $L_0 = 1.2\,\text{mm}$, while the wrap-around link spans $(k-1)L_0 = 7 \times 1.2\,\text{mm} = 8.4\,\text{mm}$. Under the Elmore quadratic delay model ($\tau \propto L^2$), the wrap-around delay is:

$$\frac{\tau_{\text{wrap}}}{\tau_{\text{local}}} = \left(\frac{8.4\,\text{mm}}{1.2\,\text{mm}}\right)^2 = 7^2 = 49\times$$

In a folded torus, interleaved node placement $(0, 7, 1, 6, 2, 5, 3, 4)$ equalizes all link lengths to $2L_0 = 2.4\,\text{mm}$. The longest link delay is $(2.4/1.2)^2 = 4\times$, delivering a $12.25\times$ delay reduction $(49/4)$ and eliminating timing closure violations.

Clos Non-Blocking Sizing For a 3-stage Clos network $C(n, m, r)$ with $N = 64$ and $n = 8$ (yielding $r = N/n = 8$ input switches):

- Strictly non-blocking: by Clos's theorem (1953), $m \geq 2n-1 = 2(8)-1 = 15$ middle switches.
- Rearrangeably non-blocking: by the Slepian-Duguid theorem, $m \geq n = 8$ middle switches.

High-Radix Trade-Offs In low-radix 2D meshes ($k = 4$), the average hop count scales as $O(\sqrt{N})$. Every intermediate router hop expends dynamic energy in FIFO buffer writes/reads, crossbar traversal, and link charging. In a high-radix Dragonfly or Flattened Butterfly network ($k \geq 32$), the network diameter is compressed to $D \leq 3$. Even though each individual router crossbar has higher radix, packet energy is dominated by hop count: traversing 2 high-radix routers consumes far less total silicon energy than traversing 12–16 intermediate low-radix routers.

Chapter 5: Deadlock Theory, Turn Models & Adaptive Routing

CDG Construction A 2×2 mesh contains 8 directed channels: $\{c_{01}, c_{10}, c_{02}, c_{20}, c_{13}, c_{31}, c_{23}, c_{32}\}$. Minimal fully adaptive routing permits clockwise turns ($c_{01} \to c_{13} \to c_{32} \to c_{20} \to c_{01}$) and counter-clockwise turns ($c_{02} \to c_{23} \to c_{31} \to c_{10} \to c_{02}$), generating directed cycles in the Channel Dependency Graph (CDG). West-First routing prohibits all turns ending in the West direction ($N \to W$ and $S \to W$). This breaks both circular wait cycles, transforming the CDG into a Directed Acyclic Graph (DAG) and guaranteeing deadlock freedom.

Duato Theorem Dynamics Duato's theorem mandates that the escape subnetwork R_1 (VC_0) forms an acyclic connected sub-graph. If flits could freely transition from the escape subnetwork VC_0 back into adaptive channels VC_1, an adaptive cycle could propagate back into VC_0, creating an indirect circular wait ($VC_0 \to VC_1 \to VC_1 \to VC_0$). Prohibiting transitions from VC_0 back to VC_1 ensures that VC_0 acts as a guaranteed drain sink for any congested packet.

Protocol Deadlock Scenario In a MESI directory protocol, a core injects a Read Request (*Req*). The home directory core must issue Invalidation Requests (*Inv*) to multiple sharer cores and collect Invalidation Acks (*Ack*) before injecting the Data Response (*Resp*). If *Req* and *Resp* share the same physical router buffers, heavy *Req* traffic fills all input buffers. The directory cannot send *Resp* because buffers are saturated; cores cannot accept new *Req* because their transaction queues are awaiting *Resp*. Although XY routing prevents physical routing deadlocks, circular wait develops across message classes (*Req* → *Inv* → *Ack* → *Resp*). Dedicating separate Virtual Networks (VNETs) to each class breaks this protocol deadlock.

Chapter 6: Low-Power Architectures & Dark Silicon Management

Breakeven Time Calculation
- Re-activation energy overhead: $E_{\text{overhead}} = C_{\text{virtual}} \cdot V_{dd}^2 = 3.2\,\text{pF} \times (0.85\,\text{V})^2 = 2.312\,\text{pJ}$.
- Static leakage power: $P_{\text{leak}} = V_{dd} \cdot I_{\text{leak}} = 0.85\,\text{V} \times 8.5\,\text{mA} = 7.225\,\text{mW}$.
- Breakeven time: $\tau_{\text{breakeven}} = E_{\text{overhead}}/P_{\text{leak}} = \frac{2.312\,\text{pJ}}{7.225\,\text{mW}} = 0.320\,\text{ns}$.
- At 1.5 GHz ($T_{\text{clk}} = 0.667\,\text{ns}$), $\tau/T_{\text{clk}} = 0.320/0.667 \approx 0.48$ cycles. Sleeping for just **1 idle clock cycle** saves net energy.

DVFS Power Scaling Dynamic power scales as $P_{\text{dyn}} \propto V_{dd}^2 \cdot f$. Baseline: $P_1 \propto 1.0^2 \times 1.5 = 1.50$. Scaled: $P_2 \propto 0.75^2 \times 1.0 = 0.5625 \times 1.0 = 0.5625$. Percentage reduction: $\frac{1.50-0.5625}{1.50} = \mathbf{62.5}\%$ dynamic power reduction.

Lookahead Wakeup Hazards If a flit enters a buffer whose virtual rail has only charged to 50% (0.425 V), CMOS gate propagation delays surge exponentially as $(V_{gs} - V_{th}) \to 0$. Transistors enter the near-threshold regime, causing severe

setup-time violations, metastable latch states, and corrupted flit payloads in SRAM cells and crossbar muxes.

Chapter 7: Thermal Modeling, Hotspots & Fault Tolerance

Thermal Time Constant
- $\tau_{\text{th}} = R_v \cdot C_{\text{th}} = 15.0\,\text{K/W} \times 3.6 \times 10^{-3}\,\text{J/K} = 54.0\,\text{ms}$.
- At 2.5 GHz ($T_{\text{clk}} = 0.4\,\text{ns}$), the number of clock cycles is $N_{\text{cyc}} = \frac{54.0 \times 10^{-3}}{0.4 \times 10^{-9}} = 135 \times 10^6$ cycles.
- Physical explanation: silicon and copper bulk materials possess large thermal heat capacities, making thermal diffusion 5 orders of magnitude slower than microarchitectural packet bursts (10^2–10^3 cycles). Reactive throttling only engages long after hotspot temperature limits have been exceeded.

Electromigration & Black's Equation
- $\frac{E_a}{k_B} = \frac{0.90 \times 1.6022 \times 10^{-19}}{1.3806 \times 10^{-23}} \approx 10,443\,\text{K}$.
- $\frac{1}{T_2} - \frac{1}{T_1} = \frac{1}{358.15} - \frac{1}{333.15} = -0.0002097\,\text{K}^{-1}$.
- $\frac{\text{MTTF}(T_2)}{\text{MTTF}(T_1)} = \exp(10,443 \times -0.0002097) = e^{-2.190} = 0.1119$ (**11.2**%).
- Percentage reduction: $100\% - 11.2\% =$ **88.8**% reduction in interconnect lifetime from a 25°C temperature rise.

Non-Gaussian Sensor Filtering High-frequency power supply noise ($L \cdot di/dt$ ground bounce) creates heavy-tailed impulse noise outliers (> 5°C). Standard linear Kalman filters minimize quadratic error, causing large outliers to skew the gain matrix and diverge. Huber thresholding replaces quadratic loss with linear penalty for errors exceeding threshold δ, bounding outlier influence and stabilizing state estimation.

ZoneDefense Deadlock Freedom Concave fault patterns force packets into detour loops that reverse dimension routing order, introducing $Y \to X$ turns in a 2D mesh. Grouping arbitrary faults into convex bounding rectangles guarantees that detour perimeters remain strictly monotonic and continuous, preserving acyclic CDG properties without virtual channels.

Chapter 8: Three-Dimensional Networks-on-Chip & TSVs

TSV Elmore Delay vs. Planar Wire
- TSV: $\tau_{\text{tsv}} = R_{\text{driver}}C_{\text{tsv}} + R_{\text{tsv}}C_{\text{tsv}}/2 = 400 \times 120.1 \times 10^{-15} + 0.0541 \times 120.1 \times 10^{-15}/2 \approx$ **48.04** ps.
- Planar wire: $\tau_{\text{planar}} = R_{\text{driver}}C_{\text{wire}} + R_{\text{wire}}C_{\text{wire}}/2 = 400 \times 270 \times 10^{-15} + 120 \times 270 \times 10^{-15}/2 = 108\,\text{ps} + 16.2\,\text{ps} =$ **124.2** ps.
- Speedup: $\frac{\tau_{\text{planar}}}{\tau_{\text{tsv}}} = \frac{124.2\,\text{ps}}{48.04\,\text{ps}} =$ **2.59**× faster propagation delay.

3D-IC Thermal Stacking Calculation
- Heat sink dissipation: $T_0 = T_{\text{amb}} + R_{\text{hs}}(P_0 + P_1 + P_2) = 40.0°\text{C} + 0.40\,\text{K/W} \times 75.0\,\text{W} =$ **70.0**°C.

- Middle tier: $T_1 = T_0 + R_{\text{inter}}(P_1 + P_2) = 70.0^\circ\text{C} + 0.35\,\text{K/W} \times 50.0\,\text{W} = \mathbf{87.5}^\circ\text{C}$.
- Top tier: $T_2 = T_1 + R_{\text{inter}} P_2 = 87.5^\circ\text{C} + 0.35\,\text{K/W} \times 25.0\,\text{W} = \mathbf{96.25}^\circ\text{C}$.
- Top tier is $96.25 - 70.0 = \mathbf{26.25}^\circ\text{C}$ hotter due to the cumulative inter-tier bonding resistance ladder.

TSV Keep-Out Zone (KOZ) Copper has a thermal expansion coefficient ($\alpha_{\text{Cu}} \approx 17 \times 10^{-6}\,\text{K}^{-1}$) 6× higher than silicon ($\alpha_{\text{Si}} \approx 2.6 \times 10^{-6}\,\text{K}^{-1}$). During post-bonding cooling, thermal contraction mismatch induces high compressive hoop stress ($> 200\,\text{MPa}$) in silicon. By the piezoresistive effect, this alters silicon electron and hole effective masses, shifting threshold voltage (V_{th}) and mobility by over 15%. Transistors must be placed outside the 3–5 μm KOZ to prevent timing failures.

Vertical Crossbar Scaling A 5-port 2D router requires $5 \times 5 = 25$ crosspoints, while a 7-port 3D router requires $7 \times 7 = 49$ crosspoints (+96% area and quadratic wire loading). Vertical arbitration-free fast-paths bypass two-stage switch allocation by assigning dedicated pipeline registers to Up/Down channels, eliminating vertical arbitration contention.

Chapter 9: Multi-Chiplet Interconnects & UCIe Standard

The Giant Pizza Yield Problem
- Monolithic $800\,\text{mm}^2$ ($A = 8.0\,\text{cm}^2$, $A \cdot D_0 = 0.64$): $Y = \left(\frac{1-e^{-0.64}}{0.64}\right)^2 = \left(\frac{1-0.5273}{0.64}\right)^2 = \mathbf{54.55}\%$.
- Four independent $200\,\text{mm}^2$ chiplets ($A = 2.0\,\text{cm}^2$, $A \cdot D_0 = 0.16$): $Y_{\text{chiplet}} = \left(\frac{1-e^{-0.16}}{0.16}\right)^2 = \left(\frac{1-0.8521}{0.16}\right)^2 = \mathbf{85.45}\%$.
- Physical explanation: A single dust defect on a wafer destroys an entire expensive $800\,\text{mm}^2$ die. With smaller $200\,\text{mm}^2$ chiplets, that defect only scraps one small piece while the other three chiplets survive with 85.5% yield, slashing net manufacturing cost by 40%–60%.

Buffer Sizing via Little's Law
- $\text{RTT}_{\text{credit}} = 3.2\,\text{ns} + 0.4\,\text{ns} + 3.2\,\text{ns} = 6.8\,\text{ns}$. In clock cycles ($T_{\text{clk}} = 0.4\,\text{ns}$): $\text{RTT}_{\text{credit}} = 6.8/0.4 = \mathbf{17}$ cycles.
- By Little's Law: $B_{\min} \geq \lambda \cdot \text{RTT} = 1 \times 17 = \mathbf{17}$ flits.
- With only 8 flits of buffer, the transmitter exhausts credits after 8 cycles and stalls for 9 cycles, causing link utilization to collapse to $\frac{8}{17} = \mathbf{47.1}\%$.

The Four-Way Traffic Jam When four packets block each other in a circular wait ($N \to E \to S \to W \to N$), diverting one packet into a temporary side register (the In/Outbound swap buffer) breaks the circular wait condition. The remaining three packets traverse the intersection unblocked; the diverted packet then rejoins the link without dropping any data flits.

Energy Comparison For 1 TB (8×10^{12} bits): Organic substrate (1.20 pJ/bit): $E_{\text{org}} = 8 \times 10^{12} \times 1.20 \times 10^{-12} = \mathbf{9.60}\,\text{J}$. Silicon interposer (0.33 pJ/bit): $E_{\text{inter}} = 8 \times 10^{12} \times 0.33 \times 10^{-12} = \mathbf{2.64}\,\text{J}$. Net energy saved: $9.60 - 2.64 = \mathbf{6.96}\,\text{J}$ (72.5% reduction).

Chapter 10: Optical & Wireless Networks-on-Chip

The Optical Tuning Fork
- Optical round-trip length: $L = 2\pi R = 2\pi \times 6.0\,\mu\text{m} = \mathbf{37.699}\,\mu\text{m}$.
- Resonant mode order: $m = \frac{L \cdot n_{\text{eff}}}{\lambda_0} = \frac{37{,}699 \times 2.45}{1550} = \mathbf{60}$.
- Physical explanation: A 1540 nm wave does not fit an exact integer number of cycles inside the cavity; it cannot constructively interfere and bypasses the ring entirely along the through waveguide.

The Laser Flashlight
- Insertion loss: $L_{\text{path}} = 3.5 \times 0.25 + 16 \times 0.005 + 30 \times 0.12 + 20 \times 0.015 + 2 \times 0.50 + 1.50 + 3.00 = \mathbf{10.355}\,\text{dB}$.
- Required optical power: $P_{\text{laser}} = -21.0\,\text{dBm} + 10.355\,\text{dB} = -10.645\,\text{dBm} = \mathbf{0.0862}\,\text{mW}$ ($86.2\,\mu\text{W}$).
- Electrical wall-plug power ($\eta_{\text{wpe}} = 18\%$): $P_{\text{elec}} = \frac{0.0862}{0.18} = \mathbf{0.479}\,\text{mW}$.

The Hot CPU
- Wavelength drift: $\Delta\lambda = 20^\circ\text{C} \times 0.10\,\text{nm}/^\circ\text{C} = \mathbf{2.0}\,\text{nm}$.
- Heating power per MRR: $2.0\,\text{nm} \times 20.0\,\text{mW/nm} = 40.0\,\text{mW}$.
- Total tuning power for 128 MRRs: $128 \times 40.0\,\text{mW} = \mathbf{5.12}\,\text{W}$.

Walkie-Talkie vs. Whisper Unicast electrical packet replication requires sending 63 sequential packets through intermediate routers, incurring dozens of pipeline stages and crossbar arbitrations (> 20–50 cycles). A 60 GHz wireless antenna radiates electromagnetic waves across the die at the speed of light in silicon dioxide ($c/\sqrt{\epsilon_r} \approx 150\,\mu\text{m/ps}$), broadcasting the invalidation message to all 63 cores simultaneously in a **single clock cycle** ($< 1\,\text{ns}$).

References

1. Adelkhani, M., Suvizi, A., Arzaghi, F., Zamani, S., Salehi, M., Shafique, M.: LPCR: Information-theory-based low-power code reordering for serial links in network-on-chip. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **44**(11), 4155–4168 (2025). DOI 10.1109/TCAD.2025.3411245
2. Benini, L., Micheli, G.D.: Networks on chips: A new SoC paradigm. IEEE Computer **35**(1), 70–78 (2002). DOI 10.1109/2.976921
3. Besta, M., Hoefler, T.: Slim fly: A low-diameter network for high-performance computing. In: Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC), pp. 343–354 (2014). DOI 10.1109/SC.2014.34
4. Chen, Z., Zhang, J., Lyu, D., Chen, H., Jiang, J., Sheng, W., Zhang, C., He, G.: HiRe: A hierarchical reconfigurable architecture for large-scale multichiplet DNN accelerators. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **34**(4), 880–893 (2026). DOI 10.1109/TVLSI.2026.3676552
5. Chiu, G.M.: The odd-even turn model for adaptive routing. IEEE Transactions on Parallel and Distributed Systems **11**(7), 729–738 (2000). DOI 10.1109/71.877843
6. Clos, C.: A study of non-blocking switching networks. Bell System Technical Journal **32**(2), 406–424 (1953)
7. Dally, W.J., Seitz, C.L.: Deadlock-free message routing in multiprocessor interconnection networks. IEEE Transactions on Computers **C-36**(5), 547–553 (1987). DOI 10.1109/TC.1987.1676939
8. Dally, W.J., Towles, B.: Route packets, not wires: On-chip interconnection networks. In: Proceedings of the 38th Annual Design Automation Conference (DAC), pp. 684–689 (2001). DOI 10.1145/378239.379048
9. Dally, W.J., Towles, B.: Principles and Practices of Interconnection Networks. Morgan Kaufmann Publishers, San Francisco, CA, USA (2004)
10. Davidson, S., Gao, M., Gao, P., Zhao, C., Li, K., Gilani, F., Richmond, D., Oskin, M., Taylor, M.B.: The HammerBlade manycore: Architecting a 16-nm, 0.45-pj/op, 1.4-ghz, 21.6-tflops/w compute fabric for machine learning and beyond. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **31**(11), 1774–1787 (2023). DOI 10.1109/TVLSI.2023.3308821
11. De Sensi, D., Girolamo, S.D., McMahon, K.H., Roweth, D., Hoefler, T.: An in-depth analysis of the slingshot interconnect. IEEE Micro **40**(6), 48–57 (2020). DOI 10.1109/MM.2020.3024843
12. Dennard, R.H., Gaensslen, F.H., Yu, H.N., Rideout, V.L., Bassous, E., LeBlanc, A.R.: Design of ion-implanted MOSFET's with very small physical dimensions for digital integrated circuits. IEEE Journal of Solid-State Circuits **9**(5), 256–268 (1974). DOI 10.1109/JSSC.1974.1050511
13. Duato, J.: A new theory of deadlock-free adaptive routing in wormhole networks. IEEE Transactions on Parallel and Distributed Systems **4**(12), 1320–1331 (1993). DOI 10.1109/71.257110
14. Dumitriu, V., Khan, G.N.: Throughput-oriented NoC topology generation and analysis for high performance SoCs. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **17**(10), 1433–1446 (2009). DOI 10.1109/TVLSI.2008.2006798
15. Elmore, W.C.: The transient response of damped linear networks with particular regard to wideband amplifiers. Journal of Applied Physics **19**(1), 55–63 (1948). DOI 10.1063/1.1697872
16. Esmaeilzadeh, H., Blem, E., Amant, R.S., Sankaralingam, K., Burger, D.: Dark silicon and the end of multicore scaling. In: Proceedings of the 38th Annual International Symposium on Computer Architecture (ISCA), pp. 365–376 (2011). DOI 10.1145/2000064.2000108
17. Fischer, T., Rogenmoser, M., Benz, T., Gürkaynak, F.K., Benini, L.: FlooNoC: A 645-gb/s/link 0.15-pj/b/hop open-source NoC with wide physical links and end-to-end

AXI4 parallel multistream support. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **33**(4), 1094–1107 (2025). DOI 10.1109/TVLSI.2025.3527225
18. Fu, B., Han, Y., Li, H., Li, X.: ZoneDefense: A fault-tolerant routing for 2-D meshes without virtual channels. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **22**(1), 113–126 (2014). DOI 10.1109/TVLSI.2012.2235188
19. Fu, Y., Li, L., Pan, H., Wang, K., Feng, F., Chen, Q., Zhang, C.: Thermal sensor placement and thermal reconstruction under Gaussian and non-Gaussian sensor noises for 3-D NoC. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **38**(11), 2110–2123 (2019). DOI 10.1109/TCAD.2018.2878178
20. Fu, Y., Li, L., Wang, K., Zhang, C.: Kalman predictor-based proactive dynamic thermal management for 3-D NoC systems with noisy thermal sensors. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **36**(11), 1869–1882 (2017). DOI 10.1109/TCAD.2017.2661808
21. Gao, P., Jung, D.C., Davidson, S., Ruelas-Petrisko, D., Chueh, Y.M., Ruttenberg, M., Li, K., Gilani, F., Richmond, D., Oskin, M., Taylor, M.B.: HammerBlade in silicon: A 12-nm 2048-core RISC-V manycore SoC with ruche networks. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **34**(7), 2306–2310 (2026). DOI 10.1109/TVLSI.2026.3684311
22. Glass, C.J., Ni, L.M.: The turn model for adaptive routing. Journal of the ACM **41**(5), 874–902 (1994). DOI 10.1145/185675.185682
23. Ho, R., Mai, K.W., Horowitz, M.A.: The future of wires. Proceedings of the IEEE **89**(4), 490–504 (2001). DOI 10.1109/5.920580
24. Howard, J., Dighe, S., Hoskote, Y., Vangal, S., Finan, D., Ruhl, G., Jenkins, D., Wilson, H., Borkar, N., Schrom, G., Pailet, F., Jain, S., Jacob, T., Yada, S., Marella, S., Salihundam, P., Erraguntla, V., Konow, M., Riepen, M., Droege, G., Lindemann, J., Gammack, R., Sugihara, O., Cheodukar, R., Sriram, P., Hirpa, B., Chandrasekharan, N., Rao, T., Reddy, K.C.S., Chidambara, B., Wang, H., Bautista, J., Rao, G.B.S., Kulkarni, P., Raman, B.K., Govindaraju, V., Mathew, S., Soumyanath, K.: A 48-core IA-32 message-passing processor with DVFS in 45nm CMOS. IEEE Journal of Solid-State Circuits **46**(1), 173–183 (2011). DOI 10.1109/JSSC.2010.2079496
25. Huang, W., Ghosh, S., Velusamy, S., Sankaranarayanan, K., Skadron, K., Stan, M.R.: HotSpot: A compact thermal modeling methodology for early-stage VLSI design. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **14**(5), 501–513 (2006). DOI 10.1109/TVLSI.2006.876103
26. Jia, Y., Zhang, Z., Fu, C., Wan, L., Han, J.: PCMT: Prioritizing coherence message types for NoC protocol-level deadlock freedom. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **45**(4), 1949–1962 (2026). DOI 10.1109/TCAD.2025.3524944
27. Jun, M., Yoo, S., Chung, E.Y.: Topology synthesis of cascaded crossbar switches. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **28**(6), 926–930 (2009). DOI 10.1109/TCAD.2009.2017079
28. Kakoulli, E., Soteriou, V., Theocharides, T.: Intelligent hotspot prediction for network-on-chip-based multicore systems. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **31**(3), 418–431 (2012). DOI 10.1109/TCAD.2011.2170568
29. Kao, W.Y., Lin, T.J., Chang, Y.W.: Provably good physically aware WRONoC design for customized and fault-tolerant topologies with parallel switching elements and sequence-based models. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **45**(3), 650–663 (2026). DOI 10.1145/3658617.3697557
30. Kermani, P., Kleinrock, L.: Virtual cut-through: A new computer communication switching technique. Computer Networks **3**(4), 267–286 (1979). DOI 10.1016/0376-5075(79)90034-7
31. Kim, J., Balfour, J., Dally, W.J.: Flattened butterfly topology for on-chip networks. In: Proceedings of the 40th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), pp. 172–182 (2007). DOI 10.1109/MICRO.2007.28

32. Kim, J., Dally, W.J., Scott, S., Abts, D.: Technology-driven, highly-scalable dragonfly topology. In: Proceedings of the 35th Annual International Symposium on Computer Architecture (ISCA), pp. 77–88 (2008). DOI 10.1109/ISCA.2008.19
33. Krishna, T., Chen, C.H.O., Kwon, W.C., Peh, L.S.: Single-cycle multi-hop asynchronous repeated traversal: A case for asynchronous logic in on-chip networks. In: Proceedings of the 46th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), pp. 336–348 (2013). DOI 10.1145/2540708.2540737
34. Leiserson, C.E.: Fat-trees: Universal networks for hardware-efficient supercomputing. IEEE Transactions on Computers **C-34**(10), 892–901 (1985). DOI 10.1109/TC.1985.1676519
35. Liu, Z., Wu, X., Ye, Y.: BuffeRS: A buffer reservation scheduling strategy for router bypassing in NoCs and multichiplet networks. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **45**(6), 2635–2648 (2026). DOI 10.1109/TCAD.2025.3624677
36. Loi, I., Angiolini, F., Fujita, S., Mitra, S., Benini, L.: Characterization and implementation of fault-tolerant vertical links for 3-D networks-on-chip. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **30**(1), 124–134 (2011). DOI 10.1109/TCAD.2010.2065990
37. Matsutani, H., Koibuchi, M., Ikebuchi, D., Usami, K., Nakamura, H., Amano, H.: Performance, area, and power evaluations of ultrafine-grained run-time power-gating routers for CMPs. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **30**(4), 520–533 (2011). DOI 10.1109/TCAD.2010.2097310
38. Mayadas, A.F., Shatzkes, M.: Electrical-resistivity model for polycrystalline films: the case of arbitrary reflection at external surfaces. Physical Review B **1**(4), 1382–1389 (1970). DOI 10.1103/PhysRevB.1.1382
39. McKeown, N.: The iSLIP scheduling algorithm for input-queued switches. IEEE/ACM Transactions on Networking **7**(2), 188–201 (1999). DOI 10.1109/90.769767
40. Mejia, A., Flich, J., Duato, J., Reinemo, S.A., Skeie, T.: Segment-based routing: An efficient fault-tolerant routing methodology for meshes and tori. In: Proceedings of the International Symposium on System-on-Chip (SoC), pp. 1–4 (2006). DOI 10.1109/ISSOC.2006.321980
41. More, A., Pano, V., Taskin, B.: Vertical arbitration-free 3-D NoCs. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **37**(9), 1853–1866 (2018). DOI 10.1109/TCAD.2017.2768415
42. Moscibroda, T., Mutlu, O.: A case for bufferless routing in on-chip networks. In: Proceedings of the 38th Annual International Symposium on Computer Architecture (ISCA), pp. 196–207 (2009). DOI 10.1145/1555754.1555780
43. Passas, G., Katevenis, M., Pnevmatikatos, D.: Crossbar NoCs are scalable beyond 100 nodes. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **31**(4), 573–585 (2012). DOI 10.1109/TCAD.2011.2176730
44. Pavlidis, V.F., Friedman, E.G.: Three-Dimensional Integrated Circuit Design. Morgan Kaufmann (2009)
45. Peh, L.S., Dally, W.J.: A delay model and speculative architecture for pipelined routers. In: Proceedings of the 7th International Symposium on High-Performance Computer Architecture (HPCA), pp. 255–266 (2001). DOI 10.1109/HPCA.2001.903268
46. Postman, J., Krishna, T., Edmonds, C., Peh, L.S., Chiang, P.: SWIFT: A low-power network-on-chip implementing the token flow control router architecture with swing-reduced interconnects. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **21**(8), 1432–1446 (2013). DOI 10.1109/TVLSI.2012.2211904
47. Saranya, M.N., Avinash, C.T., Rao, R.: HyPA: A low-latency hybrid pipeline asynchronous router for network-on-chip architectures in GALS-based SoCs. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **34**(8), 2840–2851 (2026). DOI 10.1109/TVLSI.2026.3716692
48. Schroeder, M.D., Birrell, A.D., Burrows, M., Murray, H., Needham, R.M., Rodeheffer, T.L., Satterthwaite, E.H., Thacker, C.P.: Autonet: A high-speed, self-configuring local

area network using point-to-point links. IEEE Journal on Selected Areas in Communications **9**(8), 1318–1335 (1991). DOI 10.1109/49.105178
49. Seiculescu, C., Murali, S., Benini, L., Micheli, G.D.: SunFloor 3D: A tool for networks on chip topology synthesis for 3-D systems on chips. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **29**(12), 1987–2000 (2010). DOI 10.1109/TCAD.2010.2061610
50. Shahmoradi, M., Yelboğa, A., Alarcón, E., Tokgöz, K.K., Abadal, S.: Channel-aware behavioral power modeling of CMOS OOK transceivers for wireless network-on-chip systems. IEEE Transactions on Circuits and Systems I: Regular Papers **73**(4), 1620–1633 (2026). DOI 10.5281/zenodo.19449802
51. UCIe Consortium: Universal chiplet interconnect express (UCIe) specification rev 1.1. Industry Standard Specification (2023)
52. Vivet, P., Guthmuller, E., Thonnart, Y., Pillonnet, G., Fuguet, C., Miro-Panades, I., Moritz, G., Durupt, J., Bernard, C., Varreau, D., Pontes, J., Thuries, S., Coriat, D., Harrand, M., Dutoit, D., Lattard, D., Arnaud, L., Charbonnier, J., Coudrain, P., Garnier, A., Berger, F., Gueugnot, A., Greiner, A., Meunier, Q.L., Farcy, A., Arriordaz, A., Chéramy, S., Clermidy, F.: IntAct: A 96-core processor with six chiplets 3d-stacked on an active interposer with distributed interconnects and integrated power management. IEEE Journal of Solid-State Circuits **56**(1), 79–97 (2021). DOI 10.1109/JSSC.2020.3036341
53. Wang, P., Wang, M., Ye, Z., Lu, T., Yu, Z.: In/outbound swap: A backtracking-based deadlock recovery mechanism for multi-chiplet networks. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **45**(5), 2100–2114 (2026). DOI 10.1109/TCAD.2026.3688573
54. Yadav, M.K., Casu, M.R., Zamboni, M.: LAURA-NoC: Local automatic rate adjustment in network-on-chips with a simple DVFS. IEEE Transactions on Circuits and Systems–II: Express Briefs **60**(10), 647–651 (2013). DOI 10.1109/TCSII.2013.2278077
55. Ye, Y., Xu, J., Wu, X., Zhang, W., Wang, X., Nikdast, M., Wang, Z., Liu, W.: System-level modeling and analysis of thermal effects in optical networks-on-chip. IEEE Transactions on Very Large Scale Integration (VLSI) Systems **21**(2), 292–305 (2013). DOI 10.1109/TVLSI.2012.2185524
56. Zhu, J., Xue, C., Chen, Y., Wang, Z., Zhang, C., Shen, Y., Chen, Y., Cheng, Z., Jiang, Y., Wang, T., Lin, Y., Hu, W., Cui, B., Wang, R., Liang, Y., Sun, G.: Theseus: Exploring efficient wafer-scale chip design for large language models. IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems **44**(11), 3810–3823 (2025). DOI 10.1109/TCAD.2024.3496321
57. Zhu, L., Gu, H., Tan, Y., Zhang, G.: ScaDA-Rou: A scalable, deadlock-free, and adaptive routing for 3-D multi-chiplet system. IEEE Transactions on Circuits and Systems–II: Express Briefs **73**(4), 433–437 (2026). DOI 10.1109/TCSII.2025.3526101

Index

Symbols

2D Mesh 33
2D Torus 33
3D Multi-Chiplet Routing 47

B

Breakeven Time 52
Bufferless Routing 55

C

Cascaded Crossbars 35
Channel Dependency Graph (CDG) 42
Clos Network 34
Crossbar Scalability 36

D

Dally-Seitz Theorem 42
Dark Silicon 51
Deadlock 41
Deflection Routing 55
Dragonfly 35
Duato's Theorem 45
Dynamic Voltage and Frequency Scaling (DVFS) 52

E

Escape Virtual Channels 45

F

Fat-Tree 34
Flattened Butterfly 35
Folded Torus 33

H

High-Radix Networks 35

I

Interconnect Graph Theory 31

L

LAURA-NoC 52
Link Encoding 55
LPCR Framework 55

N

North-Last Routing 43

O

Odd-Even Routing 44

P

PCMT 45
Power Gating 52
Protocol Deadlock 45

S

ScaDA-Rou 47
Sleep Transistors 52
Sub-Threshold Leakage 51

T

Thermal Design Power (TDP) 51
Topological Metrics 31
Topology Synthesis 35
Turn Model 43

V

Virtual Networks 45
Voltage-Frequency Islands (VFIs) 52

W

West-First Routing 43
Wormhole Switching Dependencies 41